



Exploring Math  with WEB REDUCE

Statistics with REDUCE



Dr. Wolfgang Lindner
Dr.W.G.Lindner@gmail.com
Leichlingen, Germany
2026

Contents

Preface

0.1	Why REDUCE for Statistics?	3
0.2	Some Hints for the Use of WEB REDUCE	6
1	Descriptive Statistics	6
1.1	<code>mean</code> - the arithmetic Mean	7
1.2	<code>var</code> - the variance	11
1.3	<code>sd</code> - the Standard deviation	13
1.4	<code>sem</code> - the Standard Error of Mean	16
1.5	<code>mad</code> - the average absolute deviation	18
1.6	<code>rms</code> - the root mean square	21
1.7	<code>median</code> - the Median	23
1.8	<code>mode</code> - the Mode	25
1.9	<code>quantile</code> - the Quantile	27
1.10	<code>moment</code> - the r^{st} Moment	30
1.11	<code>skew</code> - the Skew.ness	33
1.12	<code>kurtosis</code> - the Kurtosis	37
1.13	<code>cov</code> - the Covariance	42
2	Discrete distributions	45
2.1	Binomial distribution	45
2.2	Geometric distribution	50
2.3	Negative binomial distribution	53
2.4	Hypergeometric distribution	56
2.5	POISSON distribution	61
3	Continuous distributions	65
3.1	Normal distribution	66
3.2	Exponential distribution	72
3.3	Student's t -distribution	76
3.4	SNEDECOR's F distribution	82
3.5	Chi-Square distribution	87
3.6	PARETO distribution	92
3.7	WEIBULL distribution	95
4	Test Statistics	99
4.1	One Sample Z-Test alias GAUSS test	99
4.2	Two Sample Z-Test	101
4.3	One Sample t -Test	104
4.4	Two Sample t -Test	107
4.5	Paired t -Test alias Differences t -Test	110
4.6	Chi-Squared Test on Variance	113

4.7	F test	116
4.8	One Sample Sign Test	120
4.9	Two Sample Sign Test	123
4.10	One Sample WILCOXON Test	126
4.11	Two Sample WILCOXON Test	129
4.12	MANN–WHITNEY U test	132
4.13	PEARSON’s Chi-squared test & Contingency Tables	135
4.14	FISHER test	138
4.15	MCNEMAR test	141
5	Correlation and Bootstrap	144
5.1	PEARSON’s ρ Correlation coefficient	144
5.2	SPEARMAN’s ρ_S rank correlation coefficient	147
5.3	KENDALL’s τ rank correlation coefficient	150
5.4	ICC - Intraclass Correlation Coefficient	153
5.5	regression - the linear regression line	156
5.6	anova1 - One-way Analysis Of Variance	159
5.7	boot1 - the bootstrap method for dependent samples	162
5.8	boot2 - the bootstrap method for independent samples	165
5.9	bootCI - Confidence Interval using bootstrap	168
5.10	jackknife - The Jackknife method	172
6	Bibliography	176

0.1 Preface

This collection of small scripts shows the use of the free CAS REDUCE to implement some well known concepts and routines of elementary statistics. We focus on comprehension of statistical concepts by direct translation of mathematical formulas, qualitative simple pictures and worked examples. This makes this ‘handbook’ and the accompanying REDUCE worksheets a perfect candidate for a learning environment at an ‘action and process level’ in the sense of the APOS Theory, cf. the short summary in [12].

The statistical concepts are presented in a CAS-language, that lies between the semiprofessional slang, in which mathematical concepts are presented and roughly explained and the high precision of the formal mathematical language, which is not so easy to grasp at a first attempt. The REDUCE language allows to *make the concepts of statistics executable*, to coin math formulas and processes into ‘run’able functions/procedures and therefore to allow own experiments. The results of one’s thinking in the CAS language are immediately returned to the screen and help to check the right understanding.

Why statistics with REDUCE?

Professional and university statisticians use free software like R, PYTHON, OCTAVE or free spreadsheets like LIBREOFFICE CALC etc. Besides possible problems of installation of the software by the novice user, it is often not easy to look into the source code of the relevant procedures, because they are cluttered into diverse packages or dependency structures and are not useable stand-alone. So, here are some reasons for the use of REDUCE ...

- allows to program statistics formulas very close to their mathematical formulation,
- *each worksheet is totally independent* of other domain-specific code, its self-contained, because the *prerequisites* for running the code of the worksheet are explicitly listed in its introductory section (which is also an opportunity for learning the concepts),
- it’s easy to include comments, tests, tables and plots in the worksheet,
- it is just fun and illuminating to use REDUCE in it’s new WEB REDUCE version.

Some remarks on the content of this booklet

The order of the following suite of worksheets with definitions, checks, exercises (problems) etc. follows a standard presentation.

1 First we define the standard functions of descriptive statistics, i.e. `mean`, `var`, `std`, `sem`, `mad`, `rms`, `median`, `mode`, `madM`, `quantile`, `moment`, `skew`, `kurtosis`. Ergo, the user should be able to easily follow the calculations in standard texts like [18] or [8].

2 We implement some important statistical distributions as helper functions to enable the statistical tests in chapter §4 and §5. For each distribution presented in chapter §2 and chapter §3 we give

- the definition resp. coding in REDUCE notation for
 - the probability density functions f named *...PDF*,
 - the cumulative distribution functions F named *...CDF* and
 - the quantile functions F^{-1} named *...INV*, i.e. the INVerse of the CDF.
- often a check of the calculations against the statistics software R or MATLAB,
- a short table of the distributions values,
- a qualitative plot to have a visual impression or to verify one's result on a graph,
- solution of a prototypical problem resp. application followed by a set of exercises.

To run a REDUCE sheet click on a link starting with  View *Code*, run REDUCE.

3 Then we implement some standard functions of test statistics, i.e. *Gauss tests*, *STUDENT t -tests*, *F-test*, *ChiSquare test*, *WILCOXON tests*, etc. The user should easily follow these calculations in standard texts like [18] or [8]. All these calculations are demonstrated and recalculated as examples using REDUCE functions.

4 Last we implement some rank'ing tests like *SPEARMAN rank correlation*, *do regression* and *one-way ANOVA* and close with *bootstrap* and *jackknife* methods. Again, these calculations are demonstrated and calculated using small self-made REDUCE functions.

Some remarks on the didactical concept of this booklet

This small booklet is dedicated to the novice – w.r.t. the subject matter as well as w.r.t. the use of CAS REDUCE. Therefore, I roughly follow the educational APOS theory, cf. [12]: this means each section of this book is divided into the following 4 steps:

- M:** we start with a short *motivation* of the mathematical concept using colloquial *words*.¹
- V:** we present an adequate concrete *visualization* of the mathematical concept using a prototypical example. The user should pause and think and reflect a while about it.
- D:** we give a precise mathematical *definition* using mathematical *symbols*. This definition should be memorized along with the visualization and the prototypical example.

¹Some of these sentences are produced using AI answers. Therefore, I would appreciate any information regarding the original source of the citation and would add the appropriate credit to the quote.

E: we solve a concrete *example* often w.r.t. part **V** and prepare for the use of the CAS REDUCE. The reader will fully understand the mathematical concept in question when he can translate the mathematical definition **D** into a working code snippet. This simultaneously enables to conduct own experiments with own data, allowing to *observe the effects of own actions* in the CAS REDUCE.

Remark. The presented code in this booklet is ‘educational’ code, aimed at the novice learner who is willing to learn statistics and a CAS like REDUCE simultaneously. I have tried to program all concepts *as simply² as possible*, therefore there is *no guaranty for correct results using the presented statistical functions*. Nevertheless, our code is checked with the professional industry strength code from MATLAB (online) aka OCTAVE or the academical strength code from R.

There is currently no statistics package in REDUCE. Looking at the Wester-tests [30, p.43] or [33, p.19] we see that the whole problem group *E1* until *E16* is not yet solved with REDUCE functions. This booklet tries to fill these gaps and also present additional functions.

For inspecting or running a REDUCE script no installation is necessary, *everything runs directly online*: a click on a link like ‘code’



View Code, run REDUCE.

in this text is enough to view the corresponding script and (after copying parts of it) execute it in WEB REDUCE.³ To profit from these features, we benefit from the excellent online version of REDUCE, called WEB REDUCE, which is great work by F. J. WRIGHT and others.

The code lines were tested on an iMac 24", Apple M3, Tahoe 26.6.2 using Safari and WEB REDUCE (6745), 2-Apr-2024.

Acknowledgement. While using WEB REDUCE to write these notes I benefited greatly from the experience and expertise of Dr. Francis J. WRIGHT. He also corrected or optimized some procedures, so I am very grateful for his friendly help and support with tips and hints.⁴

Any feedback from the user is very welcome.

Wolfgang Lindner
Leichlingen, Germany
July 2026

²Use our functions with care, because I have dropped e.g. checks against input errors by the user.

³Of course, you can also use and run the code from the scripts in the REDUCE app.

⁴This booklet is based on my earlier books [10], where I use the small compact CAS EIGENMATH [23] designed for the use by physicists and the CAS Maxima version ‘Statistics - A Maxima Companion’. The interested reader may therefore look at the realizations of statistical concepts in different CAS.

0.2 Some Hints for the Use of WEB REDUCE



Read about the Limitations of WEB REDUCE.

Nevertheless, we draw attention in particular to the following facts.

1. Note that, because of the asynchronous nature of JavaScript, a `load_package` statement must always be alone in its own input.
2. Before using the `mat` matrix input operator, you need to execute `load_package matrix;` - otherwise you may get an *end-of-file* warning.
3. Any kind of comment at the end of an input command causes an *end-of-file* warning in the [Input/Output Display]. Normally, this should not irritate you - therefore we use end-of-line commands via `...% ...` if we want to comment a short command.
4. You may paste single commands or the whole `*.txt` file with code lines from our scripts into the WEB REDUCE [Input Editor] window, where you do not see any line numbers as in normal desktop REDUCE. The InputEditor window then allows free manipulations of it's content before you press [Send input].
5. Advice: submit each of the command lines (separated by blank lines) in our scripts one after another to the WEB REDUCE [InputEditor] window and then press [enter]. Observe and study the result for a good learning process. Otherwise you get all results collected as a bulk at the end of the Output window.

1 Descriptive Statistics

We collect the well known measures of tendency, dispersion and deviation. We translate the mathematical definitions into REDUCE code and demonstrate exemplary calls and examples. We define and code the standard functions of descriptive statistics, i.e. `mean`, `var`, `std`, `sem`, `mad`, `rms`, `median`, `mode`, `madM`, `quantile`, `moment`, `skew`, `kurtosis`. Ergo the user should easily follow the calculations in standard texts like [18].

All these calculations are simply done with REDUCE.

1.1 mean - the arithmetic Mean

The arithmetic mean $\bar{X}$ of a list X of numbers is the sum of all of the numbers in the list divided by their count.

Definition For a vector (list) $(x_1, x_2, \dots, x_n) \in \mathbb{R}^n$ of numbers the *arithmetic mean* $\bar{X}$ or `mean(X)` (or simply *mean* or *average*) is defined by


$$\bar{X} := \frac{1}{n} \sum_{i=1}^n x_i = \frac{x_1 + x_2 + \dots + x_n}{n} =: \text{mean}(X)$$

Example [from WIKIPEDIA] The arithmetic mean of the five values: 4, 36, 45, 50, 75 is:

$$\bar{X} = \frac{4 + 36 + 45 + 50 + 75}{5} = \frac{210}{5} = 42$$

If the five values are interpreted as hourly earnings in of five employees, then the arithmetic mean corresponds to the hourly earning, that everyone would receive if the total earnings were distributed equally among all employees.

Mental image



Use WEB REDUCE (a small case study)

To calculate the *mean* by the sum of all of the numbers in the list divided by their count, we first have to calculate their sum. We present three different possibilities to calculate the *sum*, also to introduce some concepts of REDUCE.

1. First we use the data structure *array* of REDUCE.⁵

▷ <https://reduce-algebra.sourceforge.io/web-reduce/?gnuplot>

```
array X(5);                % declare a field ('vector') of 5 elements
X(1):=4; X(2):=36$ X(3):=45$ X(4):=50$ X(5):=75$           % fill array
S := (for i:=1:5 sum X(i)); % sum the members by means of a for loop
write S;                  % display the result
```

⁵In the following code lines we could also write: `S := for i:=1:5 sum X(i);` i.e. leave out the outer `(..)` and cancel the line `write S;` – because the result will be displayed automatically.

2. Second we use the data structure *list* of REDUCE.

```
X := {4 ,36, 45, 50, 75};           % (1)
S := for each el in X sum el;      % (2)
```

Comment. In line (1) we declare the list X using the list delimiters $\{..\}$.

In (2) the sum S is built by **sum**'ing up all elements (el) of the list X . Using this Lisp-oriented language this code line reads like commands in a genuine English sentence – maybe like WOLFRAM|ALPHA. $\square$

Here is how we execute script 1 with WEB REDUCE: write the code lines in the window named *Input Editor* and press *Send input*:




Figure 1: Left: Code of script 1. in the [Input Editor] area of WEB REDUCE  Right: .. code is executed after click on [Send Input] button. The inputs and outputs are shown in the [Input/Output Display] window: with input lines in red and outputs in blue.

3. Third we abstract the *summation* and the calculation of the *mean* using a *function procedure* of REDUCE. How to write the **sum1** function using code.1 is shown in fig.1.left. We discuss here the second function **mean** using a LISP construct:⁶

```
procedure xapply(f,lst); lisp(f.cdr(lst))$           % (1)
procedure mean(lst); xapply(+, lst)/length(lst)$     % (2)
mean({4 ,36, 45, 50, 75});                           % (3)
```

⁶cf. [30, p.385, footnote]

Comment. (1) defines a lisp-helper function `xapply(f,L)` of two arguments: an operator f and a list L . This helper is used in (2) defining the function `mean` by applying `xapply` on the operator `+` and the list `lst`, giving the sum of all elements of the list. This interim result is then divided by the `length` (i.e. the number of elements in the list) to give the mean. (3) shows an example how to invoke function `mean`. – This construction of `mean` using `xapply()` has the advantage, that we may use it later with other operators like `-` or $\sqrt{}$ etc. $\square$



Here is a workflow in case longer code snippets occur and you don't want to write it directly on the Input Editor of WEB REDUCE.

- ▷ View code snippet e.g. via a link like View Code,
- ▷ then *mark* all or part of the code,
- ▷ then *paste* it into the Input Editor of WEB REDUCE and
- ▷ *press* **Send Input**.

In a nutshell:

View Code, run REDUCE

You should have the following configuration:

COMMENT

STATISTICS WITH REDUCE ch. 1.1
Dr. W.G. Lindner
Leichlingen, Germany 2026;

```
X := {4 ,36, 45, 50, 75};

procedure sum1(lst); for i := 1:length(X) sum part(X, i)$
sum1(X);

procedure xapply(f,lst); lisp(f.cdr(lst))$
procedure mean(lst); xapply(+, lst)/length(lst)$
mean({4 ,36, 45, 50, 75});

on rounded;
mean(X);
off rounded;
```

Figure 2: Left: opened source of code after invoking the link ▷ Code
Right: pasted code lines executed after click on button [Send Input]

Remark. If you copy the commands part-by-part you may simulate a DEMO process. If you copy the commands completely you get a BATCH process.

General information

General mathematical information about the concept is here ▷ WIKIPEDIA : Mean

Syntax and semantic of the function is here ▷ MATLAB : `mean`

1.1.1 Exercises

Exercise 1. (Spiegel, p.74, P 3.19 - grouped data with guessed mean A)

Calculate the arithmetic mean of the numbers 5, 8, 11, 9, 12, 6, 14, 10 choosing as the ‘guessed mean’ the number $A = 9$.



View *Solution*, run REDUCE.

Exercise 2. (mean w.r.t. a frequency)

100 students took part in a test. The marks $X = (61, 64, 67, 70, 73)$ were distributed with the frequency $f = (5, 18, 42, 27, 8)$. In other words, we had the frequency tableau $\begin{pmatrix} 61 & 64 & 67 & 70 & 73 \\ 5 & 18 & 42 & 27 & 8 \end{pmatrix}$, i.e. 5 students get 61 marks, 18 students get 64 marks etc.

Calculate the mean mark of the 100 students.

Hint: use formula $\text{meanF}(X, f) = \text{dot}(f, X) / \text{sum}(f)$ for a mean of X w.r.t. a frequency f , where $\text{dot}(A, B)$ is the well-known scalar product ‘ $\bullet$ ’ of Linear Algebra.⁷



View *Solution*, run REDUCE.

Exercise 3. Let $X = (1, 3, 6, 6, 6, 6, 7, 7, 12, 12, 17) = \begin{pmatrix} 1 & 3 & 6 & 7 & 12 & 17 \\ 1 & 1 & 4 & 2 & 2 & 1 \end{pmatrix}$.

Calculate the mean of X .

⁷Here is a hint by F.J. WRIGHT: you might like to take a look at the REDUCE LISTVECOPS package, see LISTVECOPS, which provides some useful list operations including the dot (scalar) product and more elegant syntax than using the part operator.

1.2 var - the variance

The *variance* is the mean squared deviation from the mean. Variance is a measure of how far the observed values x_i in a dataset X fall from the arithmetic mean $\mu := \bar{X}$ and is therefore a measure of spread.

Definition

- For a vector $X := (x_1, x_2, \dots, x_n)$, the (population) *variance* is defined as

$$\text{var}(X) := \frac{1}{n} \cdot \sum_{i=1}^n (x_i - \mu)^2 = \frac{(x_1 - \mu)^2 + (x_2 - \mu)^2 + \dots + (x_n - \mu)^2}{n}$$

where μ is the mean of X .

- The (sample) *variance* is defined as $\text{var1}(X) := \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)^2$.

Examples The (sample!) variance of the five values: 4, 36, 45, 50, 75 is:

$$\text{var}(X) = \frac{(4 - 42)^2 + (36 - 42)^2 + (45 - 42)^2 + (50 - 42)^2 + (75 - 42)^2}{5} = 528.4$$

Use WEB REDUCE

```
X := {4, 36, 45, 50, 75} $
var(X);                % = 528.4
var1(X);               % = 660.5
```



View *Code*, run REDUCE.

Mental image




Figure illustrating the concept of variance for two different sets of 19 numbers (0, 5, ..., 90) and (0, 37, 38, ..., 53, 90). μ denotes the mean and σ denotes the square root of the variance as a measure of the spread.

General information

General mathematical information about the concept is here ▷ [WIKIPEDIA : Variance](#)
 Syntax and semantic of the function is here ▷ [MATLAB : var](#)

1.2.1 Exercises

Exercise 4. Check the example using the built-in function *var* of the software $\mathcal{R}$.

Hint: R: `var`

Exercise 5. (example from wikipedia, *ibid.*)

For a set of numbers 10, 15, 30, 45, 57, 52, 63, 72, 81, 93, 102, 105, if this set is the *whole data population* for some measurement, then variance is the population variance 932.743 as the sum of the squared deviations about the mean of this set, divided by 12 as the number of the set members.

If the set is a *sample* from the whole population, then the unbiased sample variance can be calculated as 1017.538 that is the sum of the squared deviations about the mean of the sample, divided by 11 instead of 12.

Check this using our user-defined functions `var` and `var1`.



View *Solution*, run REDUCE.

Hint: Use the following code of CAS MAXIMA as pattern .

```
X : [10, 15, 30, 45, 57, 52, 63, 72, 81, 93, 102, 105]$

Sum(M) := apply("+", M)$
Mean(M) := Sum(M)/length(M)$
Var(M) := Sum((M-Mean(M))^2) / length(M)$
Var(X), numer;
```

1.3 sd - the Standard deviation

The *standard deviation* is the square root of the mean squared deviation from the mean. A large standard deviation indicates that the data points x_i in a dataset X can spread far from the mean $\mu := \bar{X}$ and is therefore a measure of spread.

Definition

- For a vector $X := (x_1, x_2, \dots, x_n)$, the (population) *standard deviation* is defined as

$$\textcolor{red}{sd}(X) := \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2}$$

where μ is the mean of X .

- The (sample) *standard deviation* is defined as $sd1(X) := \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)^2}$.

Mental image




Table of annual salary of employees, $\mu = \text{mean}$, $s = \text{sd}(X)$: $\triangleright$ var

▢: 68 % of the data fall inside the red region $\mu \pm s$.

Figure 4: ▢: 95 % of the data fall inside the red+green region $\mu \pm 2 * s$.

▢: 99.7 % of the data fall inside the red+green+blue region $\mu \pm 3 * s$.

$\frown$: the empirical density function

Example

The (population) standard deviation of the five values: 4, 36, 45, 50, 75 is:

$$\text{sd}(X) = \sqrt{\text{var}(X)} = \sqrt{528.4} = 22.98$$

Use REDUCE

```

on rounded$
print_precision 5$
X := {4, 36, 45, 50, 75}$
sd(X);                % = 22.98
sd1(X);               % = 25.70

```

View *code*, run REDUCE.**General information**

- General mathematical information about the concept is ▷ WIKIPEDIA : std.deviation
- Syntax and semantic of the function is in ▷ MATLAB : **std**

1.3.1 Exercises

Exercise 6. Check the example using the built-in functions *sd* of the software $\mathcal{R}$.

Hint: R: sd or R: r-sd

Exercise 7. (example from wikipedia, ibid.)

Suppose that the entire population of interest is eight students in a particular class. Their marks are the following eight values: 2,4,4,4,5,5,7,9.

- Our *sd* formula of REDUCE is valid only if the eight values with which we began form the complete population. If the values instead were a random sample drawn from some large parent population (for example, there were 8 students randomly and independently chosen from a student population of 2 million), then one divides by 7 (which is $n - 1$) instead of 8 (which is n) in the denominator of the *sd* formula, and we have to use function *sd1* of REDUCE.

Check this using our user-defined function **sd** resp. **sd1**.

View *code*, run REDUCE.

Exercise 8. (example from cf. Spiegel p. 106, P 4.10)

Find the std. deviation of 12, 6, 7, 3, 15, 10, 18, 5.

View *code*, run REDUCE.

Exercise 9. (students marks)

Find the std. deviation of the height of the students marks $X = (61, 64, 67, 70, 73)$ with frequency $f = (5, 18, 42, 27, 8)$, e.g. mark 61 was reached 5 times.

Hint: We have to adapt the definition of *std* w.r.t. frequency f .

Define e.g. 'mean w. frequency' by $meanF(X, f) = \frac{dot(f, X)}{\sum f}$.

View *code*, run REDUCE.

Exercise 10. (Z-score alias `standard.score`)

The Z-score function is used to transform a list (‘vector’) of numbers into a new list with $mean = 0$ and $std = 1$, cf. [1] or [2]

Task: Define function ‘`zscore`’ and transform the list 6, 2, 8, 7, 5 into Z-scored numbers. Check mean and `std.dev` of the transformed numbers.



View *code*, run REDUCE.

```
X := {6,2,8,7,5}$  
zscore(X);           % 0.1942, ...
```

Exercise 11. (example from `geeksforgeeks`)

Using REDUCE do Example 1 and 2 in [3].



View *code*, run REDUCE.

1.4 sem - the Standard Error of Mean

For a given sample X , the standard error of the mean $sem(X)$ equals the standard deviation divided by the square root of the sample size $dim(X)$.

In other words, the *standard error of the mean* is a measure of the dispersion of sample means around the population mean. The SEM is a measure of how much a sample mean is likely to differ from the true population mean, see example 2. This concept is heavily used in hypothesis testing and calculating confidence intervals.

Definition

For a vector $X = (x_1, x_2, \dots, x_n)$, the *standard error of the mean* is defined as

$$sem(X) := \frac{sd(X)}{\sqrt{n}}$$

where $sd(X)$ is the standard deviation of the population.

Mental image




Figure 5: Visualization of sample vs. population and their resp. means $\bar{X}$ vs. μ . μ_0 is the theoretical mean, i.e. the target value.

Examples

1. The standard error of the mean of the five values $X = (4, 36, 45, 50, 75)$ is $sem(X) = 10.28$

2. In the table below, a measurement of $N = 10$ resistance values of a production of resistors with a $\mu_0 = 100 \Omega$ ‘guaranteed’ value is shown. A reasonable assumption is that the production is normally distributed. Let us assume that the spread of the production is well known with sigma $\sigma = 0.5 \Omega$. To check whether the target value 100Ω is still being met by the production, one has to determine first of all the test statistic $T = \frac{(\bar{X} - \mu_0)}{sem(X)}$.

Calculate T for the following sample of resistors.⁸

<i>Sample of a production of 100 resistors.</i>										
$n :$	1	2	3	4	5	6	7	8	9	10
$\Omega :$	100.1	101.2	99.5	99.0	100.7	100.0	101.2	99.2	99.0	98.7

⁸This example concerning resistors quality is in BEUCHER [1, p. 219]. See chapter §4.1 of this script.

Use REDUCE for example.1:

```
X := {4, 36, 45, 50, 75} $
std(X)/ sqrt(length(X));           % =10.28
sem(X);
```



View *code*, run REDUCE.

General information

- Mathematical **Information** about the concept is $\triangleright$ WIKIPEDIA : std.error of mean
- Syntax and semantic of the function in REDUCE : – no built-in function provided –

1.4.1 Exercises

Exercise 12. (user-defined function *sem* for REDUCE)

Use our user-defined function **sem** following the mathematical definition of *sem* and calculate the std.err.mean and the test statistic *T* for the values of example.2.

- MATLAB: Here by default, the standard deviation is normalized by $n - 1$, where n is the number of observations.



View *code*, run REDUCE.

Exercise 13. (check)

Verify, that $sem([1, 2, 3, 4]) = 0.6454$.

Exercise 14. (example from soton.ac.uk)

Read: [1]. Program the right most explicit term for SE in [1] in REDUCE.

Calculate the sem (=SE) of examples 1 and 2 in [1] using REDUCE.

Exercise 15. (example from investopedia)

Do the example in [2].

1.5 mad - the average absolute deviation

The average absolute deviation *mad* of a data set is the average of the absolute deviations from a central point. It is a measure of statistical dispersion or variability. The central point can be e.g. *mean*, *median*, *mode*.

Definition

For a list $X = (x_1, x_2, \dots, x_n)$, the *mean/median/.. absolute deviation* **mad** is defined as

$$\text{mad}(X) := \frac{1}{n} \sum_{i=1}^n |x_i - m(X)|$$

where $n = \text{length}(X)$ and $m \in \{\text{mean}, \text{median}, \text{mode}, \dots\}$.

Mental image




Figure 6: Figure illustrating the concept of *mad* for the data set of the 19 numbers $X = (0, 5, 10, \dots, 90)$.

Examples

1. The *mad* w.r.t. the *mean* of the five values 4, 36, 45, 50, 75 is 17.6.
2. The *mad* w.r.t. the *median* of the five values 4, 36, 45, 50, 75 is 9.

Use REDUCE

```
L := {4,36,45,50,75};
mad(L);                      % = 17.6
```



View *code*, run REDUCE.

General information

General mathematical information about the concept is $\triangleright$ WIKIPEDIA : average abs. dev.
 Syntax and semantic of the function is here $\triangleright$ MATLAB : **mad**

1.5.1 Exercises

Exercise 16. (user-defined function *mad* for REDUCE.)

Write a function `mad` using the mathematical definition of *mad* and use it to calculate the *mad* for example 1. and 2. For 2. use *median()* instead of *mean*.



View *code*, run REDUCE.

Exercise 17. (check, WIKIPEDIA *ibid.*)

Verify, that for example, for the data set 2, 2, 3, 4, 14 the *mad* is 3.6.



View *code*, run REDUCE.

Exercise 18. (mean squared deviation)

Calculate for the contrast to the mean absolute deviation (*mad*) the ‘mean squared deviation’ `dsquare` defined by $dsquare(X) = \frac{\sum \bar{X}^2}{dim(X)}$ using REDUCE; you may use this ‘pseudo code’ (in fact runnable EIGENMATH^{online} code) as pattern:

```
X=(4,36,45,50,75)
mX    = mean(X)
Xbar  = X-mX
Xbar2 = Xbar^2
dsquare = sum(Xbar^2)/dim(X)
```



View *code*, run REDUCE.

Exercise 19. (median abs.dev. around the median)

For the sample 2, 2, 3, 4, 14 the number 3 is the median, so the absolute deviations from the median are 1, 1, 0, 1, 11 (reordered as 0, 1, 1, 1, 11) with a median of 1, in this case unaffected by the value of the outlier 14, so the median absolute deviation is 1.

Define a function *madMed()* to calculate the ‘median abs.dev. around the median’ and verify, that `madMed( 1,2,3,4,5,6,7,8,9 ) = 2`.

Hint: use the definition `madMed(X) = median(abs(X - median(X)))`.



View *code*, run REDUCE.

Exercise 20. (test vehicles of one car type, cf. Beucher p.176, P 64)

For 40 test vehicles of one car type, the fuel consumption *X* was determined in liters per 100 km of driving in the city center. This sample yielded the following results (table TX with 5 rows):

```
TX = zero(5, 8)
TX[1] = (10.1,10.6,10.9,10.0,10.4,10.5, 9.7,10.5)
TX[2] = (10.4,10.1,10.8, 9.2,10.2,10.3,10.5, 9.2)
TX[3] = (10.2,10.5, 9.4,10.2, 9.6,10.2, 9.7,10.2)
TX[4] = (10.8, 9.9,10.5,10.6, 9.8,10.7,11.2,10.8)
TX[5] = ( 9.9,10.0,10.5,10.4,11.4,10.4,10.1,10.4)
```

Divide the sample into the following five approx. equidistance characteristic classes: < 9.5 , $9.5 - 10.0$, $10.0 - 10.5$, $10.5 - 11.0$, > 11.0 .

For the following task use functions of REDUCE.

- a. Determine the list of absolute and relative frequencies.
- b. Determine the cumulative distributions corresponding to the class classification and represent this in a histogram.
- c. Collect the known ‘measures of dispersion’ i.e. mean, std.dev., variance, m.a.deviation etc. in a ‘score’ table.



View *code*, run REDUCE.

1.6 rms - the root mean square

The *root mean square* (abbrev. *rms*) of a list of values is the square root of the list's mean square.

Definition

For a vector $X = (x_1, x_2, \dots, x_n)$, the *rms* is defined as

$$\text{rms}(X) := \sqrt{\frac{1}{n} (x_1^2 + x_2^2 + \dots + x_n^2)}$$

where $n = \text{length}(X)$.

Mental image




Figure 7: Figure illustrating the concept of *rms* for the data set of the 19 numbers $X = (0, 5, 10, \dots, 90)$.

Examples

1. The *rms* of the five values 4, 36, 45, 50, 75 is 47.88.
2. The *rms* of $X = (0, 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75, 80, 85, 90)$ is 52.68.

Use REDUCE to check 1. (where there is also no build-in function for *rms*()):

```
X := {4,36,45,50,75}$
rms(X); % = 47.87901
```



View code, run REDUCE.

General information

Mathematical information about the concept is ▷ WIKIPEDIA : Root mean square

Syntax and semantic of the function is here ▷ MATLAB : rms

1.6.1 Exercises

Exercise 21. (user-defined function *rms* for REDUCE)

Write a function `rms` using the mathematical definition of *rms* in REDUCE and calculate the *rms* for example 2.



View *code*, run REDUCE.

Exercise 22. (check, WIKIPEDIA *ibid.*)

Verify, that for example, for the data set $X = (1, 2, 3, 4, 5, 6, 7, 8, 9, 10)$ the *rms*(X) is 6.204837.



View *code*, run REDUCE.

Exercise 23. (*rms* in MATLAB)

Check some of the examples in ▷ MATLAB : `rms`



View *code*, run REDUCE.

1.7 median - the Median

The *median* is the central value of an *ordered* data set and divides it into a part with small and with large data points. Therefore, the median is a suitable measure for determining the distribution of a data set. The median minimizes the total distance to all other data elements. It is the solution to the optimization problem $\min_{a \in \mathbb{R}} \sum_{\nu=1}^n |x_{\nu} - a|$.

If a data element changes, this only causes a change in the median, if the older value moves from one half of the ordered data set to the other half.

Definition

For a *sorted* vector $X = (x_1, x_2, \dots, x_n)$, the *median* $\tilde{X} = \text{median}(X)$ is defined as

$$\tilde{X} := \begin{cases} x_{(n+1)/2} & \text{if } n = \dim(X) \text{ is odd} \\ \frac{1}{2}(x_{n/2} + x_{n/2+1}) & \text{if } n = \dim(X) \text{ is even} \end{cases}$$

where $n = \text{length}(X)$.

Mental image




Figure 8: Figure illustrating the concept of *median* for the data set of the 19 numbers $X = (0, 5, 10, \dots, 90)$.

Examples

1. The median of the five values 4, 36, **45**, 50, 75 is 45; $n = \dim(X) = 5$.
2. The median of $X = (36, 4, 75, 45, 50) \xrightarrow{\text{sort}} (4, 36, \mathbf{45}, 50, 75) \xrightarrow{\sim} 45 = \text{median}(\text{sort}(X))$.

Use REDUCE to check 1. :

```
load_package assist;
X := {4, 36, 45, 50, 75}$
median(X);           % = 45
```



View *code*, run REDUCE.

General information

General mathematical information about the concept is $\triangleright$ WIKIPEDIA : Median
 Syntax and semantic of the function is here $\triangleright$ MATLAB : median

1.7.1 Exercises

Exercise 24. (user-defined function `Median`)

Write function `Median` using the mathematical definition of *median* and check your function on example 1.

Hint: you must take into account that the list X for the use of *Median* has to be sorted! You may use REDUCE's `sort(X,<)` function.



View *code*, run REDUCE.

Exercise 25. (check median for a data set)

Verify, that for example, for the data set

```
X = (1,2,3,4, 5, 6,7,8,9)
    -- 4 elem. ^ 4 elements right
```

your `Median(X)` is 5.



View *code*, run REDUCE.

Exercise 26. (*median* in MatLab and $\mathcal{R}$)

Check some of the examples in $\triangleright$ MATLAB : `median` or in $\triangleright$ R : `median`



View *code*, run REDUCE.

1.8 mode - the Mode

The *mode* is the value $\tilde{X}$ in a data set X that appears most often, i.e. the maximal value in the frequency distribution of X .

Definition

For a vector $X := (x_1, x_2, \dots, x_n)$, the *mode* is defined as

$$\text{mode}(X) := \max \text{freq}(X)$$

where *freq* is the frequency table of X (not defined here).

Mental image



Figure 9: Figure illustrating the concept of *mode* for a sets of 9 numbers (37, 38, 38, 40, 42, 42, 42, 45, 46). The most sold length of the 9 shoes was 42.

Examples

1. The *mode* of the five values 4, 36, 45, 50, 75 is 4.
2. $\text{freq}(1, 3, 6, 6, 6, 6, 7, 7, 12, 12, 17) = \begin{pmatrix} 1 & 3 & 6 & 7 & 12 & 17 \\ 1 & 1 & 4 & 2 & 2 & 1 \end{pmatrix} \Rightarrow \text{mode}((1, 3, 6, 6, 6, 6, 7, 7, \dots, 17)) = 6$
3. $\text{mode}((1, 1, 2, 4, 4)) = 1$

Use OCTAVE e.g. with SageMathCell or MATLAB^{online} to check 1. :

```
octave:1> X=[1,2,1,4,4]
octave:2> mode(X)
ans = 1
```

General information

General mathematical information about the concept is ▷ WIKIPEDIA : Mode
 Syntax and semantic of the function is here ▷ MATLAB : mode

1.8.1 Exercises

Exercise 27. (user-defined function `mode`)

There is no build-in function `mode` for REDUCE.

Write function `mode()` using the above incomplete mathematical ‘definition’ of `mode` and calculate the `mode`’s for examples 1., 2. and 3. .

Hint: you may get an idea looking at example 2.



View `code`, run REDUCE.

Exercise 28. (another user-defined function `mode`)

Write a user-defined function `mode1()` using the following ‘pseudo’code and check your code for examples 1., 2. and 3. .

```

----- DEFINITION ----- mode
mode(v, i,j,k,m,n) =      -- by G. Weigt 8/2025
do( n = 0,
  for(i,1,dim(v),
    k = 0,
    for(j,1,dim(v), test(v[i] == v[j], k = k + 1)),
    test(or(k > n, and(k == n, v[i] < m)),
      do(n = k, m = v[i]))), m )

-- The above pseudocode (aka runnable EIGENMATH code)
-- should be easily ported to REDUCE.
-----

```

- What do think about an insightful programming of the function `mode` using a list oriented programming a la REDUCE vs. a C-style program like the above?

Exercise 29. (`mode` in MatLab)

Check some of the examples in ▷ MATLAB : `mode`



View `code`, run REDUCE.

1.9 quantile - the Quantile

A *quantile* is a score at or below which a given percentage of all the scores exists, i.e., a score in the k -th percentile would be above approximately $k\%$ of all scores in its set.

[We quote ▷QUANTILE] One definition of percentile or *quantile* Q_p is that the p -th percentile of a list X of n ordered values (sorted from least to greatest) is the smallest value in the list such that no more than p percent of the data is strictly less than that value and at least p percent of the data is less than or equal to that value. This is obtained by first calculating the ordinal rank and then taking the value from the ordered list that corresponds to that rank. The ordinal rank N is calculated using the formula $N = \lceil p * 0.01 * n \rceil$.⁹

Definition (The *nearest-rank method* for a quantile)

For a vector $X := (x_1, x_2, \dots, x_n)$, the *quantile* Q_p is defined as

$$Q_p(X, p) := \begin{cases} \max(X) & \text{if } p = 100 \\ X_{\lceil p * 0.01 * n \rceil} & \text{if } p \in [0, 100[\end{cases}$$

where $n = \dim(X)$ is the length of X .

Mental image




Figure illustrating the concept of *quantile* for the sets of numbers
 Figure 10: (12, 3, 4, 56, 78, 18, 46, 78, 100): $\begin{pmatrix} 0\% & 25\% & 50\% & 75\% & 100\% \\ 3 & 12 & 46 & 78 & 100 \end{pmatrix}$, visualized in a so-called *boxplot*. $Q_{50} = 46$ is the *median*(X).

⁹https://en.wikipedia.org/wiki/Floor_and_ceiling_functions, the ceiling function maps x to the least integer greater than or equal to x , denoted $\lceil x \rceil$ or $\text{ceil}(x)$.

Example

The quantile values for the ordered list $X = (12, 3, 4, 56, 78, 18, 46, 78, 100)$ are $Q_0 = 3, Q_{25} = 12, Q_{50} = 46, Q_{75} = 78, Q_{100} = 100$.

Use REDUCE to check:

```
X:={12, 3, 4, 56, 78, 18, 46, 78, 100};

on rounded $
print_precision 5$

quantile(X, 25);    % = 12
quantile(X, 50);    % = 46
quantile(X, 75);    % = 78
quantile(X, 100);   % = 100
```



View *code*, run REDUCE.

- OCTAVE uses `prctile(X,[0,25,50,75,100])` and gives `ans = 3 10 46 78 100`.

Remark.

- Using the nearest-rank method on lists with fewer than 100 distinct values can result in the same value being used for more than one percentile.
- A percentile calculated using the nearest-rank method will always be a member of the original ordered list. This is not always the case using $\mathcal{R}$ or MATLAB/OCTAVE, which uses advanced sophisticated algorithms.
- The 100th percentile is defined to be the largest value in the ordered list, the 0th percentile to be the smallest value.

General information

- General mathematical information about the concept is $\triangleright$ WIKIPEDIA : quantile
- Syntax and semantic of the function is here $\triangleright$ MATLAB : `prctile`

1.9.1 Exercises

In CAS REDUCE, quantiles are calculated using our user-built functions from the following exercise 31. To find the p -quantile, use the naming convention `quantile(X,p)`, where p is a number in $[0, \dots, 100]$ and X is the sample list.

Exercise 30. (boxplot of Fig.10)

Reproduce the boxplot fig.10 for the data set (12, 3, 4, 56, 78, 18, 46, 78, 100), cf. `> boxplot`.

`> Solution with CAS MAXIMA.`

Exercise 31. (a user-defined function `Quantile`)

Write a user-defined function `quantile()` using the mathematical definition for Q_p using the *nearest-rank method* for a quantile and check your function on example §1.9.

Compare the result w.r.t. the CAS REDUCE, MATLAB and $\mathcal{R}$.



View code, run REDUCE.

Exercise 32. (five-number summary)

The ‘five-number summary’, consisting of the minimum, 1st Quartile, Median, 3rd Quartile, and maximum, can be computed using the *quantile* function. These five statistics break down the position and spread of a dataset.

Program the ‘five-number summary’ as a REDUCE function.



View code, run REDUCE.

Exercise 33. (example from MATLAB)

Calculate the quantile of data set X for percentage $p = 42\%$, where $X : 0.5377, 1.8339, -2.2588, 0.8622, 0.3188, -1.3077, -0.4336$.

Result: `quantile(X,42)` should be -0.1026 or -0.0424



View code, run REDUCE.

Exercise 34. (example from $\mathcal{R}$)

Reproduce the following calculation in REDUCE:

```
R> df <- c(6, 3, 2, 10, 1)
R> quantile(df)
[1] 0% 25% 50% 75% 100%
     1  2  3  6 10
R> quantile(df, .4)
[2] 2.6 -- which is NOT in the data !
```



View code, run REDUCE.

Exercise 35. (example with ‘pseudo’code)

Do the following calculation in REDUCE:

```
A = (6, 3, 2, 10, 1)
percentile(A, 40)
Qp(A,40)
boxplot(A)
```



View code, run REDUCE.

1.10 moment - the r^{st} Moment

In statistics, moments are parameters that describe the shape of a probability distribution by measuring different aspects, such as its central tendency, spread, and asymmetry. The most common moments are the first moment (mean), which indicates the center; the second moment (variance), which measures spread; skewness, the third moment, which shows asymmetry; and kurtosis, the fourth moment, which describes the peakedness or flatness of the distribution, cf. $\triangleright$ *Google search: statistics moments. AI overview.*

The *moment* concept generalizes the concepts of variance, skewness and kurtosis. The r -th raw moment of a population can be estimated using the r -th raw sample moment, applied to a sample $x_1, \dots, x_n$ drawn from the population. The *(central) moment of order r about A* , $\text{moment}(X, A, r)$ computes a sample version of a population value. The first-order central moment is zero, the second-order central moment is the variance computed using a divisor of n rather than $n - 1$, where n is the length of the data vector X .


Definition

For a data vector $X := (x_1, x_2, \dots, x_n)$, the r^{th} moment about A is defined as

$$\text{moment}(X, A, r) := \frac{1}{n} \sum_{i=1}^n (x_i - A)^r$$

where n is the length of X .

Mental image



Imagine the data distributed under the blue graph. Then the moment
 Figure 11: may be interpreted as ‘center of mass’, balanced upon the ‘expected’
 value $\triangle$, the moment e.g. the mean; cf. GUNDERSEN

Examples

1. The 1st *moment* of the five values 4, 36, 45, 50, 75 is 45, i.e. its mean.
2. The 2nd moment of the five values 4, 36, 45, 50, 75 is 2292.4.
3. The 3rd moment about $A = 4$ of the five values 4, 36, 45, 50, 75 is 111387.

Use REDUCE to check:

```
X := {4,36,45,50,75}$  
moment(X,0,1);      % = 45  
moment(X,0,2);      % = 2292.4  
moment(X,4,3);      % = 111387.2  i.e.: list=X, center=4, order=3
```



View *code*, run REDUCE.

General information

- General mathematical information about the concept ▷ WIKIPEDIA : Moment
- Syntax and semantic of the function is here ▷ MATLAB : `moment`

1.10.1 Exercises

Exercise 36. (user-defined function `moment`)

Write a function `moment(X,A,r)` using the mathematical definition for *moment* and check your function on the three examples.



View *code*, run REDUCE.

Exercise 37. (moment about the mean)

Let $X = (0.5377, 1.8339, -2.2588, 0.8622, 0.3188, -1.3077)$ be a small sample.

Let $A = \text{mean}(X)$ be the mean of X .

Calculate the $\text{moment}(X, A, 3)$ and $\text{moment}(X, 0, 3)$.

Result: -1.1143 using $\mathcal{R}$ with `center=TRUE`; -1.1274 using $\mathcal{R}$ with `center=FALSE`.



View *code*, run REDUCE.

Exercise 38. (example from Spiegel, p.129, P 5.1)

Calculate the first, second, third and fourth moments of the data set 2, 3, 7, 8, 10.

Hint: Let $A = 0$ be the ‘noncentral’ value of X .



View *code*, run REDUCE.

Exercise 39. (moments for grouped data, cf. Spiegel, p.131 P 5.6)

Students marks in a class have the following frequency table

class mark X:	61	64	67	70	73
frequency f:	5	18	42	27	8

i.e. class mark 67 comes with frequency 42.

Task: find the first four moments of the class marks.

Hint: the solution needs *meanF* for grouped data and a variant of moment w.r.t. mean of grouped data; here is the pseudocode:

```
meanF(X,f)      = dot(f,X) / sum(f)
momentMF(X,f,r) = dot(f,(X-meanF(X,f))^r) / sum(f)
```

Write `meanF(X,f)` and `momentF(X,f,r)` in REDUCE and then do `momentF(X,f,1) ...`.



View *code*, run REDUCE.

Exercise 40. (moments for grouped data, cf. Spiegel, p.137 P5.27)

Find the first four moments for the distribution $X : f$, where the

data is $X = (12, 14, 16, 18, 20, 22)$ with frequency $f = (1, 4, 6, 10, 7, 2)$.

Result: $\text{momentF}(X,f,1) = 0$; $\text{momentF}(X,f,2) = 5.97$; $\text{momentF}(X,f,3) = -0.397$;
 $\text{momentF}(X,f,4) = 89.22$



View *code*, run REDUCE.

1.11 skew - the Skew.ness

Skewness is a *measure of the asymmetry* of the data around the sample mean. If skewness is negative, the data spreads out more to the left of the mean than to the right. If skewness is positive, the data spreads out more to the right.

The skewness of the normal distribution (or any perfectly symmetric distribution) is zero. The *relative position of arithmetic mean and the median* to each other is also characterized by the skewness of a data set. If the data set is *skewed right* (steep to the left), the arithmetic mean lies to the *right* of the median. If the data set is skewed left (steep to the right), the arithmetic mean lies to the left of the median. If the data set is symmetrical, then the arithmetic mean and median are approximately the same.

Definition

For a data vector $X := (x_1, x_2, \dots, x_n)$, the *skewness* $\mathbf{skew}(X)$ is defined as

$$\begin{aligned} \text{skew}(X) &:= \frac{1}{n} \sum_{i=1}^n \left(\frac{x_i - \bar{X}}{s} \right)^3 = \frac{m_3}{s^3} \\ \text{skewness}(X) &:= \frac{\sum_{i=1}^n (x_i - \bar{X})^3}{(n-1)s^3} \end{aligned}$$

where n is the number of observations, m_3 is the 3rd moment and $s = sd(X)$ the standard deviation and $\bar{X}$ is the mean of X . – The 1st formula is used by MATLAB, the 2nd by R.

Mental image




Figure illustrating the concept of *skewness* for a distribution of numbers.
 Figure 12: The frequency distribution has positive skewness of 0.808 (skew right).
 The *mean*= 4.46 is placed *right* of the *median*= 4.

Example

The *skewness* of the five values 4, 36, 45, 50, 75 is -0.306 resp. -0.2741 .

Use `REDUCE` to check:

```
X := {4, 36, 45, 50, 75}$
skew(X);                      % = -0.306 [MatLAB resp. octave]
```



View *code*, run `REDUCE`.

Warning: if you use $\mathcal{R}$ to check you get different results w.r.t. options:¹⁰

	<code>X := {4, 36, 45, 50, 75}</code>	
	<code>skew(X)</code>	-0.306 [MatLAB]
	<code>skewR(X)</code>	-0.274 [R]
	<code>skew1(X)</code>	-0.219 [R, type=3]
	<code>skew2(X)</code>	-0.456 [R, type=2]

General information

- General mathematical information about the concept is $\triangleright$ WIKIPEDIA : Skewness
- Syntax and semantic of the function is $\triangleright$ MATLAB : **Skewness**

¹⁰the functions `skew`, `skewR`, `skew1` and `skew2` are defined in exercise 41.

1.11.1 Exercises

Exercise 41. (user-defined function `skewness` in REDUCE)Write a self-defined function `skew(X)` using the mathematical definition for *skew*.**Remark.**

1. The software $\mathcal{R}$ uses different realizations for *skewness*, which we also define in this exercise. So you get very different results. Study the code lines with care.
2. Because we have our own self-defined function *skew*, we are able to adapt their code for other features, i.e. to implement versions of *skew* which have identical values w.r.t. other statistics software. Look at the warning after the example.

View *code*, run REDUCE.**Exercise 42.** (*skewness* example from MATLAB)Let $X = (0.5377, 1.8339, -2.2588, 0.8622, 0.3188, -1.3077, -0.4336, 0.3426, 3.5784, 2.7694)$.

a. Verify using exercise.41.

You should at minimum do $\text{skew}(X) := \text{skewness}(X)$:

```

skew(X)      -- 0.1060  R: typ=1, Matlab: default
skewR(X)     -- 0.1006  R: ok
skew1(X)     -- 0.0905  R: typ=3, Matlab: ./
skew2(X)     -- 0.1258  R: typ=2, Matlab: flag = 0

```

b. Verify the results of $\mathcal{R}$ using our user-defined functions of exercise 41:

```

R> library(e1071)
skewness(X)                                -- default, we use skew1
[1] 0.09058831
skewness(X, type=3)                        -- we use skew1
[1] 0.09058831
skewness(X, type=1)                        -- we use skew
[1] 0.1060983
skewness(X, type=2)                        -- we use skew2
[1] 0.1258171

```

View *code*, run REDUCE.

1.12 kurtosis - the Kurtosis

The kurtosis, the fourth moment, describes the *peakedness or flatness* of the distribution.

Cf. ▷ *Google search: statistics moments. AI overview.*

[MatLab:] Kurtosis is a measure of how outlier-prone a distribution is. *The kurtosis of the normal distribution is 3.* Distributions that are more outlier-prone than the normal distribution have kurtosis greater than 3; distributions that are less outlier-prone have kurtosis less than 3. Some definitions of kurtosis subtract 3 from the computed value (measure only the 'excess'), so that the normal distribution has kurtosis of 0.

◦ Our kurtosis function `kurtosis(X)` does not use this convention.

Definition

For a data vector $X := (x_1, x_2, \dots, x_n)$, the *kurtosis* `kurtosis(X)` is defined as

$$\begin{aligned} \text{kurtosis}(X) &:= \frac{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{X})^4}{\left[\frac{1}{n} \sum_{i=1}^n (x_i - \bar{X})^2 \right]^2} = \frac{m_4}{m_2^2} \\ \text{kurtosisM}(X) &:= \frac{\sum_{i=1}^n (x_i - \bar{X})^4}{(n-1)s^4} \end{aligned}$$

where n is the length of X and $\bar{X}$ is the mean of X , $s = sd(X)$ the standard deviation and m_r is the r^{st} moment of X . – The 1st formula is used by R, the 2nd by MATLAB.

Mental image




Figure illustrating the concept of *kurtosis* for a set of numbers.

Figure 13: The normal distribution is noted in ‘---’, the distribution of the number set in ‘---’. The kurtosis is 2.295 resp. 2.226. See example.2.

Examples

1. The *kurtosis* of the five values 4, 36, 45, 50, 75 is 1.9129 resp. 1.7390.
2. The *kurtosis* of the sample

$X=(4,4,4,5,5,5,5,5,5,6,6,6,6,6,6,6,6,6,7,7,7,8,8,8,10,10,10,12,12,12,13,13,13)$
 with the frequency distribution $\begin{pmatrix} 3 & 6 & 9 & 3 & 3 & 3 & 3 & 3 \\ 4 & 5 & 6 & 7 & 8 & 10 & 12 & 13 \end{pmatrix}$, e.g. value 5 occurs 6 times, shown in figure.13 above is 2.295 resp. 2.226.

Use REDUCE to check example.1:

```
X := {4,36,45,50,75};
kurtosis(X);           % = 2.347 resp. 1.9129
```



View *code*, run REDUCE.

Check example.2 with $\mathcal{R}$:

```
| X=(4,4,4,5,5,5,5,5,5,6,6,6,6,6,6,6,6,6,7,7,7,8,8,8,10,10,10,12,12,12,13,13,13)
|
| kurtosis(X) | 2.2958 [R; library(moments)]
```

General information

General mathematical information about the concept is ▷ WIKIPEDIA : Kurtosis
 Syntax and semantic of the function is ▷ MATLAB : `kurtosis`

1.12.1 Exercises

Exercise 46. (user-defined function `kurtosis` in REDUCE)Write a self-defined function `kurtosis(X)` using the mathematical definition for *kurtosis*.

- a. Check your function on the examples 1. and 2.  View code, run REDUCE.
- b. This implementation of `kurtosis` gives compatibility with MATLAB's *kurtosis* with 'no flag'. Verify this with your `kurtosis(X)` w.r.t. the following commands:

```
MATLAB>> Y = [4, 5, 5, 6, 6, 6, 7, 8, 10, 12, 13]
>> kurtosis(Y)
ans = 2.2959
```

- c. Check, that our `kurtosis(X)` is compatible with $\mathcal{R}$'s *kurtosis* from library 'moments':

```
## check compatibility with R and library(moments)
R> X = c(2,3,4,4,4,5,6,7,8,9,10)
R> library(moments)
R> kurtosis(X)
[1] 1.912949
```

- The software $\mathcal{R}$ use different realizations for *kurtosis*, which we discuss in the following exercises. Be warned: you get very different results w.r.t. *kurtosis* using different software!

Exercise 47. (MATLAB's *kurtosis0* with *flag* = 0, i.e. bias-corrected equation)Reconstruct MATLAB's `kurtosis(X, flag)`. That is: When X represents a sample from a population, the *kurtosis* of X is 'biased', meaning it tends to differ from the population kurtosis by a systematic amount based on the sample size. If this is the case:

- .. set *flag* = 0 to correct for this systematic bias.
- .. this bias-corrected equation requires that X contain at least four elements.

Here is a 'pseudo'code for your help:

```
(* MatLab's kurtosis with flag=0, i.e. bias-corrected equation *)
kurtosis0(X,n) = do( n=dim(X),
                    test( n<4 , "n must be >=4",
                        (n-1)/((n-2)*(n-3))*((n+1)*kurtosis(X) -3*(n-1))+3 ))
```

- a. Write REDUCE function `kurtosis0`, which implements MATLAB's `kurtosis(X, flag)`.
- b. Check your `kurtosis0` function w.r.t. this MATLAB call:

```
MATLAB>> Y = [4, 5, 5, 6, 6, 6, 7, 8, 10, 12, 13]
>> kurtosis(Y)
ans = 2.2959
>> kurtosis(Y,0)
ans = 2.6598
```



View code, run REDUCE.

Exercise 48. (kurtosis of data with different implementations)

Calculate the kurtosis of $data = (2, 3, 4, 4, 4, 5, 6, 7, 8, 9, 10)$ and compare the result with $\mathcal{R}$:‘e1071’ and $\mathcal{R}$:‘moment’.

```
# kurtosis with R and library(moments)
R> X = c(2,3,4,4,4,5,6,7,8,9,10)
R> library(moments)
R> kurtosis(X)
      [1] 1.912949

#kurtosis with R and library(e1071)
R> X = c(2,3,4,4,4,5,6,7,8,9,10)
R> library(e1071)
R> kurtosis(X)
      [1] -1.41905      --> .. + 3 = 1.581
```



View *code*, run REDUCE.

Exercise 49. (kurtosis of a skew data set)

Let $Y = (4, 5, 5, 6, 6, 6, 7, 8, 10, 12, 13)$ be a data set.

Calculate the kurtosis of Y with different functions.



View *code*, run REDUCE.

Exercise 50. (kurtosis for experimental data of figure.13)

Let Y be the data set of fig.13:

```
4,4,4,
5,5,5,5,5,5,
6,6,6,6,6,6,6,6,6,
7,7,7,
8,8,8,
10,10,10,
12,12,12,
13,13,13
```

a. Calculate the *kurtosis* of Y and argue, which value of the different versions (e.g. `kurtosis()` or `kurtosisM()` or `kurtosisR3()` etc.) of *kurtosis* is best suited to describe the shape of Y .

b. Which connection is seen between our self-built *kurtosis* function and the kurtosis of package *descriptive* of CAS MAXIMA?



View *code*, run REDUCE.

Exercise 51. (kurtosis of a frequency table, cf. Spiegel, p.135, P 5.13)

Let $X = (61, 64, 67, 70, 73)$ be the marks with a frequency distribution $f = (5, 18, 42, 27, 8)$.

Adapt the kurtosis to median and frequency of the group and calculate kurtosis of $X : f$.

Hint: you may use this pseudocode:

```
.....
momentF(X,f,r) = dot( f, (X-meanF(X,f))^r) / sum(f)
kurtosisF(X,f) = momentF(X,f,4) / momentF(X,f,2)^2
.....
```



View *code*, run REDUCE.

Exercise 52. (a left skew data set, cf. Spiegel, data Xleft: p.126, col.5-6)

The ‘obitage data’ set is

```
X = (102,55,70,95,73,79,60,73,89,85,
      72,92,76,93,76,97,10,70,85,25,83,
      58,10,92,82,87,104,75,80,66,93,
      90,84,73,98,79,35,71,90,71,63,
      58,82,72,93,44,65,77,81,77)
```

a. Calculate the kurtosis of X using $kurtosis(X)$ and $kurtosisR3(X)$ (i.e. giving the same value as if calling $\mathcal{R}$:e1071’ - Typ 3).

b. Give a tabulated summary of the results.



View *code*, run REDUCE.

Summary of the different implementations of *kurtosis*:

To calculate KURTOSIS	<i>use package</i>	<i>use REDUCE</i>
with self-build kurtosis :		<code>kurtosis</code>
analog MATLAB:	no flag	<code>kurtosis</code>
MATLAB:	flag = 0	<code>kurtosis0</code>
analog $\mathcal{R}$:	moments	<code>kurtosis</code>
$\mathcal{R}$:	e1071, typ 1	<code>kurtosisR1</code>
$\mathcal{R}$:	e1071, typ 2	<code>kurtosisR2</code>
$\mathcal{R}$:	e1071, typ 3	<code>kurtosisR3</code>

1.13 cov - the Covariance

Covariance is a statistical measure of the joint variability of two random variables, indicating how much they change together. A positive covariance signifies that variables tend to move in the same direction, a negative covariance indicates they move in opposite directions, and a zero covariance suggests no linear relationship.¹¹

There are different methods used to compute covariance: the *population cov* dividing the sum of the distance products $(x_i - \bar{X}) \cdot (y_i - \bar{Y})$ by n vs. the *sample cov* dividing by $n - 1$.

Definition

- For two data vectors $X := (x_1, x_2, \dots, x_n)$ and $Y := (y_1, y_2, \dots, y_n)$, the (population) *covariance* $\text{cov}(X, Y)$ between X and Y is defined as

$$\text{cov}(X, Y) := \frac{1}{n} \cdot \sum_{i=1}^n (x_i - \bar{X}) \cdot (y_i - \bar{Y})$$

where n is the length of X and Y and $\bar{X}$ resp. $\bar{Y}$ is the mean of X resp. Y .

- The *sample* covariance is defined as $\text{covR}(X, Y) := \frac{1}{n-1} \cdot \dots$, e.g. in R: typ ‘pearson’.

Mental image




Figure illustrating the concept of *covariance* for a set of numbers. Left: Figure 14: positive cov. Right: negative cov. $\bullet = (\bar{X}, \bar{Y})$. $\bullet = (x_i, y_i)$ a data point. $\times \times = (x_i - \bar{X}) \cdot (y_i - \bar{Y})$ a distance product w.r.t. the mean.

¹¹cf. $\triangleright$ Google search: covariance.

Examples

The covariance of the two data sets 1, 3, 5, 10 and 2, 4, 6, 20 is 23 resp. 30.66.

Use REDUCE to check the example:

```
% (* .. user-defined functions are here .. *)
X := {1,3,5,10};
Y := {2,4,6,20};

cov(X,Y);      % = 23
covR(X,Y);     % = 30.66
```



View *code*, run REDUCE.

- Check with $\mathcal{R}$:

```
> X<-c(1,3,5,10)
> Y<-c(2,4,6,20)
> cov(X, Y)           # default:  method = "pearson"
[1] 30.66667
```

General information

General mathematical information about the concept is $\triangleright$ WIKIPEDIA : Covariance
 Syntax and semantic of the function is $\triangleright$ MATLAB : cov

1.13.1 Exercises

Exercise 53. (user-defined function `cov` in REDUCE)

Write a user-defined function `cov(X)` using the mathematical definition for *cov*.
Check your function on the example.

Remark. 1. **Warning:** There are two different methods to compute the covariance, which leads to different results w.r.t population vs. sample.

2. The implementation of *covR* gives compatibility with $\mathcal{R}$'s *cov* with 'method = pearson'.



View *code*, run REDUCE.

Exercise 54. (Calculate the covariance of two data vectors)

Let $X = (16.50, 19.16, 7.64, 8.14, 8.14, 2.48, 9.23, 4.78, 6.93, 11.91, 11.60, 13.97)$
and $Y = (16.50, 19.16, 7.64, 8.14, 8.14, 2.48, 9.23, 4.78, 6.93, 11.91, 11.60, 13.97)$.

Determine `cov(X,Y)` and `covR(X,Y)`.



View *code*, run REDUCE.

Exercise 55. (Calculate the covariance of two data vectors)

Let $a = (2, 4, 6, 8, 10)$ and $b = (1, 11, 3, 33, 5)$.

Calculate `covR(a,b)` and `cov(a,b)`.

Result: `covR = 15, with method = "pearson"`



View *code*, run REDUCE.

Exercise 56. (Calculate the covariance and interpret the result)

Task: calculate

a. `cov( (2,3,4,3), (2,3,4,3) )`

b. `covR( (2,3,4,3), (2,3,4,3) )`

c. `cov( (1,2,3,4), (4,3,2,1) )`

d. Plot a diagram a la Fig.14 for each task and interpret the results.



View *code*, run REDUCE.

Exercise 57. (example from r-bloggers cf. *cov*)

Suppose we have two vectors, x and y , representing the number of hours studied and the corresponding test scores, respectively, for a group of students.

We want to measure the covariance between these two variables. – *Solution:*

1. Create example vectors $x = (5, 7, 3, 6, 8)$ and $y = (65, 80, 50, 70, 90)$

2. Calculate the covariance `cov(x,y)` resp. `covR(x,y) = 29`.

ok: `R> covariance(x,y) [1] 29`

3. Interpretation: – In this example, the resulting covariance value will help us understand the relationship between the hours studied and the corresponding test scores. What this particular example is saying is that for every unit increase in x there is a 29 unit increase in y .

Task: do this exercise in REDUCE.



View *code*, run REDUCE.

2 Discrete distributions

We implement some discrete resp. continuous statistical distributions as helper functions to do the statistical tests in chapter 4 and 5. For each distribution presented in chapter 2 and chapter 3 we give the definition resp. coding in REDUCE notation for the probability density function f (named ‘...*PDF*’), the cumulative distribution function F (named ‘...*CDF*’ and the quantile function F^{-1} (named ‘...*INV*’) i.e. the INVerse of the CDF.

2.1 Binomial distribution

In probability theory and statistics, the binomial distribution with parameters n and p is the discrete probability distribution of the number of successes in a sequence of n independent experiments, each asking a yes-no question, and each with its own Boolean-valued outcome: success or failure. The binomial distribution is used when there are exactly two mutually exclusive outcomes of a trial, cf. WIKIPEDIA: *Binomial_distribution*

Definition Notation: $X \sim \text{Binomial}(n, p)$

- The *probability density*¹² function for the binomial distribution is defined as:

$$\text{binoPDF}(\mathbf{k}, \mathbf{n}, \mathbf{p}) := f(k, n, p) := \Pr(X = k)^{13} = \binom{n}{k} p^k (1 - p)^{n-k}$$

- The *cumulative* binomial distribution function can be expressed as:

$$\text{binoCDF}(\mathbf{k}, \mathbf{n}, \mathbf{p}) := F(k; n, p) = \Pr(X \leq k) = \sum_{i=0}^{\lfloor k \rfloor} \binom{n}{i} p^i (1-p)^{n-i}$$

- The *quantile* function (inverse cumulative distribution function) for $\text{Bin}(n, p)$ is¹⁴

$$\text{binoINV}(\alpha, n, p) := \inf\{k \in \mathbb{R} : \alpha \leq F(k; n, p)\}$$

i.e. we must find the *smallest* k such that $\alpha \leq \sum_{i=0}^{\lfloor k \rfloor} \binom{n}{i} p^i (1-p)^{n-i}$ for given α .

Examples Verified with REDUCE in exercises N° 58, 60, 63.

1. *binomial PDF*: Suppose a biased coin comes up heads with probability 0.3 when tossed. What is the probability of seeing exactly 4 heads in 6 tosses?

binoPDF(4, 6, 0.3)	0.059535
--------------------	----------

¹²We adopt the naming convention of MATLAB.

¹³ $\Pr(X=k)$: represents the probability of getting exactly k successes

¹⁴there is no closed term for the inverse binomial distribution


2. *binocdf*: The cumulative probability of a binomial distribution with 10 trials and a probability of success 0.5 for 4 successes is 0.3769.

| `binocdf(4, 10, 0.5)` | 0.3769

3. *binoinv*: Given a number of trials $n=100$, the probability of success $p=0.3$, the cumulative area of the binomial distribution $\alpha=0.7$, find the first value x such that $\alpha = 0.7 \leq \text{binocdf}(x, 100, 0.3)$. Result: 32

| `binoinv(0.7, 100, 0.3)` | 32

Graphical representation



Left figure: $\bullet$ = plot for `binopdf(k, 40, 0.5)` for $k = 0, \dots, 40$.

Check `binopdf(20, 40, 0.5) = 0.125` on the graph.

Figure 15: Right figure: $\bullet$ = plot for `binocdf(k, 40, 0.5)` for $k = 0, \dots, 40$.

Check `binocdf(20, 40, 0.5) = 0.563` on graph.

Check `binoinv(0.56, 40, 0.5) = 20` on graph via $y = 0.56 \rightarrow \downarrow 20 = k$

Remark. A representation of `binocdf(k, n, p)` in terms of the *incomplete beta function* $I_x(a, b)$ is

$$\begin{aligned} \text{binocdf}(k, n, p) &= I_{1-p}(n - k, k + 1) \\ &= (n - k) \binom{n}{k} \int_0^{1-p} t^{n-k-1} (1 - t)^k dt, \end{aligned}$$

General information

- General mathematical information about the concept is $\triangleright$ WIKIPEDIA : Binomial distr.
- Syntax and semantic of the function is here $\triangleright$ MATLAB : `binopdf`
- Online calculator *Bognar's app*

2.1.1 Exercises

Exercise 58. (user-defined function `binPDF` in REDUCE)

Write function `binPDF()` using the mathematical definition for *binPDF* in REDUCE. Check your function on example §2.1-1.



View *code*, run REDUCE.

Exercise 59. (Table and plot for `binPDF`)

- Print 10 values for the binomial PDF for $n = 10$ and $p = 1/6$ and e.g. $i = 1$ thru 10.
- Plot the Binomial PDF values from a. - if necessary by paper and pencil.
- Do a plot for the binomial PDF for $n = 40$ and $p = 0.5$.



View *code*, run REDUCE.

Exercise 60. (user-defined function `binCDF` in REDUCE)

Write function `binCDF()` using the mathematical definition for *binCDF* in REDUCE. Check your function on the example §2.1-2.



View *code*, run REDUCE.

Exercise 61. (Table and plot for `binCDF`)

- Print 10 values for the binomial CDF for $n = 10$ and $p = 1/6$ and e.g. $i = 1$ thru 10.
- Plot the Binomial CDF values from a. - if necessary by paper and pencil.
- Do a plot for the binomial CDF for $n = 40$ and $p = 0.5$.



View *code*, run REDUCE.

Exercise 62. (a continuous version of *choose* and *binCDF* via Gamma function)

Pocket calculators often have a button `nCr` (=' `n` Choose `r`') to calculate the binomial coefficient $\binom{n}{r}$, in REDUCE `binomial(n,r)`.¹⁵ We now extend this function and define a 'continuous' version of *choose*() via Gamma function Γ and also extend the cumulative binomial probability to `binCDF2`, e.g. to plot graphs of the `binCDF`:

$$nCk(n, k) = \Gamma(n+1) / (\Gamma(k+1) \cdot \Gamma(n-k+1))$$

$$binCDF2(k, n, p) = \sum_{j=0}^k nCk(n, j) \cdot p^j \cdot (1-p)^{n-j}$$

where $\Gamma(x)$ is `gamma(x)|Reduce`.

- Define `nCk(n,k)` and `binCDF2(k,n,p)` in REDUCE and check it e.g. by `binCDF2(4, 6, 0.3); %ok: 0.9890`

- Draw `binCDF2(x, 10, 1/6)` using `xrange = (0,10)` and `yrange = (0,1)`.

¹⁵Maybe in this context you want to use the symbol `choose` instead of `binomial`. Hint by F.J. WRIGHT: 'You can literally define an alias in REDUCE by using the `define` command, e.g. `define choose = binomial;` which is actually more efficient than defining a new procedure just to provide a different name for a function.'

Remark.

In REDUCE, the *gamma* function is written as *gamma*(*z*). It evaluates symbolically, numerically, and can be expanded or simplified for various inputs.

For integer arguments $n \in \mathbb{N}$, REDUCE automatically evaluates to the exact *factorial value* $n!$. For example, *gamma*(5); evaluates to $24 = 5!$. Use `on rounded; gamma(x);` to evaluate the gamma function for real numbers $\mathbb{R}$ to receive a floating-point result.



View *code*, run REDUCE.

Exercise 63. (user-defined function **binoINV** in REDUCE)

Write a user-defined function **binoINV()** using the mathematical definition for *binoINV* in REDUCE.

The inverse binomial cumulative distribution function **binoINV()** results in the minimum number of successes, such that the cumulative binomial probability for that minimum number of successes is $\geq$ the given cumulative probability. Because there is no closed formula for **binoINV()** we do a simple search to solve the following typical problem to understand the construction of **binoINV()**, cf. finding values of *binoCDF*

- a. Given a number of trials $n = 100$, the probability of success $p = 0.3$, the cumulative area of the binomial distribution is $A = 0.7$, find the value x such that $A = 0.7 \leq \text{binoCDF}(100, 0.3, x)$.

To solve 'Task: search for k with $\text{binoCDF} \geq 0.7$ for 1st time' do a simple search via a loop e.g.

```
for k:=30:35 do write binoCDF(k,100,0.3) ;
```

and read off 'ok : $k=32$ '.

Verify EXCEL's solution: $x = 32$ by call = BINOM.INV(100, 0.3, 0.7).

- b. Now program **binoINV()** using e.g. the following pseudocode¹⁶

```
k=0
loop(k=k+1, x=binoCDF(k,100,0.3), x,k, test(x>0.95, break))
return(k)
```



View *code*, run REDUCE.

¹⁶which is in fact runnable code in CAS EIGENMATH^{online}

Exercise 64. (problem from Texas Instruments, cf.: TI)

Assume the toss of a fair coin 30 times. What is the minimum number of heads you must observe such that the cumulative binomial probability for that number of observed heads is at least 0.95?

Hint:

- (idea:) set `x = binCDF(k,30,0.5)` and use a table of values (starting at 0 and increment by 1) to find when the cumulative binomial probability is at or just above the given cumulative binomial probability. This gives you a view of all values to make decisions. For this example, search in the table to find the cumulative binomial probability just larger than 0.95. So, the number of successes is 19.
- (motivation for `binoINV`, quoting TI) The results on the screen first show that the minimum number of successes to obtain at least the given cumulative binomial probability of 0.95 is 19. Next, the cumulative probability for up to 19 is computed using `binoCDF` and is approximately 0.9506314271 which meets the criteria of $0.9506314271 \geq 0.95$. Abstracting this process we get our small function `binoINV2`.

```

procedure binoINV2(alpha,n,p);
  begin scalar x; x:=0; k:=0;
    repeat << k := k+1;
      x := binoCDF(k,n,p); >>
    until x >= alpha;
    return(k);
  end;

  binoINV(0.95, 30, 0.5);

```



View *code*, run REDUCE.

2.2 Geometric distribution

A geometric distribution is a discrete probability distribution that describes the chances of achieving success in a series of independent trials, each having two possible outcomes. The geometric distribution thus helps measure the probability of success after a given number of trials. In the binomial distribution, the number of trials is fixed, and we count the number of ‘successes’. Whereas, in the geometric distribution, the number of successes is fixed, and we count the number of trials needed to obtain the desired number of successes. The geometric distribution is a discrete analog of the exponential distribution. It is discrete, i.e. existing only on the nonnegative integers.

Definition Notation: $X \sim \text{Geometric}(n, p)$

- The probability density function of a geometric distribution is defined as:

$$\text{geoPDF}(\mathbf{k}, \mathbf{p}) := \Pr(X = n) = p(1 - p)^k, \quad k = 0, 1, 2, 3, \dots \text{ and } 0 < p < 1$$

where k is number of failures before the first success and p is the probability of success on a given trial.

- The cumulative geometric distribution can be expressed as:

$$\text{geoCDF}(\mathbf{n}, \mathbf{p}) := \Pr(X \leq k) = 1 - (1 - p)^n, \quad n = 1, 2, 3, \dots$$

- The *quantile* function (inverse cumulative geometric distribution) is¹⁷

$$\text{geoINV}(\mathbf{u}, \mathbf{p}) := \left\lceil \frac{\log(1 - u)}{\log(1 - p) - 1} \right\rceil$$

i.e. we must find the *smallest* n such that $u \leq 1 - (1 - p)^n$ for given u .

Examples from MATLAB, verified with REDUCE in exercises N° 65, 67, 69.

1. *geoPDF*: A man asking for help with a probability of getting help is 0.5 .Calculate the probability hat the person experiences 5 'failures' before the first success.

$$| \quad \text{geoPDF}(5, 0.5) \quad | \quad 0.015625$$

2. *geoCDF*: The probability of getting the help is 0.6. Calculate the probability that the person will have to talk to 8 or less people to find someone who helps.

$$| \quad \text{geoCDF}(8, 0.6) \quad | \quad .9997$$


¹⁷ $\lceil \dots \rceil$ is the ceiling function.

3. *geoINV*: Suppose the probability of a five-year-old car battery not starting in cold weather is 0.03. If we want no more than a ten percent chance that the car does not start, what is the maximum number of days in a row that we should try to start the car?

| `geoINV(0.1, 0.03)`

| 3

Graphical representation



Experiment: rolling a six-sided die.

Left figure: $\bullet$ = plot for `geoPDF(k, 1/6)` for $k = 0, \dots, 10$.

Check `geoPDF(4, 1/6) = 0.08` on the graph.

Figure 16: Right figure: $\bullet$ = plot for `geoCDF(k, 1/6)` for $k = 0, \dots, 10$.

Check `geoCDF(5, 1/6) = 0.665` on graph.

Check `geoINV(0.7, 1/6) = 6` on the CDF graph via $y = 0.6 \rightarrow \downarrow 6 = k$

General information

General mathematical information about the concept is here $\triangleright$ WIKI : `Geometric_distr.`

Syntax and semantic of the function is here $\triangleright$ MATLAB : `geopdf`

○ Online calculator *Bognar's app*

2.2.1 Exercises

Exercise 65. (user-defined function `geoPDF` in REDUCE)

Write function `geoPDF()` using the mathematical definition for *geoPDF* in REDUCE. Check your function on example §2.2-1.



View *code*, run REDUCE.

Exercise 66. (Table and plot for *geoPDF*)

- Print 10 values for the geometric PDF for $k = 1..10$ and $p = 1/6$.
- Plot the geometric PDF values from a. - if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 67. (user-defined function `geoCDF` in REDUCE)

Write function `geoCDF()` using the mathematical definition for *geoCDF* in REDUCE. Check your function on the example §2.2-2.



View *code*, run REDUCE.

Exercise 68. (Table and plot for *geoCDF*)

- Print 10 values for the geometric CDF for $n = 1..10$ and $p = 1/6$.
- Plot the geometric CDF values from a. - if necessary by paper and pencil.
- Do a plot for the geometric CDF for $n = 40$ and $p = 0.5$.



View *code*, run REDUCE.

Exercise 69. (user-defined function `geoINV` using explicit formula)

Write function `geoINV()` using the mathematical definition for *geoINV* in REDUCE.



View *code*, run REDUCE.

Exercise 70. (alternative user-defined function `geoINV` using loop construct)

Write a user-defined function `geoINV()` using a loop construct.

The inverse geometric cumulative distribution function `geoINV()` results in the minimum number of successes, such that the cumulative geometric probability for that minimum number of successes is $\geq$ the given cumulative probability. Although there is a closed formula for `binoINV()` in Exercise.69, we can do a simple search to solve the typical problem of example.3 to understand the construction of `binoINV()`, see exercise 64.



Run REDUCE.

Exercise 71. (inverse probability)

Verify, that the inverse probability at $pin = 0.25, p = 0.6835$ is 3.

2.3 Negative binomial distribution

The simplest motivation for the negative binomial is the case of successive random trials, each having a constant probability p of success. The number of extra trials you must perform in order to observe a given number r of successes has a negative binomial distribution. The negative binomial distribution is a discrete, i.e. existing only on the nonnegative integers.

Definition Notation: $X \sim NB(r, p)$

- The probability density function of a negative binomial distribution is defined as:

$$\text{nbPDF}(\mathbf{k}, \mathbf{r}, \mathbf{p}) := \binom{k+r-1}{k} p^r (1-p)^k =: f(k, r, p) = \Pr(X = k)$$

where k is number of failures before the first success and p is the probability of success on a given trial.

- The cumulative negative binomial distribution can be expressed as:

$$\text{nbCDF}(\mathbf{k}, \mathbf{r}, \mathbf{p}) := \sum_{i=0}^k \binom{i+r-1}{r-1} p^r (1-p)^i =: F(k, r, p) = \Pr(X \leq k)$$

- The *quantile* function (inverse cumulative negative binomial distribution) is

$$\text{nbINV}(\alpha, \mathbf{r}, \mathbf{p}) := \inf\{k \in \mathbb{R} : \alpha \leq F(k; r, p)\}$$

i.e. we must find the *smallest* k such that $\alpha \leq \sum_{i=0}^k \binom{i+r-1}{r-1} p^r (1-p)^i$ for given α .

`nbINV(alpha, size, prob)` returns the number of trials (or failures before the *size*-th success) such that the probability of observing that many or fewer failures is at least p .

Examples from Excel, DataCamp and MatLAB, verified with REDUCE in exercises N° 72, 75, 77..

1. *nbPDF*: In quality control, if we need 3 defective units and each unit has a 10% chance of being defective, what is the probability of getting exactly 5 non-defective units before finding the third defective one?

| `nbPDF(5, 3, 0.1)` | 0.0124

2. *nbCDF*: You have to identify five people who have excellent reflexes, you know the probability that any one candidate meets this requirement is 0.25. What is the probability that you will interview a certain number of unsuitable candidates before identifying five suitable candidates.


| `nbCDF(10, 5, 0.25)` | 0.3135

3. *nbinINV*: How many times would you need to flip a fair coin to have a 99% probability of having observed 10 heads?

| `nbinINV(0.99, 10, 0.5)`

| 23

Graphical representation



Left figure: $+$ = plot for `nbinPDF((x,3,0.5))` for $x = 0, 1, 2, \dots, 10$.

Check `nbinPDF((3,3,0.5)) = 0.15` on the graph.

Figure 17: Right figure: $+$ = plot for `nbinCDF((x,3,0.5))` for $k = 0, 1, 2, \dots, 10$.

Check `nbinCDF(4,3,0.5)) = 0.77` on graph.

Check `nbinINV(0.75,3,0.5) = 4` on graph via $y = 0.75 \longrightarrow \downarrow 4 = x$

General information

General mathematical information about the concept is here $\triangleright$ `nbin.distribution`

Syntax and semantic of the function is here $\triangleright$ MATLAB : `nbinpdf`

◦ Online calculator *Bognar's app*

2.3.1 Exercises

Exercise 72. (user-defined function `nbinPDF` in REDUCE)

Write function `nbinPDF()` using the mathematical definition for *nbinPDF* and check your function on example §2.3-1.



View *code*, run REDUCE.

Exercise 73. (Table and plot for *nbinPDF*)

- Print 10 values for the negative-binomial PDF for $k = 1..10$ and $p = 1/6$.
- Plot the negative-binomial PDF values from a. - if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 74. (candy problem, cf. WIKIPEDIA)

Pat Collis is required to sell candy bars to raise money for the 6th grade field trip. Pat is (somewhat harshly) not supposed to return home until five candy bars have been sold. So the child goes door to door, selling candy bars. At each house, there is a 0.6 probability of selling one candy bar and a 0.4 probability of selling nothing.

- What's the probability of selling the last candy bar at the n -th house?
- What's the probability that Pat finishes on or before reaching the eighth house?
- What's the probability that Pat exhausts all 30 houses that happen to stand in the neighborhood?



View *code*, run REDUCE.

Exercise 75. (user-defined function `nbinCDF` in REDUCE)

Write function `nbinCDF()` using the mathematical definition for *nbinCDF* in REDUCE.

You may (alternatively) use functions from the REDUCE package `specfn`.

Check your function on the example §2.3-2.



View *code*, run REDUCE.

Exercise 76. (Table and plot for *nbinCDF*)

- Print 10 values for the negative-binomial CDF for $n = 1..10$ and $p = 1/6$.
- Plot the negative-binomial CDF values from a. - if necessary by paper and pencil.
- Do a plot for the negative-binomial CDF for $n = 40$ and $p = 0.5$.



View *code*, run REDUCE.

Exercise 77. (user-defined function `nbinINV`)

Write function `nbinINV()` using the mathematical definition for *nbinINV*.

- Check your function on example §2.3-3.
- Verify, that your function solves $\text{nbinINV}(0.99, 5, 0.25) = 38$ and $\text{nbinINV}(0.95, 10, 0.7) = 9$ correctly.



View *code*, run REDUCE.

Exercise 78. (inverse probability)

Verify, that the inverse probability at $pin = 0.25, p = 0.6835$ is 3.

2.4 Hypergeometric distribution

‘Think of an urn with two colors of marbles, red and green. Define drawing a green marble as a success and drawing a red marble as a failure. Let N describe the number of all marbles in the urn and K describe the number of green marbles, then $N - K$ corresponds to the number of red marbles. Now, standing next to the urn, you close your eyes and draw n marbles without replacement. Define X as a random variable whose outcome is k , the number of green marbles drawn in the experiment.’, cf. [7].

The *hypergeometric* distribution is discrete, i.e. existing only on the nonnegative integers.

Definition Notation: $X \sim hg(N, R, n)$

- The probability density function of a hypergeometric distribution is defined as:

$$\text{hgPDF}(N, R, n, r) := \Pr(X = r) = \frac{\binom{R}{r} \binom{N-R}{n-r}}{\binom{N}{n}}$$

where r is number of failures before the first success and p is the probability of success on a given trial.

- The cumulative hypergeometric distribution can be expressed as:

$$\text{hgCDF}(N, R, n, x) := \Pr(X \leq x) = \sum_{r=0}^x \frac{\binom{R}{r} \binom{N-R}{n-r}}{\binom{N}{n}}$$

- The *quantile* function (inverse cumulative hypergeometric distribution) is

$$\text{hgINV}(\alpha, N, R, n) := \inf\{k \in \mathbb{R} : \alpha \leq \text{hgCDF}(N, R, n, k)\}$$

i.e. we must find the *smallest* k such that $\alpha \leq \Pr(X \leq k)$ for given α .

hgINV(alpha, N,R,n): You can think of α as the probability of observing r defective items in n drawings without replacement from a group of N items where R are defective.

Examples Verified with REDUCE in exercises N° 79, 85, 89.

1. *hgPDF*: What is the probability of selecting 14 red marbles from a sample of 20 taken from an urn containing 70 red marbles and 30 green marbles?

$$| \quad \text{hgPDF}(100, 70, 20, 14) \quad | \quad 0.21409$$

2. *hgCDF*: (MatLAB) Suppose you have a lot of 100 floppy disks and you know that 20 of them are defective. What is the probability of drawing two defective floppies if you select 10 at random?

$$| \quad \text{hgCDF}(100, 20, 10, 2) \quad | \quad 0.6812$$

3. *hgINV*: (MatLAB) Suppose you are the Quality Assurance manager for a floppy disk manufacturer. The production line turns out floppy disks in batches of 1,000. You want to sample 50 disks from each batch to see if they have defects. You accept 99% of the batches if there are no more than 10 defective disks in the batch. What is the maximum number of defective disks should you allow in your sample of 50?

| `hgINV(0.99, 1000, 10, 50)` | 3

Graphical representation




Figure 18: National lottery: from 49 including 6 red take 6 and get x red.
Left figure: $\bullet$ = plot for $\text{hgPDF}(49, 6, 6, x)$ for $x = 0, 1, 2, \dots, 6$.
 Check $\text{hgPDF}(49, 6, 6, 2) = 0.139$ on the graph.
Right figure: $\bullet$ = plot for $\text{hgCDF}(49, 6, 6, x)$ for $x = 0, 1, 2, \dots, 6$.
 Check $\text{hgCDF}(49, 6, 6, 1) = 0.84$ on graph.
 Check $\text{hgINV}(0.75, 49, 6, 6) = 1$ on graph via $y = 0.75 \rightarrow \downarrow 1 = x$

General information

[†1] General mathematical information is here $\triangleright$ hypergeometric. distr.
 Syntax and semantic of the function is here $\triangleright$ MATLAB : `hygepdf`
 ○ Online calculator *Bognar's app*

2.4.1 Exercises

◦ Unfortunately, the parameter list for the set of hypergeometric functions varies from one software package to another; therefore, we will address this peculiarity in the exercises – specifically regarding $\mathcal{R}$, MATLAB, and Excel. We have the

LEXICON	<i>this book</i>	<i>distrib</i>
follows convention of	MATLAB	$\mathcal{R}$
parameter list	N, R, n, r	$x, n1, n2, n$
mental model	urn	

Exercise 79. (user-defined function **hgPDF** in REDUCE)

Write function **hgPDF()** using the mathematical definition for *hgPDF*.

Check your function on example §2.4-1.



View *code*, run REDUCE.

Exercise 80. (Table and plot for hgPDF)

a. Calculate the 7 values for the hypergeometric PDF of Fig.18.left.

b. Plot the hypergeometric PDF values from a. - if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 81. (national lottery)


In a national lottery box there are 49 numbered balls.

6 balls are drawn.

a. What is the probability to predict 4 numbers correct of the 6?

b. Calculate the probabilities for exact $k = 0, \dots, 6$ 'red balls' (successes).

c. Plot the hypergeometrical PDF for the lottery outcomes.



View *code*, run REDUCE

Exercise 82. (example from MATLAB, hygepdf)

Reproduce

```
MATLAB> p = hygepdf(0:5,100,20,10)
p = 0.0951  0.2679  0.3182  0.2092  0.0841  0.0215
```



View *code*, run REDUCE.

Exercise 83. (WIKIPEDIA ibid.: working example)

The WIKIPEDIA definition of the hypergeometric PDF use another parameter sequence:

```
hgPDFwiki(r,N,R,n) := hgPDF(N,R,n,r)
```

We are interested in calculating the probability of drawing r red marbles in n draws, given that there are R red marbles out of a total of N marbles.

◦ Assume that there are 5 red and 45 green marbles in the urn. Standing next to the urn, you close your eyes and draw 10 marbles without replacement.

What is the probability that exactly 4 of the 10 are red?



View *code*, run REDUCE.

Exercise 84. (example with MS EXCEL, cf. Excel)

MS EXCEL use yet another parameter sequence for the hypergeometric PDF:

`hgPDFexcel(x, n, k, m) := hgPDF(m,k,n,x)`

in words: `hgPDFexcel(x,n,k,m)` is the probability of getting x successes from a sample of size n , where the size of the population is m of which k are successes.

EXAMPLES:

a. A bag contains 12 balls, 8 red and 4 blue. You reach into the bag and pick 3 balls at random without replacement.

What is the probability that at least 2 of the balls will be blue?

b. A warehouse contains 500 used computers. A random sample of 100 of these is tested and three of them are found to be defective.

What is the most likely percentage of defective computers in the warehouse?



View *code*, run REDUCE.

Exercise 85. (user-defined function `hgCDF` in REDUCE)

Write function `hgCDF()` using the mathematical definition for *hgCDF*.

Check your function on the example §2.4-2.

- `hgCDF(N,R,n, x)` computes the hypergeometric CDF at the value in x using the corresponding size of the population N , the number of the desired items in the population R , and the number of samples drawn n .



View *code*, run REDUCE.

Exercise 86. (table and plot of the hypergeometric CDF)

a. Print 9 values for the cumulative hypergeometric CDF for $n = 1..9$ and the parameter set $N = 100, R = 20, n = 10$.

b. Plot the hypergeometric CDF values from a. - if necessary by paper and pencil.

c. Plot the values of the hypergeometric CDF for the lottery outcomes, see Fig.18.right.



View *code*, run REDUCE.

Exercise 87. (defective floppy disks, cf. `hygecdf`)

Suppose you have a lot of 100 floppy disks and you know that 20 of them are defective.

What is the probability of drawing zero to two defective floppies if you select 10 at random?



View *code*, run REDUCE.

Exercise 88. (draw at least 3 numbers right)

What is the probability to draw at least 3 numbers right by drawing of 6 out of 49?



View *code*, run REDUCE.

Exercise 89. (user-defined function **hgINV** in REDUCE)

Write function **hgINV()** using the mathematical definition for *hgINV*.

Check your function on example §2.4-3.



Solution *code*, run REDUCE.

- This function finds the n .th quantile, that is if $Pr(x \leq k)$ is given, it finds k .
- **hgINV**(α, N, R, n) returns the smallest integer x such that the hypergeometric CDF evaluated at x equals or exceeds Pr . You may think of α as the probability of observing x defective items in n drawings without replacement from a group of N items, where R are defective.

Exercise 90. (Table and plot for hgINV)

a. Print 10 values for the hypergeometric INV for $\alpha = 0.1 \dots 1$ *step* 0.1

and $N = 100, R = 10, n = 50$.

b. Plot the hypergeometric INV values from a. - if necessary by paper and pencil.

c. Plot the values of the hypergeometric INV for the lottery outcomes, see Fig.18.right.



View *code*, run REDUCE.

Exercise 91. (A manufacturer produces shirts, cf. real-statistics)

A manufacturer produces shirts in batches of 3,000. They decide to sample 50 shirts from each batch and count the number of defective shirts in each sample. If they want to make sure that 99.5% of the batches have no more than 20 defects,

What is the maximum number of defects that they should allow in each sample?



View *code*, run REDUCE.

Exercise 92. (example from $\mathcal{R}$)

Reproduce in REDUCE:

```
R> qhyper(p, m, n, k, lower.tail = TRUE, log.p = FALSE)
R> qhyper(p = 0.90, m = 5, n = 11, k = 4)
[1] 2
```

Warning: watch the other parameter list.



View *code*, run REDUCE.

2.5 POISSON distribution

The Poisson distribution can be applied to systems with a large number of possible events, each of which is rare. The Poisson probability *density* function lets you obtain the probability of an event occurring within a given time or space interval exactly k times if on average the event occurs λ times within that interval.

The POISSON distribution is discrete, i.e. existing only on the nonnegative integers.

Definition Notation: $X \sim \text{Poisson}(k, \lambda)$

- The probability density function of a POISSON distribution is defined as:

$$\text{poissonPDF}(\mathbf{k}, \lambda) := \frac{\lambda^k e^{-\lambda}}{k!} = f(k, \lambda) = \Pr(X = k)$$

where k is number of failures before the first success and λ is the probability of success on a given trial.

- The cumulative POISSON distribution can be expressed as:

$$\text{poissonCDF}(\mathbf{x}, \lambda) := \sum_{k=0}^x \frac{\lambda^k}{k!} e^{-\lambda} = F(x; \lambda) = \Pr(X \leq x)$$

- The *quantile* function (inverse cumulative POISSON distribution) is

$$\text{poissonINV}(\alpha, \lambda) := \inf\{x \in \mathbb{R} : \alpha \leq \sum_{k=0}^x \frac{\lambda^k}{k!} e^{-\lambda}\}$$

i.e. we must find the *smallest* k such that $\alpha \leq \Pr(X \leq k)$ for given α .

`poissonINV`(α, λ) returns the smallest value k such that the Poisson CDF evaluated at k equals or exceeds p , using mean parameters in `lambda`.

Examples Verified with REDUCE in exercises N° 93, 98, 100.

1. *poissonPDF*: What is the probability of making 2 sales in a week if the average sales rate is 3 per week?

$$| \quad \text{poissonPDF}(2, 3) \quad | \quad 0.2240$$


2. *poissonCDF*: (MatLAB) A computer hard disk manufacturing facility performs random tests of individual hard disks. The policy is to shut down the manufacturing process if an inspector finds more than four bad sectors on a disk. Assuming that on average a disk has two bad sectors, find the probability of a manufacturing process shutdown after the first inspection.

$$| \quad 1 - \text{poissonCDF}(4, 2) \quad | \quad 0.0526$$

3. *poissonINV*: If the average number of defects is two, what is the 95th percentile of the number of defects?

| `poissonINV(0.95, 2)` | 5

Graphical representation



Plot of POISSON distribution for parameter $\lambda = 2.5$

Left figure: • = plot of `poissonPDF(k, 2.5)` for $k = 0, \dots, 10$.

Check `poissonPDF(2.5, 3) = 0.21` on the graph.

Right figure: • = plot of `poissonCDF(k, 2.5)` for $k = 0, \dots, 10$.

Check `poissonCDF(2.5, 3) = 0.75` on graph.

Check `poissonINV(0.7, 2.5) = 3` on the CDF via $y = 0.7 \rightarrow \downarrow 3 = k$

Figure 19:

General information

- General mathematical information about the concept is here ▷ WIKI : Poisson.distr.
- Syntax and semantic of the function is here ▷ MATLAB : `poisspdf`
- Online calculator *Bognar's app*

2.5.1 Exercises

Exercise 93. (a user-defined function `poissonPDF` in REDUCE)

Write a user-defined function `poissonPDF()` using the mathematical definition for *poissonPDF*. Check your function on example §2.5-1.



View *code*, run REDUCE.

Exercise 94. (Table and plot for `poissonPDF`)

- Calculate the 11 values of the `poissonPDF` shown in fig.19.left.
- Plot the `poisson PDF` values from a., if necessary by paper and pencil. See fig.19.



View *code*, run REDUCE.

Exercise 95. (examples from WIKIPEDIA cf. *poisson*)

Solve the examples ('Beispiele') in [ibid.]



View *code*, run REDUCE.

Exercise 96. (examples from Hermann [7, p.42])

- On average, a book contains one typo error on every page. What is the probability that there are two typographical errors on a given page? (solution: 18.4 %)
- What is the probability that at a party with 25 people, everyone will have different birthdays?

Hint: 25 people gives $\binom{25}{2} = 300$ birthday-pairs. So $\lambda = 300/365$. (Solution: 56%)



View *code*, run REDUCE.

Exercise 97. (example from Hermann [7, p.42])

In a city, there are an average of two accidents per day.

What is the probability of more than four accidents per day?

Solution: $\lambda = 2$. $\Pr(X_i > 4) = 1 - \Pr(X_i \leq 4) = \dots = 5.2\%$



View *code*, run REDUCE.

Exercise 98. (user-defined function `poissonCDF` in REDUCE)

Write a user-defined function `poissonCDF()` using the mathematical definition for *poissonCDF*. Check your function on the example §2.5-2.

- An alternative implementation uses the incomplete Gamma function.



View *code*, run REDUCE.

Exercise 99. (table and plot of the `poisson CDF`)

- Calculate the 11 values of the `poissonCDF` shown in fig.19.right.
- Plot the `poisson CDF` values from a., if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 100. (user-defined function `poissonINV` in REDUCE)

Define a function `poissonINV()` using the mathematical definition for `poissonINV`.

Check your function on example §2.5-3.



View *code*, run REDUCE.

Remark.

- This function is used to find the n .th quantile, that is if $Pr(x \leq k)$ is given, it finds k .
- `X=poissonINV(alpha, lambda)` returns the smallest value X such that the Poisson CDF evaluated at X equals or exceeds P , using mean parameter in `lambda`. [ibid.]

Exercise 101. (Table and plot for `poissonINV`)

- Print 10 values for the poisson INV for $\alpha = 0.1 \dots 1$ *step* 0.1 and $\lambda = 2.5$.
- Plot the poisson INV values from a. - if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 102. (defects in a machine part, cf. `poissinv`)

Suppose the average number of defects in a machine part is two.

- Use the inverse of the Poisson distribution to determine the 95th percentile of the number of defects.
- Calculate the median number of defects per machine part.



View *code*, run REDUCE.



You should now be able to implement other discrete distributions along the above examples using CAS REDUCE.



3 Continuous distributions

Regarding the inverse distributions and quantiles x_α of the continuous cumulative distribution function F we present in the following sections, we have the mathematical relation

$$F(x_\alpha) = \alpha \implies x_\alpha = F^{-1}(\alpha)$$

where the value x_α is referred to as the α -quantile of F . Calculating this quantile requires the inverse function F^{-1} . Since the cumulative distribution function F itself is typically computed using special functions or series expansions, *there are no exact methods for determining the inverse function*; instead, one must rely on approximation methods. In principle, calculating a quantile of a distribution function F could be reduced to solving a transcendental equation using Newton's iteration method.¹⁸

$$x_{n+1} = x_n - \frac{F(x_n)}{F'(x_n)}$$

However, this iteration is often numerically unstable because the probability density function $f = F'(x)$ in the denominator can become arbitrarily small, cf. [7, p.69].

The pattern is to be seen here, cf. section §3.4 for details.¹⁹

The important point is that **inverse Student's t is not obtained by simply algebraically inverting the hypergeometric expression**. In general, there is no elementary closed form for

$$t = F_\nu^{-1}(p).$$

So in REDUCE 3.8 you need to solve

$$F_\nu(t) - p = 0$$

numerically.

3. Newton iteration

A particularly convenient fact is that the derivative of the CDF is simply the t-density:

$$f_\nu(t) = \frac{\Gamma((\nu+1)/2)}{\sqrt{\nu\pi} \Gamma(\nu/2)} \left(1 + \frac{t^2}{\nu}\right)^{-(\nu+1)/2}.$$

Thus Newton's method is

$$t_{k+1} = t_k - \frac{F_\nu(t_k) - p}{f_\nu(t_k)}.$$

Claude AI: Some info about STUDENTS's t distribution

¹⁸In REDUCE one calls function `num.solve()` from package `numeric` (which is autoloaded).

¹⁹Read here and in the following REDUCE 3.8 as REDUCE. Version 3.8 was the last commercial version of REDUCE, released in 2004. This version number does not matter w.r.t. the AI question, but served to distinguish the name REDUCE from the `reduce()` function in other programming languages.

3.1 Normal distribution

In probability theory and statistics, the *Normal Distribution*, also called the *Gaussian Distribution* by physicists, is the most significant continuous probability distribution. Social scientists refer to it as the *bell curve*, because of its curved flaring shape. A large number of random variables are either nearly or exactly represented by the normal distribution, in every physical science and economics.

- In a normal distribution, the mean, median and mode are equal.
- The normal distribution curve is symmetric w.r.t. the centre.

The *standard normal* distribution is one of the forms of the normal distribution. It occurs when a normal random variable has mean $\mu = 0$ and a standard deviation $\sigma = 1$.

Definition Notation: $X \sim \mathcal{N}(\mu, \sigma)$

- The probability *density* function of the *normal* distribution is defined by²⁰

$$\phi(\mathbf{x}, \mu, \sigma) := \frac{1}{\sqrt{2\pi} \cdot \sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}} := \Pr(X = x) =: \text{normPDF}(\mathbf{x}, \mu, \sigma)$$

- The *cumulative* normal distribution is defined by²¹

$$\Phi(x, \mu, \sigma) := \frac{1}{2} + \frac{1}{2} \cdot \text{erf}\left(\frac{x - \mu}{\sigma\sqrt{2}}\right) := \Pr(X \leq x) =: \text{normCDF}(\mathbf{x}, \mu, \sigma)$$

- The *inverse* normal distribution ('quantile function') is defined by

$$\Phi^{-1}(p, \mu, \sigma) := \mu + \sigma \cdot \Psi^{-1}(p) = \mu + \sigma\sqrt{2} \text{erf}^{-1}(2p - 1) = \text{normINV}(\mathbf{p}, \mu, \sigma)$$

with $p \in (0, 1)$, using the *inverse standard* normal function

$$\Psi^{-1}(p) = \sqrt{2} \text{erf}^{-1}(2p - 1)$$

Remark. REDUCE revision 7401 provides `erfinv` and `erfcinv`. We implement Φ^{-1} also using Φ and the defining relation $x = \Phi^{-1}(p)$ with $\Phi(x) = p$ and a search method..

²⁰we adopt the notation convention of MATLAB, i.e. `normXXX` with $\text{XXX} \in \{\text{PDF}, \text{CDF}, \text{INV}\}$

²¹using the build-in error-function `erf`

Examples Verified with REDUCE in exercises N° 103, 105, 114.

1. *normPDF*: A sample of 30 students has an average test score of 78 with a standard deviation of 12. Assuming the distribution of test scores is normal, what is the probability that the sample mean score is greater than 82?

| `normPDF(82, 78, 12)` | 0.0314


2. *normCDF*: A factory produces 100 Ω resistors, distributed via $N(100, 6^2)$. What proportion of the resistors have a maximum deviation of 10% from the mean value? Cf. [7, p. 50]

| `normCDF(110, 100, 6) - normCDF(90, 100, 6)` | 0.9044

3. *normINV*: Find the height at which 80% of the population falls below in a normal distribution with mean 170 and standard deviation 5.

| `normINV(0.80, 170, 5)` | 174.2

Graphical representation



Standard normal distribution with parameters $\mu = 0$ and σ equal to 1.

Left figure: `λ` = plot of `stdnormPDF(x)`.

Figure 20: Check `stdnormPDF(1) = 0.24` on the graph.

Right figure: `λ` = plot of `stdnormCDF(x)`.

Check `stdnormCDF(1) = 0.84` on graph.

Check `stdnormINV(0.6) ≈ 0.25` on the CDF via $y = 0.6 \rightarrow \downarrow 0.25 = x$

General information

- General mathematical information about the concept is here $\triangleright$ Normal.distribution
- Syntax and semantic of the function is here $\triangleright$ MATLAB : `normpdf`
- Online calculator *Bognar's app*

3.1.1 Exercises

Exercise 103. (user-defined function `normPDF` in REDUCE)Write function `normPDF()` and make use of the mathematical definition for $\phi(x, \mu, \sigma)$.

Check your function on example §3.1-1.



View code, run REDUCE.

Exercise 104. (Table and plot for [standard] normal PDF)a. Reproduce the following $\mathcal{R}$ session in REDUCE, where `dnorm` is the density of the normal distribution, i.e. our normal PDF.

```
R> # Create 10 numbers between 0 and 2 incrementing by 0.1:
R> x <- seq(0, 1, by = .1)
R> # Choose the mean as 0 and standard deviation as 1:
R> dnorm(x, mean = 0, sd = 1)
[1] 0.39894228 0.39695255 0.39104269 0.38138782 0.36827014 ..
```

b. Plot the PDF of the Standard Normal distribution `normPDF(x, 0, 1)` for $xrange = (-5, 5)$ and $yrange = (0, 1)$.

View code, run REDUCE.

Exercise 105. (user-defined function `normCDF` in REDUCE)A user-defined function `normCDF()` make use of the mathematical definition for $\Phi(x, \mu, \sigma)$, using the `erf` function of the special function package `specfn` of REDUCE.

Check your function on example §3.1-2.



View code, run REDUCE.

Exercise 106. (table and plot of the stdnormal CDF)a. Calculate the 7 values of the `stdnormCDF` at $-3, -2, \dots, 2, 3$ shown in fig.20.right.

b. Plot the stdnormal CDF values from a., if necessary by paper and pencil.



View code, run REDUCE.

Exercise 107. (examples from MATLAB)Construct a table for `normCDF` for $x = -2, -1, 0, 1, 2$ and $\mu = 2; \sigma = 1$ and verify the MATLAB session:

```
MatLAB> x = [-2,-1,0,1,2]; mu = 2; sigma = 1;
> p = normcdf(x,mu,sigma)
p = 1x5
    0.0000    0.0013    0.0228    0.1587    0.5000
```



View code, run REDUCE.

Exercise 108. (sugar in packs from [7, p.50])A manufacturer fills sugar in packs, which have a distribution of $\mathcal{N}(1000, 12^2)$.

What percentage of the packs meet the printed minimum weight of 980 g?



View code, run REDUCE.

Exercise 109. (examples from $\mathcal{R}$)

a. Calculate the probability of a value less than or equal to 1.96 in a standard normal distribution. $\triangleright$ `R> pnorm(1.96) [1] 0.9750021`

b. Determine the probability of a value less than or equal to 2 in a normal distribution with mean 1 and standard deviation 0.5? $\triangleright$ `R> pnorm(2, mean = 1, sd = 0.5) [1] 0.9772499`



View *code*, run REDUCE.

Exercise 110. (resistors from Hermann [7, p.51])

A factory produces 100 Ω resistors, distributed via $\mathcal{N}(100, 6^2)$.

What proportion of the resistors have a maximum deviation of 10% from the mean value?



View *code*, run REDUCE.

Exercise 111. (assortment of screws from [7, p.42])

From an assortment of screws distributed via $\mathcal{N}(10, 0.3^2)$ in a all screws larger than 10.5 mm or smaller than 9.5 mm are removed.

What is the probability that one of the remaining screws is smaller than 10.3 mm?



View *code*, run REDUCE.

Exercise 112. (men's height from [7, p.50])

When measuring men, the mean height was 1.73 m with a spread of 7 cm. Men between the heights of 166 cm and 180 cm are therefore considered to be of 'normal' height.

What is the probability that a man is shorter than 1.85 m according to this distribution?



View *code*, run REDUCE.

Exercise 113. (alternative definition of standard normal distribution using `erfc(z)`)

`erfc(z)` is the so-called Complementary Error Function. This function is used by MATLAB to build `normcdf()`, cf. `normcdf`.

a. Write function `stdnormCDF1` via `erfc(z)`. Use the facts

$$\begin{aligned}\text{stdnormCDF1}(x) &= 1/2 * \text{erfc}(-x/\text{sqrt}(2)) \\ \text{normCDF1}(x, \mu, \sigma) &= \text{stdnormCDF1}((x - \mu)/\sigma)\end{aligned}$$

b. Solve exercise.109 using `normCDF1`.



View *code*, run REDUCE.

Exercise 114. (user-defined function `normINV` in REDUCE)

To code function `normINV()`, first implement $\text{erf}^{-1} \equiv \text{erfINV}$ and then use the mathematical definition from p.66. We show three different solutions.

Check your function on example §3.1-1.



a. View *code* found by Claude AI, run REDUCE.



b. View *code* for a user-defined numeric inverse by F.J. WRIGHT, run REDUCE.



c. View *code* using `num_solve` from package `numeric`, run REDUCE.

Remark.

- This function finds the n^{th} quantile, that is if $P(x \leq k)$ is given, it finds k .
- The inverse normal distribution `normINV` allows you to find the value of x of a normal distribution given a probability p , the mean μ and the standard deviation σ of the distribution. Simply put, it works backward from the area under the normal distribution curve to find the x -value corresponding to that area.
- There is no simple algebraic formula; instead, one may use an iterative search technique to find the x , that satisfies `p = normCDF(x,mu,sigma)`. The standard normal version is a special case where the mean is 0 and the standard deviation is 1.

Exercise 115. (Table and plot for `std.normINV`)

- Print 9 values for the `std.normINV` for $\alpha = 0.1 \dots 0.9$ step 0.1.
- Plot the `std.normINV` values from a., if necessary by paper and pencil.

View *code*, run REDUCE.**Exercise 116.** (MATLAB: `norminv`)

Compute the inverse of cdf values evaluated at the probability values in p for the normal distribution with mean $\mu = 2$ and standard deviation $\sigma = 1$.

```
p = 0:0.25:1; mu = 2; sigma = 1;
norminv(p,mu,sigma)
ans = -Inf    1.3255    2.0000    2.6745    Inf
```

View *code*, run REDUCE.**Exercise 117.** (Compare with R `qnorm` list)

Reproduce with REDUCE:

```
R> # Create 10 numbers between 0 and 1 incrementing by 0.1.
R> x <- seq(0, 1, by = .1)
R> # Choose the mean as 0 and standard deviation as 1.
R> qnorm(x, mean = 0, sd = 1)
[1] -Inf -1.2815516 -0.8416212 -0.5244005 -0.2533471  0.0 ..
```

View *code*, run REDUCE.**Exercise 118.** (confidence interval, cf. MATLAB)

Find an interval that contains 95% of the values from a standard normal distribution.

```
Matlab> x = norminv([0.025    0.975])
x = 1x2    -1.9600    1.9600
Interval = [-1.96, 1.96]
```

Use REDUCE.

View *code*, run REDUCE.

Exercise 119. (some Quantile of normal/GAUSS distribution)

Verify, that for $X = (0.9, 0.95, 0.975, 0.99)$ we get the following quantile of the normal aka GAUSS distribution: $R = (1.281552, 1.644854, 1.959964, 2.326348)$ in $\mathcal{R}$.



View *code*, run REDUCE.

Exercise 120. (The 68-95-99.7 rule)

Let us determine the area under the standard normal curve for ± 1 standard deviation, for ± 2 standard deviations, and for ± 3 standard deviations.

Remark. Awesome, we just confirmed the Empirical Rule, also known as the "68-95-99.7" rule, which relates to the CHEBYSHEV's theorem.

For a bell-shaped distribution the 3 rules are, that approximately

- 68% of the observations lie within 1 std.deviation of the mean,
- 95% of the observations lie within 2 std.deviation of the mean,
- 99.7% of the observations lie within 3 std.deviation of the mean.



View *code*, run REDUCE.

The exponential distribution is a probability distribution that is used to model the time we must wait until a certain event occurs.

Definition Notation: $X \sim \text{Exp}(\lambda)$

- $$\text{expPDF}(\mathbf{x}, \lambda) := f(x, \lambda) = \Pr(X = x) = \begin{cases} \lambda e^{-\lambda x} & : x \geq 0, \\ 0 & : x < 0 \end{cases}$$

- The *cumulative* Exponential distribution is

$$\text{expCDF}(\mathbf{x}, \lambda) := F(x, \lambda) = \Pr(X \leq x) = \begin{cases} 1 - e^{-\lambda x} & : x \geq 0, \\ 0 & : x < 0 \end{cases}$$

- $$\text{expINV}(\mathbf{p}, \lambda) := F^{-1}(p, \lambda) = \frac{-\ln(1-p)}{\lambda}, \quad 0 \leq p < 1$$

1. *expPDF*: Compute the density of the observed value 5 in the exponential distribution specified by mean 3.

expPDF(5,1/3)	0.0629
---------------	--------

2. *expCDF*: The lifespan of a light bulb is exponentially distributed and averages 1,000 hours. What is the probability, that it will last 800 hours?


1 - expCDF(800, 1/1000)	0.4493
-------------------------	--------

3. *expINV*: (MatLAB) Assume that the lifetime of light bulbs are exponentially distributed with a mean of 700 hours. Find the median lifetime.

expINV(0.50, 1/700)	485.2
---------------------	-------

²²The ‘quantile’ function of a distribution is the inverse of the cumulative distribution function.

Graphical representation



Exponential distributions with different rates λ .

Left figure: λ = plot of $\text{expPDF}(x, \lambda)$, $\lambda = 1; 1.5; 5$.

Check $\text{expPDF}(1) = 0.24$ on the graph.

Figure 21: **Right** figure: λ = plot of $\text{expPDF}(x, \lambda)$, $\lambda = 1; 1.5; 5$.

Check $\text{expPDF}(2, 1) \approx 0.14$ on the graph.

Check $\text{expINV}(0.8, 1) \approx 1.6$ on the CDF via $y = 0.8 \rightarrow \downarrow 1.6 = x$

General information

- General mathematical information about the concept is here $\triangleright$ Exponential.distribution
- Syntax and semantic of the function is here $\triangleright$ MATLAB : `explpdf`
- Online calculator *Bognar's app*

3.2.1 Exercises

Exercise 121. (user-defined function `expPDF` in REDUCE)Write function `expPDF()` using the mathematical definition for *expPDF*.`expPDF` returns the pdf of the exponential distribution with rate L , evaluated at x .

a. Check your function on example §3.2-1.

b. Verify: `R> dexp(5,1/2)` [1] 0.0410425

View code, run REDUCE.

Exercise 122. (Table and plot for `expPDF`)a. Calculate the 11 values of the `expPDF` shown in fig.21.left for $\lambda = 5$.b. Plot the `expPDF` values from a., if necessary by paper and pencil. See fig.21.

View code, run REDUCE.

Exercise 123. (lifetime of a product; from [1, p.68])The lifetime of a product with the parameter $L = 1/2$ is exponentially distributed.

What is the probability that the product will fail in a period between three and five years?



View code, run REDUCE.

Exercise 124. (user-defined function `expCDF` in REDUCE)Write function `expCDF()` using the mathematical definition for *expCDF*.Check your function on the example `expCDF(4,2)`.

View code, run REDUCE.

Exercise 125. (table and plot of the `expCDF`)a. Calculate the 11 values of the `expCDF` shown in fig.21.right.b. Plot the `expCDF` values from a., if necessary by paper and pencil.

View code, run REDUCE.

Exercise 126. (lifespan of a light bulb, cf. [7, p.63])

The lifespan of a light bulb is exponentially distributed and averages 1,000 hours.

What is the probability, that it will last 800 hours?



View code, run REDUCE.

Exercise 127. (mean time between failures, cf. [7, p.63])

A printer has a mean time between failures ('MTBF') of 2,000 hours.

What is the probability that the first 500 hours of the trip will be trouble-free?



View code, run REDUCE.

Exercise 128. (refrigeration unit, cf. [7, p.63])

A refrigeration unit has a mean exponentially distributed lifespan of ten years.

What is the probability that it will last another five years, if it survives the first five years?



View code, run REDUCE.

Exercise 129. (function `expINV` in REDUCE)

Write function `expINV()` using the mathematical definition for `expINV`.

Check your function on example §3.2-3.



View *code*, run REDUCE.

Remark.

- MATLAB: The result x is the value such that an observation from an exponential distribution with parameter m will fall in the range $[0, x]$ with probability p .
- the quantile function (aka inverse CDF) is given by $F^{-1}(u) = -(1/L) * \ln(1 - u)$, where $0 < u < 1$. This inverse function takes a uniform random variable (between 0 and 1) as input and produces a value that would be expected from the exponential distribution.
- This function finds the n^{th} quantile, that is if $P(x \leq k)$ is given, it finds k .
- `x = expINV(p,lambda)` returns the smallest value x such that `expCDF` evaluated at x equals or exceeds p , using mean parameter in `lambda`.

Exercise 130. (Table and plot for `expINV`)

a. Print 10 values for the `expINV` for $p = 0.1 \dots 1$ *step* 0.1 and $\lambda = 2.5$.

b. Plot the `expINV` values from a., if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 131. (median lifetime, from `expinv`)

Assume that the lifetime of light bulbs are exponentially distributed with a mean of 700 hours. – Find the median lifetime using `expinv`.

Verify:

```
MATLAB> expinv(0.50,700)
ans = 485.2030
.. Half of the light bulbs will burn out within the first 485 hours of use.
```

```
R> qexp(0.5, 1/700)
[1] 485.203
```



View *code*, run REDUCE.

3.3 Student's t -distribution

As we know, the normal distribution assumes two important characteristics about the dataset: a large sample size and knowledge of the population standard deviation. However, if we do not meet these two criteria, and we have a small sample size (i.e. the sample size is 30 or less than 30) or an unknown population standard deviation, then we use the t -distribution. It is similar to the standard normal distribution ('Z'-distribution), but it has heavier tails. The t -score represents the number of standard deviations the sample mean is away from the population mean. ▷ g4g.

Definition Notation: $X \sim T(t, \nu)$

- The probability density function of Student's t -distribution is defined as:

$$\text{tPDF}(\mathbf{t}, \nu) := \frac{1}{\sqrt{\nu} B\left(\frac{1}{2}, \frac{\nu}{2}\right)} \left(1 + \frac{t^2}{\nu}\right)^{-(\nu+1)/2}$$

where ν is the number of degrees of freedom, $t \in (0, 1)$ is a probability value and B is the beta function.

- The *cumulative* probability of Student's t -distribution can be expressed as

$$\text{tCDF}(\mathbf{t}, \nu) := I\left(\frac{t + \sqrt{t^2 + \nu}}{2\sqrt{t^2 + \nu}}, \frac{\nu}{2}, \frac{\nu}{2}\right)$$

using the regularized incomplete beta function $I(x, a, b) := I_x(a, b)$.

- The *quantile* function (inverse cumulative Student's t -distribution) is approximately

$$\text{tINV}(\alpha, \nu) \approx \sqrt{\nu \cdot \exp(c \cdot u_\alpha^2)} \quad , \text{ where } c := \frac{\nu - \frac{5}{6}}{(\nu - \frac{2}{3} + \frac{1}{10\nu})^2}$$

where u_α is the α -quantile of the standard normal distribution.

We use this approximation formula by PREIZER & PRATT, [7, p. 70] in an exercise. ²³

The approximation error for $0.5 < \alpha < 0.99$ and $\nu \geq 3$ is maximal 0.08.

Instead we implement Student's t inverse function using Student's t CDF via the defining relation $F := \Pr(X \leq k)$ and $x = F^{-1}(p, \nu)$ with $F(x, \nu) = p$ and a classic search method.

²³See also G. W. HILL: ACM Algorithm 396. <https://dl.acm.org/doi/10.1145/355598.355599>.
A formula for the quantile function of the t -distribution does not exist in a closed form.

Graphical representation




Figure 22: Student's t distributions with different degrees of freedom $df := \nu$.
Left figure: λ = plot of $\mathbf{tPDF}(x, \nu)$, $\nu = 1; 2; 5$.
 Check $\mathbf{tPDF}(0, 1) = 0.3$ on the graph.
Right figure: λ = plot of $\mathbf{tCDF}(x, \nu)$, $\nu = 1; 2; 5$.
 Check $\mathbf{tCDF}(0, 1) \approx 0.5$ on the graph.
 Check $\mathbf{tINV}(0.2, 1) \approx -1$ on the CDF via $y = 0.2 \rightarrow \downarrow -1 = x$

Examples Verified with REDUCE in exercises N° 132, 136, 139.

1. *tPDF*: (MatLAB) The mode of the Student's t distribution is at $x = 0$. Compute the pdf at the mode for degree of freedom 3.

<code> tPDF(0,3)</code>	<code> 0.3675</code>
--------------------------	-----------------------
2. *tCDF*: (MatLAB) Determine the probability that an observation from the Student's t distribution with degrees of freedom 99 falls on the interval $[10, \dots, \text{Inf}]$.

<code> 1 - tCDF(10,99)</code>	<code> 0</code>
--------------------------------	------------------
3. *tINV*: (MatLAB) Compute the 99th percentile of the Student's t distribution for 3 degrees of freedom.

<code> tINV(.99, 3)</code>	<code> 4.5407</code>
-----------------------------	-----------------------

General information

- General mathematical information about the concept is here $\triangleright$ Student-t.distribution
- Syntax and semantic of the function is here $\triangleright$ MATLAB : `tpdf`
- Online calculator *Bognar's app*

3.3.1 Exercises

Exercise 132. (user-defined function **tPDF** in REDUCE)

Write function **tPDF()** for **student_t** distribution using the mathematical definition for *tPDF* and the special function $Beta(a, b)$ from **specfn** of REDUCE.

There is no simple term. Check your function on example §3.3-1.



View *code*, run REDUCE.

Exercise 133. (Table and plot for **tPDF**)

a. Calculate the 11 values of the **tPDF** shown in fig.22.left.

b. Plot the **tPDF** values from a.



View *code*, run REDUCE.

Exercise 134. (degree of freedom table for Student's t distribution from MATLAB)

The value of the pdf at the mode is an increasing function of the degrees of freedom.

Let the mode of the Student's t distribution at x be 0.

Compute the pdf at the mode for degrees of freedom 1 to 6 with REDUCE.

```
MATLAB> tpdf(0,1:6)
ans = 0.3183 0.3536 0.3676 0.3750 0.3796 0.3827
```



View *code*, run REDUCE.

Exercise 135. (t distribution converges to the standard normal distribution, cf. poisson)

The t distribution converges to the standard normal distribution as the degrees of freedom approach infinity.

Compute the difference between the pdf's of the standard normal distribution and the Student's t distribution pdf with 30 degrees of freedom.



View *code*, run REDUCE.

Exercise 136. (user-defined function **tCDF** in REDUCE)

Write function **tCDF()** for cumulate **student_t** distribution using the mathematical definition for *tCDF* and the special function **iBeta(a,b,x)**. There is no simple term.

Check your function on the example §3.3-2.



View *code*, run REDUCE.

Exercise 137. (table and plot of the student-t-CDF)

a. Calculate the 11 values of the **tCDF** shown in fig.22.right.

b. Plot the poisson CDF values from a., if necessary by paper and pencil.



View *code*, run REDUCE.

Exercise 138. (example from MATLAB, cf. *tcdf*)

Determine the probability that an observation from the Student's t distribution with degrees of freedom 99 falls on the interval $[10 \text{ Inf}]$.



View *code*, run REDUCE.

Exercise 139. (user-defined function `tINV` in REDUCE)

Write function `tINV()` for inverse `student_t` distribution using the mathematical definition for `tINV`. This function finds the n^{th} quantile, that is if $Pr(x \leq k)$ is given, it finds k . There is no simple term. You may use REDUCE's function `num_solve()`, cf. `num_solve`. Check your function on example §3.3-3.



Solution *code*, run REDUCE.

Exercise 140. (user-defined search methods for `tINV`)

The important point is that **inverse Student's t is not obtained by simply algebraically inverting the hypergeometric expression**. In general, there is no elementary closed form for

$$t = F_{\nu}^{-1}(p).$$

So in REDUCE 3.8 you need to solve

$$F_{\nu}(t) - p = 0$$

numerically.

3. Newton iteration

A particularly convenient fact is that the derivative of the CDF is simply the t-density:

$$f_{\nu}(t) = \frac{\Gamma((\nu+1)/2)}{\sqrt{\nu\pi} \Gamma(\nu/2)} \left(1 + \frac{t^2}{\nu}\right)^{-(\nu+1)/2}.$$

Thus Newton's method is

$$t_{k+1} = t_k - \frac{F_{\nu}(t_k) - p}{f_{\nu}(t_k)}.$$

inverse Student t Distribution using hypergeometric function in reduce 3.8

🕒 Nachdenken läuft...

Background

The Student-t CDF has a clean closed form in terms of the Gaussian hypergeometric function ${}_2F_1$:

$$F(t;\nu) = \frac{1}{2} + \frac{t \Gamma\left(\frac{\nu+1}{2}\right)}{\sqrt{\pi\nu} \Gamma\left(\frac{\nu}{2}\right)} {}_2F_1\left(\frac{1}{2}, \frac{\nu+1}{2}; \frac{3}{2}; -\frac{t^2}{\nu}\right)$$

The **inverse** (quantile function) has *no* closed form in terms of ${}_2F_1$ or any other standard special function — it's defined implicitly as the root of $F(t;\nu) = p$. So “inverse t via hypergeometric function” in practice means: use the ${}_2F_1$ form of the CDF (and its derivative, the PDF) inside a numeric root-finder (Newton's method).

Figure 23: Some info about STUDENTS's t distribution from Claude AI.

Implement Student's t inverse function `tinvt` using Student's cumulative distribution $F = \text{tCDF}$ via the defining relation $x = F^{-1}(p, \nu)$ with $F(x, \nu) = p$ and a classic search method, i.e. a *bisection* method or a NEWTON-RAPHSON iteration.

Check your function on example §3.3-3.



a. Solution *code*, bisection method found by *ChatGPT AI*, run REDUCE.



b. Solution *code*, **Newton** iteration; optimized by F.J. WRIGHT, run REDUCE.

Exercise 141. (approximation for **tINV** after PREIZER & PRATT)

Because there is no simple term for the inverse **student_t** distribution **tINV()**, implement the following approximation after PREIZER & PRATT, cf. [7, p.71].

Here is C code for the formula of p.76. Translate it to REDUCE .

```
double tinv(double x,int f)
/* Berechnung der inversen t-Verteilung
   mittels Approximation von Peizer & Pratt fuer f>=3
   Eingangsparameter: Wahrscheinl.0<x<1 vom Typ double
                     Freiheitsgrad f vom Typ int
   Ausgangsparameter: - keine -
   Rueckgabewert : Quantil vom Typ double
   benoetigt Funktion norminv() */
{
  double q,c,t;
  double PI = 3.1415926535898;
  double norminv();
  /* Error: 5.75 %/
  if (f==1) return tan(PI*(x-0.5));
  if (f==2)
    { t = 2.*x-1.;
      return sqrt(2.)*t/sqrt(1.-t*t);
    }
  q = norminv(x);
  c = f-2./3.+1./(10.*f); /* 1.Phase für c : (5.3) Nenner
  c = (f-5./6.)/(c*c); /* 2.Phase für c : Gr. (6.3) */
  t = sqrt(f*exp(c*q*q)-f);
  return (x>0.5)? t: -t;
}
```

Write **tINV()** in REDUCE and test it w.r.t.

```
MATLAB> tinv(.95,50)      ans = 1.6759
R> qt(.95,50)             [1] 1.675905
```



View code, run REDUCE.

Exercise 142. (approximation for **tINV** by HILL & CORNISH-FISHER expansion)

View the code, make it running and check this approximation for the inverse **student_t** distribution **tINV()** w.r.t. the approximation by HASTINGS in the exercise before.

Test with this reference values:

```
MATLAB> tinv(.95,50)      ans = 1.6759
R> qt(.95,50)             [1] 1.675905
```



View code, run REDUCE.

Exercise 143. (Table and plot for **tINV**)

a. Print 10 values for the **tINV** for $\alpha = 0.1 \dots 1$ step 0.1 and $f = 5$.

b. Plot the **tINV** values from a.



View code, run REDUCE.

Exercise 144. (the 99th percentile of the Student's t distribution, cf. `tinvt`)
 Compute the 99th percentiles of the Student's t distribution for 1 to 6 degrees of freedom (df), i.e. verify

```
MATLAB> percentile = tinvt(0.99,1:6)
31.8205  6.9646  4.5407  3.7469  3.3649  3.1427
```



View *code*, run REDUCE.

Exercise 145. (Quantiles of the t distribution, cf. [7, p.74ff])
 Produce a table for the quantiles of t-distribution for $f = 1..20$ and $X = (0.9, 0.95, 0.975, 0.99)$.
 The result should be

quantils of t-distribution for f=1..20				
f	0.9	0.95	0.975	0.99
1	3.07768	6.31375	12.7062	31.8205
2	1.88562	2.91999	4.30265	6.96456
3	1.63774	2.35336	3.18245	4.5407
4	1.53321	2.13185	2.77645	3.74695
5	1.47588	2.01505	2.57058	3.36493
6	1.43976	1.94318	2.44691	3.14267
7	1.41492	1.89458	2.36462	2.99795
8	1.39682	1.85955	2.306	2.89646
...				

Compare the run times for the calculation of these 80 values w.r.t. the *bisection method* in exercise.140a, the two *Newton iterations* in exercise.140b and the *num_solve version* of exercise.139.



View *code*, run REDUCE.

3.4 SNEDECOR's F distribution

The confidence interval for the ratio of two variances requires the use of the probability distribution known as the F-distribution.

In particular, this distribution arises from ratios of sums of squares when sampling from a normal distribution, and is important in estimation and in hypothesis testing in the two-sample normal model. The distribution function and its quantile function do not have simple, closed-form representations. ▷ [g4g]

Definition Notation: $X \sim F(d_1, d_2)$

- The probability density function of the F -distribution is defined as

$$\text{fPDF}(\mathbf{x}, \mathbf{d}_1, \mathbf{d}_2) := \frac{\sqrt{\frac{(d_1 x)^{d_1} \cdot d_2^{d_2}}{(d_1 x + d_2)^{d_1 + d_2}}}}{x \cdot \text{B}\left(\frac{d_1}{2}, \frac{d_2}{2}\right)}$$

where d_1, d_2 are the degrees of freedom, x is a probability value and B is the **beta** function.

- The *cumulative* probability of the F -distribution can be expressed as

$$\text{fCDF}(\mathbf{x}, \mathbf{d}_1, \mathbf{d}_2) := I_{\frac{d_1 x}{d_1 x + d_2}}\left(\frac{d_1}{2}, \frac{d_2}{2}\right)$$

where I_x is the regularized *incomplete* beta function.

- The *quantile* function (i.e. inverse cumulative the F -distribution) is approximately

$$\text{fINV}(\alpha, f_1, f_2) \approx \text{inverse}\left(\Phi\left(\frac{x^{1/3} \cdot \left(1 - \frac{2}{9f_2}\right) - \left(1 - \frac{2}{9f_1}\right)}{\sqrt{\frac{2}{9f_1} + x^{2/3} \cdot \frac{2}{9f_2}}}\right)\right)$$

i.e. we have to invert the $y = \Phi(x)$ expression of the standard normal distribution to $x = \Phi^{-1}(y)$, so that we can find the *smallest* k such that $\alpha \leq \Pr(X \leq k)$ for given α . We do not use this approximation formula by E. PAULSON, [7, p. 72], for the implementation, but try it in an exercise.

Instead we implement SNEDECOR's F inverse function using the F CDF via the defining relation $F := \Pr(X \leq k)$ and $x = F^{-1}(p, d_1, d_2)$ with $F(x, d_1, d_2) = p$ and search methods.

Graphical representation




Figure 24: SNEDECOR's F distribution with different degrees of freedom d_i .
Left figure: λ = plot of $\text{FPDF}(x, d_1, d_2)$.
 Check $\text{FPDF}(0, 1) = 0.3$ on the graph.
Right figure: λ = plot of $\text{FCDF}(x, d_1, d_2)$.
 Check $\text{FCDF}(1, 5, 2) \approx 0.4$ on the graph.
 Check $\text{FINV}(0.6, 100, 100) \approx 1$ on the CDF via $y = 0.6 \rightarrow \downarrow 1 = x$

Examples Verified with REDUCE in exercises N° 146, 149, 151.

1. $fPDF$: Calculate the density of the F distribution with $df1 = 10$, $df2 = 20$ at the value $x = 1.2$.
 λ | `fPDF(1.2, 10, 20)` | 0.5626
2. $fCDF$: Calculate the area under the F -curve for the interval $[0, 1.5]$ with $d1 = 10$, $d2 = 20$.
 λ | `fCDF(1.5, 10, 20)` | 0.7890
3. $fINV$: Let X be an F random variable with 4 numerator degrees of freedom and 5 denominator degrees of freedom. What is the upper 5th percentile?
 λ | `fINV(0.95, 4, 5)` | 5.1921

General information

- General mathematical information about the concept is here $\triangleright$ wiki : *F distribution*
- Syntax and semantic of the function is here $\triangleright$ MATLAB : `fpdf`
- Online calculator *Bognar's app*

3.4.1 Exercises

Exercise 146. (user-defined function **fPDF** in REDUCE)

Write function **fPDF()** for the **f** distribution using the mathematical definition for **fPDF**; it needs the special function $Beta(a, b)$ or $\Gamma(x)$.

There is no simple term. Check the function on example §3.4-1.



View *code*, run REDUCE.

Exercise 147. (Table and plot for **fPDF**)

a. Calculate the 9 values of the **fPDF** plot ('red') shown in fig.24.left.

b. Plot the **fPDF** values from a.



View *code*, run REDUCE.

Exercise 148. (examples from MATLAB, cf. **fpdf**)

Verify the example from MATLAB:

```
MatLAB> y = fpdf(1:6,2,2)
y = 0.2500  0.1111  0.0625  0.0400  0.0278  0.0204
```



View *code*, run REDUCE.

Exercise 149. (user-defined function **fCDF** in REDUCE)

Write function **fCDF()** for the **f** distribution using the mathematical definition for **fCDF**.

It needs the special function $iBeta(a, b, x)$, which is autoloaded from package **specfn**.

There is no simple term. Check your function on example §3.5-2.



View *code*, run REDUCE.

Exercise 150. (table and plot of the **fCDF**)

a. Calculate the 9 values of the **fCDF** shown in fig.24.left.

b. Plot the **fCDF** values from a.



View *code*, run REDUCE.

Exercise 151. (user-defined function **fINV** in REDUCE)

The implementing of inverse distribution functions is notoriously difficult, so let artificial intelligence - such as Google AI, Claude, Google Gemini, or ChatGPT - help you. Therefore we're conducting a small case study and show several variations regarding the definition of **fINV**. Think about the use of **iBeta(a,b,x)**.

Calculate the function value **fINV(p, df1, df2)** for the **f** distribution

a. using step-by-step commands



View *code*, run REDUCE

b. defining a function **fINV(p, df1, df2)**.



View *code*, run REDUCE

Check your construction on example §3.4-3.

Remark.

- This construction finds the n .th quantile, that is: if $P(x \leq k)$ is given, it finds k .
- `x = fINV(alpha, lambda)` returns the smallest value x such that the `fCDF` evaluated at x equals or exceeds p .

Exercise 152. (Table and plot for `fINV`)

- Print 10 values for the `fINV` for $\alpha = 0.1 \dots 1$ step 0.1 and $df_1 = 5$ and $df_2 = 2$.
- Plot the `fINV` values from a.

View *code*, run REDUCE.**Exercise 153.** (Table from MatLAB)

Reproduce the following table with REDUCE:

```
octave:1> pkg load statistics
octave:2> X      = [0.25, 0.5, 0.75, 0.999]
octave:3> finv(X,10,20)
ans = 0.6564    0.9663    1.3995    5.0752
```

View *code*, run REDUCE.**Exercise 154.** (NEWTON-RAPHSON iteration method for `fINV`)

- AI info:

To compute the **Inverse F-Distribution** (critical values for a given cumulative probability p and degrees of freedom d_1, d_2) using **Newton's method** in REDUCE 3.8, we utilize the cumulative distribution function (CDF) expressed via the normalized (regularized) incomplete beta function, which is natively available through REDUCE's `specfn` package as `iBeta`. SourceForge

The Newton iteration updates the estimate x according to:

$$x_{n+1} = x_n - \frac{F(x_n; d_1, d_2) - p}{f(x_n; d_1, d_2)}$$

where:

- $F(x)$ is the F-distribution CDF given by $iBeta\left(\frac{d_1}{2}, \frac{d_2}{2}, \frac{d_1 x}{d_1 x + d_2}\right)$.
- $f(x)$ is the F-distribution probability density function (PDF), se as the derivative of the CDF.

+ Frag Gemini

Um die **inverse F-Verteilung** ($F^{-1}(p, d_1, d_2)$) mittels **Newton-Verfahren** (Newton-Raphson-Iteration) im Computeralgebrasystem **REDUCE 3.8** zu berechnen, löst man das Nullstellenproblem

$g(x) = I_x\left(\frac{d_1}{2}, \frac{d_2}{2}\right) - p = 0$, wobei I_x die regulierte unvollständige Beta-Funktion ist. blog.pkh.me

Da die Ableitung der kumulierten Verteilungsfunktion (CDF) genau der Wahrscheinlichkeitsdichte (PDF) entspricht, lautet die Iterationsvorschrift für das Newton-Verfahren:

$$x_{n+1} = x_n - \frac{\text{CDF}(x_n) - p}{\text{PDF}(x_n)}$$

Figure 25: Some info about SNEDECOR's F distribution from Claude AI.

To calculate the inverse F-distribution (quantile function) `fINV` using Newton's method in REDUCE, one numerically solves the equation

$$I_x(0.5 \cdot d_1, 0.5 \cdot d_2) - p = 0$$

where I_x represents the regularized incomplete beta function (the cumulative F-distribution). Since the derivative of the cumulative distribution function `fCDF` corresponds exactly to the probability density function `fPDF`, the Newton iteration formula is:

$$x_{n+1} = x_n - \frac{F_{\text{cdf}}(x_n) - p}{F_{\text{pdf}}(x_n)}$$

- a. Do an AI question e.g. 'Inverse F Distribution using Newton iteration in CAS Reduce' given to *Google AI*, *Claude*, *Gemini* or *ChatGPT* and try to get the presented code running.



View *code*, run REDUCE

- b. Test your function `fINV_Newton` w.r.t. the table of exercise.153.
- c. **Attention:** $df1 > 6$ and $df2 > 4$ should be fulfilled!

Check: `fINV_Newton(0.95,8,1)` w.r.t. `R> qf(0.95,8,1)` [1] 238.8827

Check: `fINV_Newton(0.6,2,9)` w.r.t. `R> qf(0.6,2,9)` [1] 1.016246

w.r.t. `octave:1> finv(0.6, 2, 9)` ans = 1.0162



View *code*, run REDUCE.

Exercise 155. (an approximation for `fINV` after PAULSON)

Because there is no simple term for the inverse `f` distribution `fINV()` implement the approximation after PAULSON, cited in the definition above on p.82.

Write `fINV_Paulson()` in REDUCE and test it w.r.t. exercise.153.



View *code*, run REDUCE.

3.5 Chi-Square distribution

The χ^2 -distribution is the distribution of a sum of the squares of independent standard normal random variables. The chi-squared distribution is used in the common chi-squared tests for goodness of fit of an observed distribution to a theoretical one, the independence of two criteria of classification of qualitative data, and in finding the confidence interval for estimating the population standard deviation of a normal distribution from a sample standard deviation. ▷ [WIKIPEDIA]

Definition Notation: $X \sim \chi^2(k)$

- The probability density function of the *Chi-squared distribution* χ^2 is defined by

$$\text{chiPDF}(\mathbf{x}, \mathbf{k}) := \Pr(X = k) = \begin{cases} \frac{x^{k/2-1} e^{-x/2}}{2^{k/2} \Gamma\left(\frac{k}{2}\right)} & : x > 0 \\ 0 & : \text{otherwise.} \end{cases}$$

where k is the number of degrees of freedom, and Γ is the Gamma function.

- The *cumulative* probability of the χ^2 -distribution can be expressed by

$$\text{chiCDF}(\mathbf{x}, \mathbf{k}) := \Pr(X \leq k) = P\left(\frac{k}{2}, \frac{x}{2}\right)$$

where $P(s, t)$ is the *regularized incomplete gamma* function.

- The *quantile* function (inverse cumulative of χ^2 -distribution) is approximately

$$\text{chiINV}(\alpha, k) \approx k \cdot \left(1 - \frac{2}{9k} + u_\alpha \cdot \sqrt{\frac{2}{9k}}\right)^3$$

i.e. it finds the *smallest* k such that $\alpha \leq \Pr(X \leq k)$ for given α , where u_α is the α -quantile of the standard normal distribution.

We use this approximation formula by WILSON & HILFERTY, [7, p. 73] in an exercise.

Instead we implement the χ inverse function using the χ CDF via the defining relation $F := \Pr(X \leq k)$ and $x = F^{-1}(p, k)$ with $F(x, k) = p$ and a classic search method..

Graphical representation




Figure 26: χ^2 -distribution with different degrees of freedom k . ▷wiki.
Left figure: λ = plot of `chiPDF(x, k)`, $k = 1; 2; 3; 4; 6; 9$.
 Check `chiPDF(2, 4) = 0.18` on the graph.
Right figure: λ = plot of `chiCDF(x, ν)`, $\nu = 1; 2; \dots$.
 Check `chiCDF(4, 4) ≈ 0.6` on the graph.
 Check `chiINV(0.6, 4) ≈ 4` on the CDF via $y = 0.6 \rightarrow \downarrow 4 = x$

Examples Verified with REDUCE in exercises N° 156, 158, 161.

We use the notational convention of MATLAB.

1. *chiPDF*: Let X follow a χ^2 -distribution with 3 *df*. What is the density if the value of X is 2?
 | `chi2PDF(2, 3)` | 0.2075
2. *chiCDF*: Let X be a chi-square random variable with 3 degrees of freedom.. Compute the probability $P(0.35 \leq X \leq 7.81)$.
 | `chi2CDF(7.81, 3) - chi2CDF(0.35, 3)` | 0.9002
3. *chiINV*: Find the 95th percentile for the chi-square distribution with 10 degrees of freedom.
 | `chi2INV(0.95, 10)` | 18.307

General information

- General mathematical information about the concept is here ▷ WIKI : Chi squared
- Syntax and semantic of the function is here ▷ MATLAB : `chi2`
- Online calculator *Bognar's app*

3.5.1 Exercises

Exercise 156. (user-defined function `chi2PDF` in REDUCE)

Write function `chi2PDF()` for the density of χ^2 using the mathematical definition.

You may use the special function $\Gamma(x) \stackrel{\text{REDUCE}}{:=} \text{Gamma}(x)$.

Check your function on example §3.5-1.



View *code*, run REDUCE.

Exercise 157. (Table and plot for `chi2PDF`)

a. Do the following calculation in REDUCE:

```
"      x      chi2PDF(x, 6)      Chi-squared table for df=6"
T = zero(2,9)
for(i,1,9, T[1,i] = 0.1*i)
for(i,1,9, T[2,i] = chi2PDF(0.1*i ,6))
transpose(T)
```

b. Calculate the values of the `chi2PDF` plot ('red') shown in fig.26.left.

c. Plot the `chi2PDF` values from a.



View *code*, run REDUCE.

Exercise 158. (user-defined function `chi2CDF` in REDUCE)

Write function `chi2CDF` for the χ^2 distribution using the mathematical definition.

You need the special binary function *normalised incomplete gamma* $P(a, x) := \text{GammaP}(a, x)$, named `iGamma(a, x)` in REDUCE. There is no simple term.

Check your function on example.2.



View *code*, run REDUCE.

Exercise 159. (table and plot of the `chi2CDF`)

a. Calculate the 9 values of the `chi2CDF`($x, 4$), $x = 0..8$ *step* 1, shown in fig.26.right.


b. Plot the `chi2CDF` values from a.

c. Check with REDUCE:²⁴ `octave> gammainc(3,2) ans = 0.8009`



View *code*, run REDUCE.

d. Plot the normalised incomplete gamma function 'GammaP' for $k = 3$, cf. `gammainc`



²⁴`gammainc` is the MatLab notation for `GammaP` alias `iGamma`.

Exercise 160. (χ^2 table for $df = 3$)

Do the following calculation in REDUCE:

```
Values for chi^2 with 3 degrees of freedom, i.e. df=3
x   Chi^2(x)   P value (df=3)
x   Chi^2cdf   P-value
1   0.198748   0.801252
2   0.427593   0.572407
4   0.738536   0.261464
8   0.953988   0.046012
16  0.998866   0.001134
32  0.999999   0.000001
```



View *code*, run REDUCE.

Exercise 161. (user-defined function χ^2 INV in REDUCE)

Think of the use of `iGamma(a,b)` of the REDUCE package `specfn` (autoloaded).

Calculate the function value `chi2INV(p,df)` for the χ^2 -distribution

a. using step-by-step commands



View *code*, run REDUCE

b. defining a function `chi2INV(p,df)`.



View *code*, run REDUCE

Check your construction on example §3.5-3.

Exercise 162. (bisection search method for `chi2INV`)

To implement the χ^2 inverse function using `chi2CDF` and a stable bisection search method, write `chiINV_bisection()` in REDUCE and test it w.r.t.

```
MATLAB> chi2inv(0.95,10)      ans = 18.3070
R> qchisq(x = 0.95, df = 10) [1] 18.30704
```



View *code*, run REDUCE.

Exercise 163. (χ^2 quantiles for $k = 10$)

Let $P = (0.9, 0.95, 0.975, 0.99)$. – Calculate `chi2INV(x, 10)` for p in P .

Result: (15.987, 18.307, 20.483, 23.209).



View *code*, run REDUCE.

Exercise 164. (Table from MatLAB)

Reproduce the following table with REDUCE:

```
octave> chi2inv(0.1:0.1:1.0, 10)
ans = 4.8652 6.1791 7.2672 8.2955 9.3418 10.4732 11.7807 13.4420
      15.9872 Inf (is our right limit b=20 ;)
```



View *code*, run REDUCE.

Exercise 165. (NEWTON-RAPHSON iteration method for `chi2INV`)

- AI info:

1. Mathematical Background

The cumulative distribution function (CDF) of the chi-squared distribution for a random variable x with k degrees of freedom is defined as:

$$F(x; k) = \text{iGamma}\left(\frac{k}{2}, \frac{x}{2}\right)$$

To find the inverse chi-squared distribution (i.e., finding x for a given probability p and degrees of freedom k), you must solve the equation:

$$\text{iGamma}\left(\frac{k}{2}, \frac{x}{2}\right) = p$$

This means $x = 2 \cdot \text{iGamma}^{-1}\left(\frac{k}{2}, p\right)$. Because REDUCE does not provide a built-in inverse for `iGamma`, you can implement a numerical root-finding approach (such as bisection or using REDUCE's numeric solver) combined with `iGamma`.

Newton Iteration

The cumulative distribution function (CDF) of the chi-squared distribution for a random variable x with k degrees of freedom is given by the normalized lower incomplete gamma function:

$$F(x; k) = \text{iGamma}\left(\frac{k}{2}, \frac{x}{2}\right) = p$$

To find x for a given probability p and degrees of freedom k , we define our objective function to root-find:

$$f(x) = \text{iGamma}\left(\frac{k}{2}, \frac{x}{2}\right) - p = 0$$

The derivative of $f(x)$ with respect to x is derived using the chain rule and the definition of the incomplete gamma function:

Figure 27: Left: concept for the use of `num_solve` for `chi2INV`, see a. and b.
Right: with $f'(x) = \frac{d}{dx}[\text{iGamma}(k/2, x/2)] = \frac{1}{2} \cdot \left(\frac{x}{2}\right)^{k/2-1} \cdot e^{-x/2} \cdot \Gamma(k/2)$

To calculate the inverse χ^2 -distribution (quantile function) $(\chi^2)^{-1}$ use Newton's method in REDUCE with special function `iGamma`, cf. exercise 161.

- Do an AI question e.g. 'Inverse chi squared Distribution using Newton iteration in CAS Reduce' given to Google AI, Claude, Gemini or ChatGPT, study the presented code and try to get it running.  [View code](#), run REDUCE
- Test your function `chi2INV_Newton` w.r.t. the tables of exercise 163 and 164.
- Check: `chi2INV_Newton(0.95,10)` w.r.t. `R> qchisq(0.95, df = 10)` [1] 18.30704
Check: `chi2INV_Newton(0.6,10)` w.r.t. `R> qf(0.6,10)` [1] 10.47324



[View code](#), run REDUCE.

Exercise 166. (approximation for `chi2INV` after WILSON & HILFERTY)

Because there is no simple term for the inverse $(\chi^2)^{-1}$ distribution `chi2INV()`, implement the approximation after WILSON & HILFERTY, cf. [7, p.73], cited in the definition above. Write `chi2INV_WH()` in REDUCE and test it w.r.t. `R> qchisq(0.6, df=10)` [1] 10.47324.



[View code](#), run REDUCE.

3.6 PARETO distribution

▷ [g4g:] The PARETO distribution is a power-law probability distribution that is used in description of social, quality control, scientific, geophysical, actuarial, and many other types of observable phenomena; the principle originally applied to describing the distribution of wealth in a society, fitting the trend that a large portion of wealth is held by a small fraction of the population. The PARETO principle or “80:20 rule” stating that 80% of outcomes are due to 20% of causes was named in honour of PARETO.

▷ [MATLAB] Fitting a parametric distribution to data sometimes results in a model that agrees well with the data in high density regions, but poorly in areas of low density. For unimodal distributions, such as the normal or Student’s t, these low density regions are known as the “tails” of the distribution. One reason why a model might fit poorly in the tails is that by definition, there are fewer data in the tails on which to base a choice of model, and so models are often chosen based on their ability to fit data near the mode.

If X is a random variable with a Pareto distribution, then the probability that X is greater than some number x , i.e., the ‘survival’ function (also called ‘tail’ function) is given by $\frac{\alpha \cdot x_m^\alpha}{x^{\alpha+1}}$. Here x_m is the (positive) minimum possible value of X , and α is a positive parameter. The Pareto distribution is characterized by this *scale* parameter x_m and this *shape* parameter α , which is known as the ‘tail index’.

Definition Notation: $X \sim \text{Pareto}(\alpha, x_m)$

- The probability density function of the PARETO-distribution is defined as

$$\text{paretoPDF}(x, \mu, \alpha) := \Pr(X = k) = \begin{cases} \frac{\alpha \cdot \mu^\alpha}{x^{\alpha+1}} & : x \geq \mu \\ 0 & : x < x_m \end{cases}$$

where μ is the *scale* parameter and the *shape* parameter is α .

- The *cumulative* probability of a PARETO distribution²⁵ with parameters α and x_m is

$$\text{paretoCDF}(\mathbf{x}, \mu, \alpha) := \Pr(X \leq k) = \begin{cases} 1 - \left(\frac{\mu}{x}\right)^\alpha & : x \geq x_m \\ 0 & : x < x_m \end{cases}$$

- The *quantile* function (inverse cumulative PARETO -distribution) is²⁶

$$\text{paretoINV}(p, \mu, \alpha) := \mu \cdot (1 - p)^{-\frac{1}{\alpha}} =: x_p$$

where p is a probability value and x_p is the p ’s quantile.

²⁵The CDF function for the Pareto distribution returns the probability that an observation from a Pareto distribution with the shape parameter α and the scale parameter x_m , is less than or equal to x .

²⁶This formula takes a probability p with $(0 < p < 1)$ and returns the value of x at which the CDF is equal to p .

Graphical representation




Figure 28: PARETO distributions with different shape parameters α . ▷wiki.
Left figure: λ = plot of `paretoPDF(x,  $\alpha$ )`, $\alpha = 1; 2; 3$.
 Check `paretoPDF(1, 1) = 1` on the graph.
Right figure: λ = plot of `paretoCDF(x,  $\alpha$ )`, $\alpha = 1; 2; 3$.
 Check `paretoCDF(2, 1)  $\approx$  0.5` on the graph.
 Check `paretoINV(0.6, 1)  $\approx$  3` on the CDF via $y = 0.6 \longrightarrow \downarrow 3 = x$

Examples Verified with REDUCE in exercises N° 167, 169, 171.

In the case where the shape parameter is $\alpha = \log 45 = 1.160964$, we get the famous Pareto principle, aka the 80-20 rule, which states that 80% of the outcomes are due to 20% of the causes.

E.g. 20% of the workers do 80% of the work. 80% of the wealth is owned by 20% of the people.

▷ We adopt the $\mathcal{R}$ convention (x, μ, α) , whereas `pareto`_{distrib(Maxima)} use (x, α, μ) !

1. *paretoPDF*: If X follows a Pareto distribution with shape parameter $\alpha = 5$ and scale parameter $\mu = 2$, find the probability density of the distribution at $x = 3$.

| `paretoPDF(3, 2, 5)` | 0.2194

2. *paretoCDF*: Suppose X follows a Pareto distribution with shape parameter $\alpha = 2.5$ and scale parameter $\mu = 1$. What is the probability that $X \geq 5$?

| `1 - paretoCDF(5, 1, 2.5)` | 0.0178

3. *paretoINV*: Calculate the 25'th percentile of a Pareto distribution with parameters location=1 and shape=2.

| `paretoINV(0.25, 1, 2)` | 1.1547

General information

- General mathematical information about the concept is here ▷ WIKI : Pareto
- Syntax and semantic of the function is here ▷ MATLAB : `gppdf`

- Online calculator *Bognar's app*

3.6.1 Exercises

Exercise 167. (user-defined function `paretoPDF` in REDUCE)

Write function `paretoPDF` for the *Pareto* density using the mathematical definition for *paretoPDF* in REDUCE. Check the function on example §3.6-1.

 View code, run REDUCE.

Exercise 168. (Table and plot for `paretoPDF`)

- Calculate 5 values of the `paretoPDF` plot ('red') shown in fig.28.left.
- Plot the `paretoPDF` values from a.
- MATLAB: Compute the density of the observed value 3 in the Pareto distributions with scale parameter 2 and shape parameters 1 through 5.

 View code, run REDUCE.

Exercise 169. (user-defined function `paretoCDF` in REDUCE)

Write function `paretoCDF` for the *cumulative pareto* distribution using the mathematical definition. - Check your function on example §3.6-2.

 View code, run REDUCE.

Exercise 170. (table and plot of the `paretoCDF`)

- Check the two values of the `paretoCDF` shown in fig.28.right.
- Plot the `paretoCDF` shown in fig.28.right.
- Calculate the cdf's of a Pareto distribution with parameters location=2 and shape=1, evaluated at 3, 4, and 5.

 View code, run REDUCE.

Exercise 171. (user-defined function `paretoINV` in REDUCE)

Write function `paretoINV` for the inverse *pareto* distribution using the mathematical definition in REDUCE. - Check your function on example §3.6-3.

 View code, run REDUCE.

Remark.

- `paretoINV()` solves the task: find x , given $y = p$, shape m , scale a , s.t. $\text{paretoINV}(x, m, a) = p = y$.
- The result is the value such that an observation from an Pareto distribution with parameters a, m will fall in the range $[0, x]$ with probability p .

Exercise 172. (Table and special value for `paretoINV`)

- Print 10 values for the `paretoINV` for $p = 0.1 \dots 1$ step 0.1 and $\mu = 2.5$ for $\alpha = 2$.
- MATLAB: Determine the 25'th percentile of a Pareto distribution with parameters location=1 and shape=1.

 View code, run REDUCE.

3.7 WEIBULL distribution

The WEIBULL distribution is a continuous probability distribution. It models a broad range of random variables, largely in the nature of a time to failure or time between events. Examples are maximum one-day rainfalls and the time a user spends on a web page. ▷ [WIKIPEDIA]

Definition Notation: $X \sim Weibull(\lambda, k)$

- The probability density function of a WEIBULL random variable is

$$\text{weibullPDF}(\mathbf{x}, \lambda, \mathbf{k}) := \Pr(X = k) = \begin{cases} \frac{k}{\lambda} \left(\frac{x}{\lambda}\right)^{k-1} e^{-(x/\lambda)^k} & : x \geq 0, \\ 0 & : x < 0, \end{cases}$$

where $k > 0$ is the *shape* parameter and $\lambda > 0$ is the *scale* parameter.

- The cumulative WEIBULL distribution is

$$\text{weibullCDF}(\mathbf{x}, \lambda, \mathbf{k}) := F(p, \lambda) = \Pr(X \leq k) = \begin{cases} 1 - e^{-(x/\lambda)^k}, & x \geq 0, \\ 0, & x < 0. \end{cases}$$

- The quantile function (inverse cumulative WEIBULL distribution) is

$$\text{weibullINV}(\mathbf{p}, \lambda, \mathbf{k}) := F^{-1}(p, \lambda) = \lambda(-\ln(1 - p))^{\frac{1}{k}}$$

Examples Verified with REDUCE in exercises N° 173, 176, 178.

▷ We adopt the $\mathcal{R}$ convention (x, λ, k) , whereas `weibull`_{*distrib(Maxima)*} use (x, k, λ) !

1. (MatLAB) Compute the density of the observed value 3 in the Weibull distribution with unit scale and shape.

```
| weibullPDF(3,1,1) | 0.0497
```

2. (MatLAB) What is the probability that a value from a Weibull distribution with parameters $a = 0.15$ and $b = 0.8$ is less than 0.5?

```
| weibullCDF(0.5, 0.15, 0.8) | 0.9272
```

3. (MatLAB) The lifetimes (in hours) of a batch of light bulbs has a Weibull distribution with parameters $a = 200$ and $b = 6$. Find the median lifetime of the bulbs.

```
| weibullINV(0.5, 200, 6) | 188.15
```

General information

- General mathematical information about the concept is here ▷ WIKI : Weibull
- Syntax and semantic of the function is here ▷ MATLAB : `wblpdf`
- Online calculator *Bognar's app*

Graphical representation




Figure 29: WEIBULL distributions with different parameters λ, k .
Left figure: $\lambda = 1, k = 0.5; \dots; 5$.
 Check `weibullPDF(1, 1, 5) = 1.8` on the graph.
Right figure: $\lambda = 1, k = 0.5; \dots; 5$.
 Check `weibullCDF(1, 1, 5)  $\approx$  0.6` on the graph.
 Check `weibullINV(0.8, 1, 5)  $\approx$  1.2` on the CDF via $y = 0.8 \rightarrow \downarrow 1.2 = x$

3.7.1 Exercises

Exercise 173. (user-defined function `weibullPDF` in REDUCE)

Write function `weibullPDF()` for the density of the WEIBULL distribution using the mathematical definition for `weibullPDF`. - Check your function on example §3.7-1.



View `code`, run REDUCE.

Exercise 174. (Table and plot for `weibullPDF`)

- Calculate the 9 values of the `weibullPDF` plot ('green') shown in fig.29.left.
- Plot the `weibullPDF` values from a.



View `code`, run REDUCE.

Exercise 175. (example from MATLAB, cf. `wblpdf`)

Compute the density of the observed value 3 in the Weibull distributions with scale parameter 2 and shape parameters 1 through 5.



View `code`, run REDUCE.

Exercise 176. (user-defined function `weibullCDF` in REDUCE)

Write function `weibullCDF` for the cumulative Weibull distribution using the mathematical definition for `weibullCDF`. - Check your function on example §3.7-2.



View `code`, run REDUCE.

Exercise 177. (table and plot of the `weibullCDF`)

a. Calculate the 9 values of the `weibullCDF` ‘green’ shown in fig.29.right.

b. Plot the `weibullCDF` values from a.

 View *code*, run REDUCE.

Exercise 178. (user-defined function `weibullINV`)

Write function `weibullINV` for the Weibull-CDF inverse using the mathematical definition for `weibullINV`. - Check your function on example §3.7-3.

 View *code*, run REDUCE.

Exercise 179. (Table and plot for `weibullINV`)

a. Print 10 values for the `weibullINV` for $p = 0.1 \dots 1$ step 0.1, $\lambda = 1$ and $k = 5$.


Check `weibullINV(0.8,1,5)` on the graph of fig.29.b.

 View *code*, run REDUCE.

b. Plot the `weibullINV` values from a.

c. Reproduce the following plot with REDUCE.

Quantils for Weibull(.,2,2)



Exercise 180. (1st optional parameterization used by WIKIPEDIA: `weibull1`)

Wikipedia use the transformation $b := \lambda^{-k}$, e.g. [6], p.12, 64 wiki:weibull.

a. Write the three weibull-distribution functions `weibull1XXX` w.r.t. Wikipedia.

b. Solve: A particular switch has a Weibull distributed lifetime with $k := \alpha = 0.01$ (1/year) and $\beta = 2$. What guarantee period can be given for the switch if the survival probability during this time is to be 99%?

 View *code*, run REDUCE.

Exercise 181. (2nd optional parameterization: `weibull2`)

In wiki2 the shape parameter k is the same as in the standard case, while the scale parameter λ is replaced with a rate parameter $\beta = 1/\lambda$.

a. Write the three weibull-distribution functions `weibull12XXX` w.r.t. Wikipedia.

b. Calculate $1 - \text{weibullCDF2}(1, 0.01, 2)$

 View *code*, run REDUCE.

Exercise 182. (example from BEUCHER)

Translate the following MATLAB/OCTAVE snippet to REDUCE, cf. [1].

Watch out: $(a, b) = (b, T)$:

```
x=(0:1:10);           % values for x:  0 until 10 step 1
b = 0.444;  T=0.767;   % parameter of Weibull-distribution
a = 1/(T^b);           % adaption of parameters to MATLAB
                        % calculate the Weibull-density

pdfwb = wblpdf(x, a, b)
pdfwb =
Columns 1 through 6:
    Inf  1.6312e-01  7.8813e-02  4.8763e-02  3.3664e-02  2.4765e-02
Columns 7 through 11:
    1.9002e-02  1.5028e-02  1.2160e-02  1.0018e-02  8.3746e-03
```



View *code*, run REDUCE.

4 Test Statistics


We now show, how the distributions of chapter 3 are used to implement diverse parameter tests. We translate the mathematical concepts and definitions in REDUCE code and demonstrate exemplary calls and examples.

A: Parameter Tests We do all tests first *semi-automatic* in REDUCE, so that the user can follow the essential steps and can do the steps also by hand. This *procedural* way is followed by a *fully automatic solution* using plain REDUCE.

4.1 One Sample Z-Test alias GAUSS test

The z-test is a parametric hypothesis test used to determine whether a sample data set comes from a normal distributed population with a particular mean and a known standard deviation. The one sample GAUSS-test tests the sample for a certain mean value μ_0 . For this, the parameters of the normal distribution must *not* be estimated from the sample.

Mental image



Visualization of Z-test for a Hypothesized Mean μ_0 :

Figure 30: • : sample X with sample mean $\bar{X}$.

• : population with population mean μ .

Procedure *One sample Z-test alias GAUSS -test*

1. **Assumptions** The sample $X = (x_1, \dots, x_n)$ must be $\mathcal{N}(\mu, \sigma^2)$ distributed.
2. **Null hypothesis** $H_0 : \mu = \mu_0$ (two-sided)
3. **Test statistics** Calculate the normal distributed test value

$$T := \frac{\sqrt{n}}{\sigma} \cdot (\bar{X} - \mu_0).$$

4. **Decision** Reject H_0 , if $|T| > u_{1-\alpha/2}$.
 u_α is the α -quantile of the standard normal distribution, i.e. `stdnormalINV`(α).

Remark. one-sided test $\mu > \mu_0$: $T > u_{1-\alpha}$ resp. one-sided test $\mu \leq \mu_0$: $T < u_\alpha$

Example (*female blood pressure*)

The female blood pressure of a certain population is known to follow Gaussian (alias: normal) distribution with mean 124.6 and standard deviation 14.5 measured in units of mmHg. In order to test the effect of a food product on the female blood pressure, a clinical trial was performed in which 12 female volunteers of this population consumed the product for 3 months and their blood pressures were measured at the end. The readings are as follows:

	Sample of female blood pressure											
$n :$	1	2	3	4	5	6	7	8	9	10	11	12
mmHg	141.5	152.3	121.2	123.0	151.6	124.8	138.9	137.4	145.6	135.6	135.4	121.5

Let $\alpha = 0.05$ be the probability of rejecting the null hypothesis.

Can we conclude from this data, that the population mean of the data set from which these random observations are drawn is not equal to (ie., different from) 124.6?

Solution:

```
| X:={141.5, 152.3, 121.2,123.0,151.6,124.8,138.9,137.4,145.6,135.6,135.4,121.5};
| %..... mu0      sigma alpha type
| ztest(X, 124.6, 14.5, 0.05, 0);
|
```

Z	quantil	CI.left	CI.right	p
2.6598	1.96	127.53	143.94	0.0078189



View solution step-by-step: `code`, run REDUCE.



View solution via function: `code`, run REDUCE.

General information

General mathematical information about the concept is here [▷ WIKIPEDIA : Z-Test](#)
 Syntax and semantic of the implementation is here [▷ MATLAB : Z-Test](#)

4.1.1 Exercises

Exercise 183. (example at CountBio)

Study this example at [▷ CountBio](#).

Exercise 184. (example at FU Berlin)

Study using REDUCE the example [▷ FU Berlin script.5](#)

Exercise 185. (One Sample z-test in $\mathcal{R}$)

Do using REDUCE the 'Example of One Sample z-test in R' [▷ Tutorial](#).

Exercise 186. (One Sample z-test of a sample with given parameters)

Suppose, a sample of $n = 50$ items has $\bar{x} = 105$, $\mu = 100$ and $\sigma = 15$.

Determine the z-test on the significance level $\alpha = 0.05$.




View `code`, run REDUCE.

4.2 Two Sample Z-Test

Two sample Z-Test for two samples with two means and two known variances is to test the null hypothesis that there is no difference between the means of the two independent samples. The assumptions are: 1. Normal but independent populations. 2. Variances for the populations are known.

Mental image



Visualization of Z-test for a Hypothesized Mean μ_0 :
 Figure 31: $\bullet$: samples X and Y with sample means $\bar{X}$ resp. $\bar{Y}$
 $\bullet$: population with population mean μ .

Procedure *Two sample Z-test*

1. **Assumptions** $X = (x_1, \dots, x_m) \sim \mathcal{N}(\mu_1, \sigma_1^2)$ and $Y = (y_1, \dots, y_n) \sim \mathcal{N}(\mu_2, \sigma_2^2)$ normal distributed. X and Y independent.
2. **Null hypothesis** $H_0 : \mu_1 = \mu_2$ (two-sided)
3. **Test statistics** Calculate the normal distributed test value

$$T := \frac{\bar{X} - \bar{Y}}{\sqrt{\frac{\sigma_1^2}{m} + \frac{\sigma_2^2}{n}}}.$$

4. **Decision** Reject H_0 , if $|T| > u_{1-\alpha/2}$.
 u_α is the α -quantile of the standard normal distribution, i.e. `stdnormalINV`(α).

Remark. one-sided test of $H_0 : \mu_1 > \mu_2 \Rightarrow T > u_{1-\alpha}$
 one-sided test of $H_0 : \mu_1 \leq \mu_2 \Rightarrow T < u_\alpha$

1. The Z-score T (aka the ‘test statistics’) represents the number of standard deviations that the difference between the two sample means is from zero.
2. The ‘critical values’ are based on the standard normal distribution and are used to determine whether the calculated z-score is statistically significant. If the calculated Z-score is greater than the critical value, the null hypothesis is rejected, and the alternative hypothesis is accepted.
3. For a two-tail test with a significance level of $\alpha = 0.05$, the critical value is 1.96. You can find the critical values using the `REDUCE` function `stdnormINV`($\alpha/2$) for the two-tail test.

Example *two samples drawn from a population.*


The figure show two samples X and Y , blindly drawn from a box with balls numbered (30) until (70) (the ‘population’). Check using e.g. a QQ-plot that X and Y are approximately normal distributed.

Calculate (μ_1, σ_1) and (μ_2, σ_2) .

Let $\alpha = 0.05$ be the probability of rejecting the null hypothesis $\mu_1 = \mu_2$.

Calculate the value ‘test statistic’ T and the critical value.

Decide, whether the null hypothesis can be accepted or must be rejected.



Solution:

```
X := 47,37,46,38,35,38,40,37;$
Y := 59,56,61,59,65,56,57,50;
%.... X,Y, muX, muY, sigmaX, sigmaY, alpha, altHyp
ztest2(X,Y, 0, 0, 4.11552, 2.78388, 0.05, 0);
```

Z	Z.score	CI.left	CI.right	p
-10.318	1.96	-21.568	-14.682	0.0



View step-by-step *code*, run REDUCE.



View call of ztest2 *code*, run REDUCE.

General information

General mathematical information about the concept is here [▷ WIKIPEDIA : Z-Test](#)

Syntax and semantic of the implementation is here [▷ MATLAB : Z-Test](#)

4.2.1 Exercises

Exercise 187. Study the example at `> statkat`.

Use `REDUCE` to follow and control the argumentation.

Exercise 188. (*Sperling data* samples from `> openeducator`.)

We have the *Sperling data* samples from US vs. Sweden:

```
US =(69.12,66.88,74.82,67.00,69.12,65.00,71.00,66.76,72.12,72.94,
      69.18,66.18,64.94,71.76,70.12,71.00,71.88,65.24,70.06,71.94,
      72.12,66.88,73.82,74.00,71.18,67.88,65.94,68.88,68.00,75.12)
```

```
Sw =(74.56,71.89,73.00,67.78,72.22,68.00,73.56,75.00,68.22,69.00,
      68.00,72.00,73.56,72.56,75.00,68.33,71.67,72.44,75.00,71.89,
      72.00,70.00,69.22,74.44,68.00,73.89,70.00,70.44,70.22,73.33)
```

with population means 69.98 vs. 70.43 and stdDevs 3.12 vs. 2.44.

a. Do a step-by-step two-sample-z-test analog example.4.2.1

b. Do the two-sample-z-test for US vs. Sw using the function `ztest2()` from example.4.2.2.

Exercise 189. (Two-sampled z-test for two medications `> g4g`.)

A researcher wants to compare the effectiveness of two different medications for reducing blood pressure. Medication A is tested on 50 patients, resulting in a mean reduction of 15 mmHg with a standard deviation of 3 mmHg. Medication B is tested on 60 patients, resulting in a mean reduction of 13 mmHg with a standard deviation of 4 mmHg.

At a 1% significance level, is there a significant difference between the two medications?

a. Do a step-by-step two-sample-z-test analog example.4.2.1.

b. Do the two-sample-z-test for US vs. Sw using the function `ztest2()` from example.4.2.2

.

Exercise 190. (Example from `REAL STATISTICS > real-statistics`.)

Here are the data

```
X=(82.67,90.11,89.20,119.15,83.01,93.61,88.42,97.02,126.11,127.96,
  89.03,94.51,93.32,89.26,110.36,92.52,112.87,64.05,80.06,74.13,
  109.13,81.59,94.99,101.34,104.82,106.92,80.50,106.31,85.46,103.69)
```

```
Y=(106.18,100.86,129.85,100.30,87.56,96.87,112.57,148.36,131.62,
  114.60,95.47,108.66,83.32,117.64,96.90,66.46,87.80,115.52,102.34,
  97.01,117.51,115.64,97.22,131.04,101.58,103.80,111.99,119.34,95.10,
  107.62)
```

Do the two-sample-z-test following the text.

Exercise 191. (Examples from `www.countbio.com`, `countbio`)

Do example 1 and example 2 in `www.countbio.com`.

4.3 One Sample t -Test

The one sample t -test is a parametric hypothesis test used to determine whether a sample data set X comes from a normal distributed population with a particular mean μ_0 . In contrast to the one-sample Z -test, the standard deviation σ of the population is estimated using the sample's standard deviation s .

Mental image




Figure 32: Visualization of one-sample- t -test for a hypothesized mean μ_0 : [▷ wiki](#)

Procedure *One sample t -test*

1. **Assumptions** independent sample $X = (x_1, \dots, x_n)$ is $\mathcal{N}(\mu, \sigma^2)$ distributed.
2. **Null hypothesis** $H_0 : \mu = \mu_0$ (two-sided)
3. **Test statistics** Calculate the t -distributed test value with $n - 1$ degrees of freedom.

$$T := \frac{\sqrt{n}}{s} \cdot (\bar{X} - \mu_0) \quad \text{where } s := sd(X).$$

4. **Decision** Reject H_0 , if $|T| > t_{1-\alpha/2; n-1} = \mathfrak{tINV}(1 - \alpha/2, n - 1)$.
 $t_{\alpha, \nu}$ is the α -quantile of the t -distribution, i.e. $\mathfrak{tINV}(\alpha, \nu)$.

Remark. one-sided test $\mu > \mu_0$: $T > t_{1-\alpha, n-1}$ resp. one-sided test $\mu \leq \mu_0$: $T < t_{\alpha, n-1}$

Example *Students weight in Europe, cf. geo.fu-berlin.*

A students data set consists of 8239 rows, each of them representing a particular student, and 16 columns, each of them corresponding to a variable/feature related to that particular student. We examine the average weight of a random sample of students from the students data set and compare it to the average weight of all European adults. WALPOLE et al. (2012) published data on the average body mass (kg) per region, including Europe. They report the average body mass for the European adult population to be 70.8 kg. We therefore set μ_0 , the population mean, accordingly to $\mu_0 = 70.8$. Further, we take a random sample (X) with a sample size of $n = 9$. The sample consists of the weights in kg of 9 randomly picked students from the students data set.

	Sample of students weight in Europe								
$n :$	1	2	3	4	5	6	7	8	9
kg:	64.4	68.5	64.8	58.9	64.5	68.6	68.7	62.9	73.5

The null hypothesis H_0 states that the average weight of students equals the average weight of European adults as reported by WALPOLE. In other words, there is no difference between the mean weight of students and the mean weight of European adults. Let $\alpha = 0.05$ be the significance level of rejecting the null hypothesis.

- Compute the value of the test statistic T and determine the critical value.
- Conclude, whether the null hypothesis H_0 is rejected or accepted.

Solution:

```
| X := {64.4 , 68.5 , 64.8 , 58.9 , 64.5 , 68.6 , 68.7 , 62.9 , 73.5}$
| mu0 := 70.8;$
| alpha := 0.05$
| ttest(X, mu0, alpha, 0);
```

["both: reject H0"]

(	"Significance level:"	0.05	)
	"Degrees of freedom:"	8	
	"Test statistic:"	-3.3458	
	"Critical value:"	-2.306	



View step-by-step solution [code](#), run REDUCE.



View call of TTEST [code](#), run REDUCE.

General information

General mathematical information about the concept is here [▷ WIKIPEDIA : t-Test](#)
 Syntax and semantic of the implementation is here [▷ MATLAB : t-test](#)

4.3.1 Exercises

Exercise 192. (p-value and confidence interval)
is to calculate for the example “Students weight in Europe”.

Exercise 193. (Example from statstutorial, cf. screws)
A random sample of screws have weights

29.99, 30.01, 29.97, 30.11, 29.99, 30.02.

- a. Calculate a 95% confidence interval for the population’s mean weight.
Assume the population is distributed as $\mathcal{N}(\mu, \sigma^2)$.
- b. Do a step-by-step one-sample-t-test analog example.4.3.1
- c. Do the one-sample-t-test using the function `ttest()` from example.4.3.2.

◦ Result: If we sampled many times, our interval would capture the true mean weight 95% of the time; thus, we are 95% confident that the true mean weight of all screws will fall between 29.96 and 30.07.

Exercise 194. (example from $\mathcal{R}$) Study the example at its source $\mathcal{R}$ vignette.

Exercise 195. (example from [7, p.81, p.96])
This is a normal $\mathcal{N}(100, 3^2)$ -distributed sample of length 100.


X = (96.3, 95.8, 91.7, 91.6, 98.7, 105.0, 99.4, 101.1, 92.9, 97.2,
110.6, 94.9, 101.2, 92.3, 106.1, 96.5, 107.9, 101.5, 99.0, 97.1,
90.8, 101.1, 101.0, 100., 106.9, 101.2, 95.7, 92.6, 101.1, 104.9,
101.1, 104.5, 105.5, 114.1, 91.8, 95.8, 93.5, 98.2, 98.8, 98.5,
108.1, 99.9, 105.1, 95.6, 98.3, 96.5, 94.9, 101.3, 94.7, 103.7,
99.0, 98.5, 94.4, 104.7, 93.2, 97.3, 104.3, 101.3, 96.9, 107.0,
95.0, 101.0, 103.4, 100.5, 100.2, 102.3, 95.1, 95.8, 102.9, 95.0,
102.7, 101.0, 105.5, 97.9, 104.0, 103.3, 97.7, 88.5, 95.5, 100.5,
102.3, 101.0, 100.0, 106.0, 102.8, 106.7, 106.6, 99.2, 112.2, 100.1,
99.3, 100.9, 99.8, 96.0, 97.9, 93.0, 93.7, 97.0, 95.8, 99.9)

Use the std.Dev $\sigma = 4.1$ and check the hypothesis $H_o : \mu = 100$.

4.4 Two Sample t -Test

The *two-sample t -test* tests two normally distributed samples for the same mean value. In contrast to the two-sample Z -test, only the sample variances are used.

Mental image



Visualization of two-sample- t -test for 2 populations: ▷ wiki

Figure 33: Left: the two samples X_1 and X_2 with sample means $\bar{x}_1$ and $\bar{x}_2$.

Right: the two populations with their normal probability distribution.

Procedure *Two Sample t -test*

1. **Assumptions** independent samples $X = (x_1, \dots, x_n)$ is $\mathcal{N}(\mu_1, \sigma_1^2)$ and $Y = (y_1, \dots, y_n)$ is $\mathcal{N}(\mu_2, \sigma_2^2)$ distributed.
2. **Null hypothesis** $H_0 : \mu_1 = \mu_2$ (two-sided)
3. **Test statistics** T is approximately t -distributed

$$T := \frac{\bar{X} - \bar{Y}}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}} \quad \text{where } s_1 := sd(X), n_1 = dim(X) \dots$$

if the number of degrees of freedom f following WELCH is chosen as

$$f := \frac{(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2})^2}{\frac{(\frac{s_1^2}{n_1})^2}{n_1-1} + \frac{(\frac{s_2^2}{n_2})^2}{n_2-1}}$$

4. **Decision** Reject H_0 , if $|T| > t_{1-\alpha/2; f} = \mathbf{tINV}(1 - \alpha/2, f)$.
 $t_{\alpha, \nu}$ is the α -quantile of the t -distribution, i.e. $\mathbf{tINV}(\alpha, \nu)$.

Remark. one-sided test $\mu_1 > \mu_2$: $T > t_{1-\alpha, f}$ resp. one-sided test $\mu_1 \leq \mu_2$: $T < t_{\alpha, f}$

4.4.1 Exercises

Exercise 196. (OCTAVE *t*-test tutorial)

Reproduce the following tutorial in REDUCE.

First generate some toy data

```
N = 20;
d = [randn(N,1) randn(N,1)+0.7];    \% Each column is one dataset

d =
-0.269943    0.727796
 0.234784    0.656962
 0.286618    0.107063
-0.246098   -0.796628
-0.365948    0.844111
-0.830328   -0.988360
 0.555346    1.202737
-0.431598    0.282257
 0.025352    1.049475
 1.304216   -0.617381
-0.511775    0.190412
-0.184626    0.299169
-0.335331    0.386240
-0.149371    0.809719
-0.038682    1.536842
-0.281365   -0.632293
-1.210282    0.250702
-0.557427    1.395904
-0.203508    0.083293
 0.050512    0.206742
```

Then verify the following calculation

```
% Paired-samples t-test
octave:6> [h,p,ci,stats] = ttest(d(:,1), d(:,2));
octave:7> h
  h = 1
octave:9> p
  p = 0.017092
octave:10> ci
  ci = -0.9143
        -0.1011
octave:11> stats
  stats = scalar structure containing the fields:
    tstat = -2.6133
    df = 19
    sd = 0.8688
```


Exercise 197. Follow the example at ‘MATLAB.help’ using REDUCE : ▷ ttest2/MATLAB

Exercise 198. Follow this example using REDUCE : ▷ theopeneducator

4.5 Paired t -Test alias Differences t -Test

The *Differences t -Test* checks two dependent and normally distributed samples for the same mean value. As common with all tests of the t -tests group, it uses the t -distribution to calculate the test statistics, i.e. the *critical* resp. p -value.

Mental image



Visualization of paired- t -test for 2 dependent samples

Figure 34: above: the sample X with sample mean μ_X .

bottom: the sample Y with sample mean μ_Y .

Procedure *Paired alias Differences t -test*

1. **Assumptions** dependent samples $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$, where $d := (x_1 - y_1, \dots, x_n - y_n)$ is $\mathcal{N}(\mu_d, \sigma_d^2)$ distributed.
2. **Null hypothesis** $H_0 : \mu_d = 0$ (two-sided)
3. **Test statistics** T is for H_0 approximately t -distributed

$$T := \frac{\frac{1}{n} \cdot \sum_{i=1}^n d_i}{\frac{1}{n(n-1)} \cdot \sqrt{\sum_{i=1}^n d_i^2 - \frac{1}{n} \cdot (\sum_{i=1}^n d_i)^2}}$$

with $df = n - 1$ degrees of freedom.

4. **Decision** Reject H_0 , if $|T| > t_{\alpha/2; n-1} = \mathbf{tINV}(\alpha/2, n-1)$.
 $t_{\alpha, \nu}$ is the α -quantile of the t -distribution, i.e. $\mathbf{tINV}(\alpha, \nu)$.

Remark. one-sided test $\mu_1 > 0$: $T > t_{1-\alpha; n-1}$ resp. one-sided test $\mu_1 \leq 0$: $T < t_{\alpha; n-1}$

Example *Monthly Rainfall.*

We use the example in [7, p. 99] The monthly rainfall in millimeters over the course of two years is considered. We assume that the differences in rainfall are normally distributed. The null hypothesis is the assumption of the same rainfall with a 5% probability of error.

		<i>Rainfall in mm</i>											
<i>month :</i>	1	2	3	4	5	6	7	8	9	10	11	12	
A:	52.4	37.4	41.2	71.5	51.3	21.6	17.7	21.2	41.3	32.0	53.0	61.4	
B:	47.8	42.0	33.2	41.0	29.5	28.1	17.4	21.5	41.4	51.3	49.5	53.7	

We have $\alpha = 0.05$ as the significance level of rejecting the null hypothesis.

◦ *Solution* with REDUCE :

```
| X := {52.4,37.4,41.2,71.5,51.3,21.6,17.7,21.2,41.3,32.0,53.0,61.4}$
| Y := {47.8,42.0,33.2,41.0,29.5,28.1,17.4,21.5,41.4,51.3,49.5,53.7}$
| tdiff(X,Y,0.05);
| alternative hypothesis: true mean difference is not equal to 0
```

..... accept H0

T: 1.022
p-value: 0.32874



View step-by-step solution *code*, run REDUCE.



View ttest2 in action *code*, run REDUCE.

General information

General mathematical information about the concept is here [▷ WIKIPEDIA : t-Test](#)

Syntax and semantic of the implementation is here [▷ MATLAB : t-test](#)

4.5.1 Exercises

Exercise 199. (Lowering cholesterol levels, cf. `ttest`)

To test a new therapy for lowering cholesterol levels, cholesterol levels are measured in ten subjects before and after treatment. The following measurement results were obtained:

Vor der Behandlung:	223	259	248	220	287	191	229	270	245	201
Nach der Behandlung:	220	244	243	211	299	170	210	276	252	189
Differenz:	3	15	5	9	-12	21	19	-6	-7	12

LEXICON	<i>German</i>	<i>English</i>
	Vor der Behandlung	before treatment
	Nach der Behandlung	after treatment

Do the example 2 ('Beispiel 2') w.r.t. lowering cholesterol levels.

Exercise 200. (example from $\mathcal{R}$)

Check our example alternatively with $\mathcal{R}$.

```
R> X = c(52.4,37.4,41.2,71.5,51.3,21.6,17.7,21.2,41.3,32.0,53.0,61.4)
> Y = c(47.8,42.0,33.2,41.0,29.5,28.1,17.4,21.5,41.4,51.3,49.5,53.7)
> result <- t.test(X, Y, paired = TRUE)
> print(result)
```

Paired t-test


```
data: X and Y
t = 1.022, df = 11, p-value = 0.3287
alternative hypothesis: true mean difference is not equal to 0
95 percent confidence interval:
-4.383924 11.983924
sample estimates:
mean difference
3.8
```

4.6 Chi-Squared Test on Variance

The Chi-Squared Test on Variance is a statistical test used to compare observed results with expected results. The purpose of this test is to determine if a difference between observed data and expected data is due to chance, or if it is due to a relationship between the variables you are studying. ▷ AI

The χ^2 -variance test checks a normally distributed sample for a given variance σ_0^2 .

Mental image



Visualization of Chi-Squared Test on Variance, cf. ▷ FOLEY, enhanced.
Figure 35: LCL = Lower Confidence Level; UCL = Upper CL; chisq = `chi2INV`.
The values shown are from the solution of the example, see below.

Procedure *Chi-Squared Test on Variance*

1. **Assumptions** a normal distributed sample $X = (x_1, \dots, x_n)$ with $\mathcal{N}(\mu, \sigma^2)$.
2. **Null hypothesis** $H_0 : \sigma^2 = \sigma_0^2$ (two-sided)
3. **Test statistics** T is for H_0 χ^2 -distributed with $df = n - 1$ degrees of freedom.

$$T := \frac{(n-1) \cdot s^2}{\sigma_0^2}$$

where s is the sample standard deviation, σ_0 is the hypothesized standard deviation.

4. **Decision** Reject H_0 , if $T > \chi_{1-\alpha/2; n-1}^2 = \text{tINV}(1 - \alpha/2, n - 1)$.
 $\chi_{\alpha, \nu}^2$ is the α -quantile of the χ^2 -distribution, i.e. $\text{tINV}(\alpha, \nu)$.

Remark. one-sided test $\sigma^2 > \sigma_0^2$: $T > \chi_{1-\alpha; n-1}^2$ resp. one-sided test $\sigma^2 \leq \sigma_0^2$: $T < \chi_{\alpha; n-1}^2$.

Remark. The further the ratio $\frac{s^2}{\sigma_0^2}$ deviates from 1, the more likely you are to reject the null hypothesis.

Example *The size of prey.*

We use the example in ▷ M. FOLEY's *RPubs*: The size of prey (millimeters) of two species of net-casting spiders, deinopis (X) and menneus (Y) are sampled for 10 spiders of each species.

What is the difference in the variance of the prey of the two species?

The null hypothesis is the assumption of the same mean with a 5% probability of error.

The size of prey (millimeters) of two species

X: 12.43 11.71 14.41 11.05 9.53 11.66 9.33 11.71 14.35 13.81

We have $\alpha = 0.05$ as the significance level of rejecting the null hypothesis.

◦ *Solution* of this example in REDUCE:

```
| X:={12.43, 11.71, 14.41, 11.05, 9.53, 11.66, 9.33, 11.71, 14.35, 13.81}$
| vartest(X, 0.05, 1);
```

..... reject H0				
T=chi2	p	alpha	ICL	uCL
29.208	0.0011956	0.05	1.5354	10.816

◦ Result: the p-value = 0.001196 < $\alpha = 0.05$, so reject H0 that $s^2 = \sigma^2$.



View step-by-step solution *code*, run REDUCE.



View vartest in action *code*, run REDUCE.

General information

General mathematical information about the concept is here ▷ WIKIPEDIA : χ^2 -test
 Syntax and semantic of the implementation is here ▷ MATLAB : **vartest**

4.6.1 Exercises

Exercise 201. (the NIST example, cf. GEAR)

A chi-square test was performed for the GEAR.DAT data set. The observed variance for the 100 measurements of gear diameter is 0.00003969 (the standard deviation is 0.0063). The following are the data used for the chi-square test for the variance example. We will test the null hypothesis that the true variance is equal to 0.01.

a. Do the example from NIST with the following first 20 items from a. of the data GEAR.DAT. The following data is the gear diameter:

```
1.006  0.996  0.998  1.000  0.992
0.993  1.002  0.999  0.994  1.000
0.998  1.006  1.000  1.002  0.997
0.998  0.996  1.000  1.006  0.988
```

b. Calculate the test statistic value for the sample of 20 items and compare it with the value of 0.3903 for the whole sample.
c. Check, whether the test statistic value is much smaller than the lower critical value, so we also may reject the null hypothesis and conclude that the variance is not equal to 0.01.

Exercise 202. (example from $\mathcal{R}$)

Check our solution using the `varTest(.)` function of R.


```
R> library(EnvStats)
> X <- c(12.43, 11.71, 14.41, 11.05, 9.53, 11.66, 9.33, 11.71, 14.35, 13.81)
> alpha <- 0.05
> sigma = 1.00
> result <- varTest(X, "two.sided", 0.95, sigma.squared = 1)
> print(result)
```

Null Hypothesis:	variance = 1	[=sigma.squared]
Alternative Hypothesis:	True variance is not equal to 1	
Test Name:	Chi-Squared Test on Variance	
Estimated Parameter(s):	variance = 3.245299	
Data:	X	
Test Statistic:	Chi-Squared = 29.20769	
Test Statistic Parameter:	df = 9	
P-value:	0.001195555	
95% Confidence Interval:	LCL = 1.535407	
	UCL = 10.816103	

4.7 F test

The F -test tests two normally distributed independent samples for equal variance σ^2 .

Mental image



Visualization of F -test on equal variances, F_α is the α -quantile of F .

Figure 36: The values shown are from the solution of the example, see below.

$1 - \alpha/2$ gives the area spanned inside between the limits $|\cdots|$.

Procedure F -test for equal variances

1. **Assumptions** normal distributed samples $X = (x_1, \dots, x_m)$ with $\mathcal{N}(\mu_1, \sigma_1^2)$ and $Y = (y_1, \dots, y_n)$ with $\mathcal{N}(\mu_2, \sigma_2^2)$
2. **Null hypothesis** $H_0 : \sigma_1^2 = \sigma_2^2$ (two-sided)
3. **Test statistics**

$$T := \frac{s_1^2}{s_2^2}$$

T is for H_0 F -distributed with $df_1 = m - 1$ and $df_2 = n - 1$ degrees of freedom, where s_i are the sample's standard deviations.

4. **Decision** Reject H_0 , if $T < F_{\alpha/2, m-1, n-1} =: \text{fINV}(\alpha/2, m-1, n-1)$.
 $F_{\alpha, \nu}$ is the α -quantile of the F -distribution, i.e. $\text{fINV}(\alpha, \nu)$ in REDUCE.

Remark. one-sided test $\sigma_1^2 > \sigma_2^2$: $T > F_{1-\alpha; m-1, n-1}$ resp. one-sided test $\sigma_1^2 \leq \sigma_2^2$: $T < F_{\alpha; m-1, n-1}$.

Remark. Note, that $F_{\alpha; m-1, n-1}$ is the *critical value* of the F distribution with $m-1$ and $n-1$ degrees of freedom and a significance level of α .

Remark. It should be noted that the F -test is *not robust*, i.e. it is very sensitive to small deviations from the normal distribution.

Example *Groundwater sulfate concentrations.*

We use the example from ▷ M. GIMOND : Groundwater sulfate concentrations are monitored at a contaminated site over the course of a year. Those concentrations are compared to ones measured at background sites for the same time period. We seek to compare the concentration of sulfates between background sites and a contaminated well (data taken from MILLARD et al., p. 418). Did the two samples have equal variances?

The concentrations of sulfate (in ppm) for both sites are as follows:

Groundwater sulfate concentrations in ppm

Contaminated:	600	590	590	630	610	630
Background:	560	530	570	490	510	550 550 530

◦ *Solution* of this example in REDUCE:

```
| X := {560, 530, 570, 490, 510, 550, 550, 530}; %-- Background
| Y := {600, 590, 590, 630, 610, 630}; %-- Contaminated
| varTest(X, Y, 0.05); %-- F test to compare two variances
|
```

... accept H0			
T	p-value	Cl.l	Cl.h
2.1163	0.42634	0.30882	11.185

Result: the $p\text{-value} = 0.4263 > \alpha = 0.05$, and with such a high p , we cannot reject the null hypothesis and therefore state that the variances between both populations are the same.



View step-by-step solution *code*, run REDUCE.



View varTest in action *code*, run REDUCE.

General information

General mathematical information about the concept is here ▷ WIKIPEDIA : F -test
 Syntax and semantic of the implementation is here ▷ MATLAB : `vartest2`

4.7.1 Exercises

Exercise 203. (Prices of shares, cf. f-test)

The prices of two shares A and B over eleven months at the beginning of each month are

A: 44.32 49.28 65.12 69.44 59.20 70.56
60.60 87.16 102.60 97.96 93.00

B: 59.35 54.24 59.61 58.68 67.28 60.28
65.93 73.63 82.84 76.01 78.55

Verify: Your risk when investing in shares B of Bernhard is significantly lower than when investing in shares A of Albert.

Exercise 204. (groundwater example in $\mathcal{R}$)

Check the following solution using `var.test(..)`²⁸ function of base $\mathcal{R}$.

```
R> Backg <- c(560, 530, 570, 490, 510, 550, 550, 530)
> Conta <- c(600, 590, 590, 630, 610, 630)
> var.test(Backg, Conta, alternative="two.sided")

      F test to compare two variances

data:  Backg and Conta
F = 2.1163,  num df = 7,  denom df = 5,  p-value = 0.4263
alternative hypothesis: true ratio of variances is not equal to 1
95 percent confidence interval:
      0.3088156   11.1853404
sample estimates:
ratio of variances
      2.116337
```

²⁸GIMOND, *ibid.*: “Note that the `var.test()` computes the F ratio using the first variable name in the list as the numerator. For example, had we reversed the order of variables (i.e. `var.test(Conta, Backg, alternative=“two-sided”)`), the returned F value would be the inverse of the original F value, or $1/2.12 = 0.47$. The p value would have stayed the same however.”

B: Parameter tests

We quote HERMANN[7, p.134]:

“In contrast to the parameter tests of the last Chapter, non-parametric tests do not require the presence of a normal distribution of the data or a specific parameter. The parameterfree tests are therefore based on a nominal or ordinal scale of the data as they are usually given in sociology, pedagogy and psychology.

Since a nominal or ordinal scale, in contrast to interval scales, do not allow a mean concept, only combinatorial arrangements such as rank or sign distributions, iterations ('runs') or information statistics can be evaluated. This makes parameterfree test methods universally applicable, but leads, for example, to information loss with normally distributed data.”

Hereinafter, we do all tests first in a *semi-automatic* way, so that the user can follow the essential steps and can do these steps also by hand. This *procedural* way is accompanied by a *fully automatic solution* using functions in CAS REDUCE.


We discuss using REDUCE the

1. Sign tests
2. WILCOXON tests
3. MANN–WHITNEY U test
4. PEARSON χ^2 test
5. FISHER test
6. MCNEMAR test

4.8 One Sample Sign Test

The *sign test* checks the symmetry of a sample with respect to a central value, e.g., the median, by counting the signs that result from the differences to the median.

Mental image



Visualization of F -test on equal variances, F_α is the α -quantile of F .

Figure 37: The values shown are from the solution of the example, see below.

$1 - \alpha/2$ gives the area spanned inside between the limits $|\cdots|$.

Procedure *One-Sample sign-test*

1. **Assumptions** ordinal distributed independent sample $X = (x_1, \dots, x_m)$ with continuous CDF symmetric to the median $x_{0.5}$
2. **Null hypothesis** $H_0 : x_{0.5} = x_0$ (two-sided)
3. **Test statistics**

$$T := \sum_{i=1}^n y_i \quad \text{where} \quad y_i = \begin{cases} 1, & x_i - x_0 \geq 0, \\ 0, & x_i - x_0 < 0. \end{cases}$$

T for H_0 is $\sum_{k=0}^n B(k, n, 0.5)$ -distributed, where $B(k, n, p)$ is the density of the binomial distribution.

The T -Statistic the number of positive differences between the data and the hypothesized median $x_{0.5}$.

4. **Decision** Reject H_0 , if $T < B_{\alpha/2} =: \text{binINV}(\alpha/2)$ or $T > B_{1-\alpha/2}$.
 B_α is the α -quantile of the binomial B -distribution, i.e. $\text{binINV}(\alpha)$ in REDUCE.

Example *Average age of men at first marriage.*

HERRMANN [7, p.135] gives the following example of a one-sample sign test: The following table shows the average age of men at first marriage (Federal Republic of Germany, 1971-1984):

	average age of men at first marriage													
year :	1971	1972	1973	1974	1975	1976	1977	1978	1979	1980	1981	1982	1983	1984
age:	26.0	25.6	25.5	25.6	25.3	25.6	25.7	25.9	26.0	26.1	26.3	26.6	26.9	27.0

The table is to be examined for a trend.

Check, if the age of 1st marriage has median 26.

Use the one-sample-sign-test to analyze the experiment.

◦ *Solution* of this example with REDUCE :

```
| X := {26.0, 25.6, 25.5, 25.6, 25.3, 25.6, 25.7,
|      25.9, 26.0, 26.1, 26.3, 26.6, 26.9, 27.0}$
| signtest(X, 26.0, "both") ;
```

s	p
5	0.77441

Interpretation:

5 differences with a positive sign,

i.e. 5 data elements have a value > median,

probability p for $s = 5$ is $0.7744 > 0.05 = \alpha$

Result: The null hypothesis, $H_0 : \text{median} = 26$, can be rejected.



View step-by-step solution *code*, run REDUCE.



View signTest in action *code*, run REDUCE.

General information

General mathematical information about the concept is here ▷ WIKIPEDIA : Sign-test

Syntax and semantic of the implementation is here ▷ MATLAB : `signtest`

4.8.1 Exercises

Exercise 205. (Remark and surplus)

We can do the decision about the H_0 in the ‘age’ example using alpha-quantiles of the binomial distribution. Look up our function `binoinv()`. Then do:

```
...
alpha := 0.05 $
nt     := 12 $      %-- n total
T_     := 5 $       %-- T_ = s test statistic
qt1 := binoinv(alpha/2, nt, 0.5);      %-- alpha quantile
qt2 := binoinv(1-alpha/2, nt, 0.5);    % -- alpha quantile
%      5      3          5      9
if (T < qt1) or (T > qt2) then write("reject H0") else write("accept H0");
```

Exercise 206. (Check example ‘Average age of men..’ with OCTAVE)

We use OCTAVE function `signtest()`.

```
octave:1> pkg load statistics
octave:2> X=[26.0, 25.6, 25.5, 25.6, 25.3, 25.6, 25.7,
            25.9, 26.0, 26.1, 26.3, 26.6, 26.9, 27.0];
octave:3> [p,h,stats] = signtest(X,26)
ans =
      p      = 0.7744
      h      = 0
      stats =
              zval = NaN
              sign = 5
              (wL: i.e. 5 data elements have a value > median)
```

Exercise 207. (example by SPRENT from WIKIPEDIA)

Do the cited example of a sign test for the median of a single sample in REDUCE.

Exercise 208. (example by $\mathcal{R}$)

Reproduce the following example from the $\mathcal{R}$ documentation in REDUCE:

```
R> x <- c(7.8, 6.6, 6.5, 7.4, 7.3, 7., 6.4, 7.1, 6.7, 7.6, 6.8)
R> SIGN.test(x, md = 6.5)
% Verify:
s = 9, p-value = 0.02148
alternative hypothesis: true median is not equal to 6.5
95 percent confidence interval:
 6.571273 7.457455
sample estimates: median of x
```

4.9 Two Sample Sign Test

The *Two Sample Sign test* tests for two paired samples for the same central tendency by counting the signs that result from the differences between corresponding data items.

Mental image




Figure 38: Visualization of *sign2*-test on two samples X and Y with $d_i = x_i - y_i$.
 ●: sample $X = (7.8, 6.6, 6.5, 7.4, 7.3, 7.0, 6.4, 7.1, 6.7, 7.6, 6.8)$.
 ●: sample $Y = (7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8)$.

Procedure *Two-Sample sign-test*

1. **Assumptions** two samples $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$ with independent differences $D = (x_1 - y_1, \dots, x_n - y_n)$
2. **Null hypothesis** $H_0 : Pr(X < Y) = Pr(X > Y)$ (two-sided)
3. **Test statistics**

$$T := \sum_{i=1}^n d_i \quad \text{where} \quad d_i = \begin{cases} 1, & x_i - y_i > 0, \\ 0, & x_i - y_i < 0. \end{cases}$$

T for H_0 is $\sum_{k=0}^T B(k, n, 0.5)$ -distributed, where $B(k, n, p)$ is the density of the binomial distribution.

The T -Statistic the number of positive differences $d_i > 0$ between corresponding data items in X and Y .

4. **Decision** Reject H_0 , if $T < B_{\alpha/2} =: \text{binINV}(\alpha/2)$ or $T > B_{1-\alpha/2}$.
 B_α is the α -quantile of the binomial B -distribution, i.e. $\text{binINV}(\alpha)$ in REDUCE.

Example *IQ of twin pairs.*

HERRMANN [7, p.142] gives the following example of a two-sample sign test: Given the following observation table of weights and intelligence quotients (IQ) of single twins after CHURCHILL & WILLERMAN. Test the null hypothesis that the IQ scores of the twins are independent of their birth weight.

IQ_ℓ designate the lighter twin pairs and IQ_h the heavier twin pairs.

	<i>IQ of twin pairs</i>													
<i>pair:</i>	1	2	3	4	5	6	7	8	9	10	11	12	13	14
IQ_h :	97	79	100	100	100	124	95	80	91	108	91	90	104	119
IQ_ℓ :	97	70	101	106	85	123	84	70	84	106	97	90	92	104

Use the two-sample-sign-test to analyze the experiment.

◦ *Solution* and explanation of this example with REDUCE :

```
| IGh := {97,79,100,100,100,124,95,80,91,108,91,90,104,119}$
| IQl := {97,70,101,106,85,123,84,70,84,106,97,90,92,104}$
| signtest2(IGh, IQl, both);
```

Interpretation:

$T = 9$ differences with a positive sign,
i.e. 9 data pairs have a difference greater 0.

Probability for $T = 8$ is $p = 0.1459 > 0.05 = \alpha$

s	p
9	0.146

Result: The null hypothesis, H_0 : 'IQ independent of weight', can not be rejected with a 5 % probability of error.



View step-by-step solution *code*, run REDUCE.



View signTest2 in action *code*, run REDUCE.

General information

General mathematical information about the concept is here $\triangleright$ WIKIPEDIA : Sign-test
Syntax and semantic of the implementation is here $\triangleright$ MATLAB : `signtest`

4.9.1 Exercises

Exercise 209. (Fig. 38) Do the sign test for the data of Fig.38, i.e.

X=(7.8, 6.6, 6.5, 7.4, 7.3, 7.0, 6.4, 7.1, 6.7, 7.6, 6.8)

Y=(7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8)

- Check the differences on the plot.
- How to determine the test statistic T on the graph?

Exercise 210. (Example from MATLAB, cf. `signtest`)

Do the example from the MATLAB documentation in REDUCE using the data from section ‘Medians of Paired Samples’.

- Test the hypothesis of zero median for the difference between the two paired samples using the exact methods.
- z-Statistic: For a large sample, `signtest` uses the z-statistic to approximate the p-value. Program the z-statistic in REDUCE.
- Check the z-statistic for the data of a.

Exercise 211. (example from OCTAVE)


Check the following solution of the example 4.9 using `signtest` function of `statistics` package of OCTAVE.

```
octave:1> pkg load statistics
octave:2> IQh = [97,79,100,100,100,124,95,80,91,108,91,90,104,119];
          IQl = [97,70,101,106,85,123,84,70,84,106,97,90,92,104];
octave:3> [p,h,stats] = signtest(IQh, IQl)
ans =
      p      = 0.1460
      h      = 0
      stats = scalar structure containing the fields:
              zval = NaN
              sign = 9
              (wL: i.e. 9 data pairs have a difference greater 0)
```

4.10 One Sample WILCOXON Test

The *One Sample Sign* test only registers the signs of the differences to the median. The *One Sample Sign WILCOXON* test also considers the absolute values of the differences. These values are sorted and ranked, and the *rank sum* of the positive differences is calculated. This sum represents the WILCOXON test statistic W .

Mental image



Visualization of WILCOXON-One-Sample test on the sample X .

Figure 39: ●: sample $X = (7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8)$.

⋄: items with index $i = 2, 4, 5, 8, 10$ have positive distance to the median.

Procedure *One-SampleWILCOXON-test*

1. **Assumptions** one sample $X = (x_1, \dots, x_n)$ symmetrical w.r.t. the median.
2. **Null hypothesis** $H_0 : \text{median}(X) = x_0$ (two-sided)
3. **Test statistics** with $d_i := x_i - x_0$

$$W := \sum_{i=1}^n c_i \cdot \text{Rank}(|d_i|) \quad \text{where} \quad c_i := \begin{cases} 1, & x_i - x_0 > 0, \\ 0, & x_i - x_0 < 0. \end{cases}$$

4. **Decision** Reject H_0 , if $W \leq w_{\alpha/2}$ or $W \leq w_{1-\alpha/2}$.
 w_α is the critical value of the tabulated WILCOXON-distribution, $\triangleright$ U-test

Example *median length of pygmy sunfish*

We quote the example from PennState Eberly College at ▷ Example 2.2.

Let X_i denote the length, in centimeters, of a randomly selected pygmy sunfish, $i = 1, \dots, 10$. If we obtain the data set 5.0 3.9 5.2 5.5 2.8 6.1 6.4 2.6 1.7 4.3 can we conclude that the median length of pygmy sunfish differs significantly from 3.7 centimeters?

Solution step-by-step and maybe by hand using 3.7 for the hypothesized median.

1. We construct a table with the item numbers No_i , the items X_i itself, the items centered around 3.7 i.e. $X_i - 3.7$, their absolute values $|X_i - 3.7|$, the ranked absolute values R_i and the signed ranked absolute values *signed* R_i :

No	1	2	3	4	5	6	7	8	9	10
Xi	5	3.9	5.2	5.5	2.8	6.1	6.4	2.6	1.7	4.3
Xi-3.7	1.3	0.2	1.5	1.8	-0.9	2.4	2.7	-1.1	-2	0.6
Xi-3.7	1.3	0.2	1.5	1.8	0.9	2.4	2.7	1.1	2	0.6
Rank Ri	5	1	6	7	3	9	10	4	8	2
signed Ri	5	1	6	7	0	9	10	0	0	2

2. $W = \text{sum}(\text{signed } R_i) = 40$

3. Check the ▷ **w-table**; with $n = 10$, a small sample size, the upper and lower percentiles of the Wilcoxon signed rank statistic is: $n = 10 : \Pr(T \geq W = 40) = 0.116$
4. Therefore, our P-value is $2 \times 0.116 = 0.232$. Because our P-value is large, we cannot reject the null hypothesis.

◦ *Solution* and explanation of this example with REDUCE

```
| X := {5.0, 3.9, 5.2, 5.5, 2.8, 6.1, 6.4, 2.6, 1.7, 4.3}$
| signrank(X, 3.7);
```

W	T
40.0	1.2741



View step-by-step solution *code*, run REDUCE.



View signRank in action *code*, run REDUCE.

General information

General mathematical information about the concept is here ▷ WIKIPEDIA : signrank
Syntax and semantic of the implementation is here ▷ MATLAB : **signrank**.

4.10.1 Exercises

Exercise 212. (pygmy sunfish, example from psu.edu)Check the solution of example.4.10 using `wilcox.test()` function of R.

```
R> X=c(5.0, 3.9, 5.2, 5.5, 2.8, 6.1, 6.4, 2.6, 1.7, 4.3)
R> >   wilcox.test(X, mu = 3.7)

      Wilcoxon signed rank exact test

data:  X
V = 40, p-value = 0.2324
alternative hypothesis: true location is not equal to 3.7
```

Interpretation: Based on the results of the test, (at the significance level of 0.05) we reject the null hypothesis.

Exercise 213. (Example at Fig.39)

Calculate the WILCOXON test statistic W for the sample X with hypothesized median $x_0 = 7$, where

$$X = (7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8)$$

from figure.35.

Exercise 214. (placebo-controlled clinical trial, cf. nbisweden)

Let's imagine we are a part of team analyzing results of a placebo-controlled clinical trial to test the effectiveness of a sleeping drug. We have collected data on 10 patients when they took – a sleeping drug and when they took a placebo.

The hours of sleep recorded for each study participant:

```
X=(6.1, 6.0, 8.2, 7.6, 6.5, 5.4, 6.9, 6.7, 7.4, 5.8)    -- drug
Y=(5.2, 7.9, 3.9, 4.7, 5.3, 7.4, 4.2, 6.1, 3.8, 7.3)  -- placebo
```

Before we investigate the effect of drug, a senior statistician ask us: 'Is the median sleeping time without taking the drug significantly less than the recommended 7 h of sleep?'

For the *solution* we therefore consider only Y .

4.11 Two Sample WILCOXON Test

The *Two Sample WILCOXON Test* tests for two paired samples X and Y for the same central tendency by counting the rank sums of their differences. It is a non-parametric alternative to the paired t-test, used when you have paired or dependent data, such as two measurements from the same individual.

Mental image




Figure 40: Visualization of Two Sample WILCOXON-test on two samples X and Y
 ●: sample $X = (7.8, 6.6, 6.5, 7.4, 7.3, 7.0, 6.4, 7.1, 6.7, 7.6, 6.8)$.
 ●: sample $Y = (7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8)$.
 ●: For pair 6 we have $d_6 = x_6 - y_6 = 7 - 6 = +1$.

Procedure *Two-Sample WILCOXON-test*

1. **Assumptions** two samples $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$ with differences $D = (x_1 - y_1, \dots, x_n - y_n)$
2. **Null hypothesis** $H_0 : \text{median}(D) := \tilde{d} = 0$ (two-sided)
3. **Test statistics** W with $d_i := x_i - \tilde{x}$ is calculated by

$$W := \sum_{i=1}^n z_i \cdot \text{Rank}(|d_i|) \quad \text{where} \quad z_i := \begin{cases} 1, & x_i - \tilde{d} > 0, \\ 0, & x_i - \tilde{d} < 0. \end{cases}$$

4. **Decision** Reject H_0 , if $W \leq w_{\alpha/2}$ or $W \leq w_{1-\alpha/2}$.
 w_α is the critical value of the tabulated WILCOXON-distribution, $\triangleright$ U-test

Example *Figure.40*

We do the two-sample WILCOXON-test for the data of figure.40.

We have to test the null hypothesis that the median of the differences $X - Y$ is 0, i.e. to test if the median is the same for the first and second sample.

Solution along the rows of the ‘rank’ table:

1	2	3	4	5	6	7	8	9	10	11	item number
7.8	6.6	6.5	7.4	7.3	7	6.4	7.1	6.7	7.6	6.8	sample X
7	7.6	6.5	9	7.3	6	4.4	8.1	4.2	7.8	5.8	sample Y
0.8	-1	0	-1.6	0	1	2	-1	2.5	-0.2	1	differences $d_i := x_i - y_i$
4	6.5	1.5	9	1.5	6.5	10	6.5	11	3	6.5	rank'ing of $ d_i $
↓					↓	↓		↓		↓	

1. rank'sums: $T^+ = 4 + 6.5 + 10 + 11 + 6.5 = 38$ ²⁹ and analog $T^- = 6.5 + 9 + 6.5 + 3 = 25$.

2. test statistic $T := \min(T^+, T^-) = \min(38, 25) = 25$.³⁰

3. tabulated critical value $w_{crit}(0.05; 9) = 6$. $[n = 11 - 2 = 9, \text{ because of 'ties' } (7.3, 7.3) \text{ and } (6.5, 6.5)]$

4. Because $T = 25 > 6 = w_{crit}$, accept H_0 .

5. Check the plausibility of the decision by Figure.35.

◦ *Solution* and explanation of this example with REDUCE

```
| X := {7.8, 6.6, 6.5, 7.4, 7.3, 7.0, 6.4, 7.1, 6.7, 7.6, 6.8};
| Y := {7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8};
| signrank2(X,Y);
```

Wp	Wm	T
38	25	25



View step-by-step solution *code*, run REDUCE.



View signrank2 in action *code*, run REDUCE.

General Information

General mathematical information about the concept is here $\triangleright$ WIKIPEDIA : rank test
 Syntax and semantic of the implementation is here $\triangleright$ MATLAB : `signrank(x,y)`

²⁹If one or more differences $x_i - y_i = 0$, which is the case here, the corresponding observations are not included in the test, i.e. these pairs are deleted from the sample. Then we get the same value as OCTAVE: $W = T^+ = \text{signedrank} = 28$.

³⁰In contrast to the often used $T := \min(T^+, T^-)$, in our definition in the text only T^+ is used to define $W := T^+$. This seems to be the same choice as in OCTAVE/MATLAB.

4.11.1 Exercises

Exercise 215. (sleep study, cf.psu.edu)

Going back to our sleep study, now we are ready to examine whether there is enough evidence to reject a null hypothesis of median of the differences between the paired observations is equal to 0.

a. Check the solution using `wilcox.test()` function of $\mathcal{R}$.

```
##### check using R - our W is R's V ;)
R> X=c(6.1, 6.0, 8.2, 7.6, 6.5, 5.4, 6.9, 6.7, 7.4, 5.8)
R> Y=c(5.2, 7.9, 3.9, 4.7, 5.3, 7.4, 4.2, 6.1, 3.8, 7.3)
R>
R> wilcox.test(x = Y, y = X, alternative = "two.sided", mu = 0,
               paired = TRUE, exact = F)

Wilcoxon signed rank test with continuity correction
data:  Y and X
V = 15, p-value = 0.2213
alternative hypothesis: true location shift is not equal to 0
```

b. Do the test using REDUCE.

Exercise 216. (Example at Fig.40 using OCTAVE)


Check our solution using `signrank` function of MATLAB/OCTAVE.

```
octave:1> pkg load statistics
octave:2> X = [7.8, 6.6, 6.5, 7.4, 7.3, 7.0, 6.4, 7.1, 6.7, 7.6, 6.8];
octave:3> Y = [7.0, 7.6, 6.5, 9.0, 7.3, 6.0, 4.4, 8.1, 4.2, 7.8, 5.8];
octave:4> [p,h,stats] = signrank(X,Y)
           p = 0.5625
           h = 0
           stats = [...]
                signedrank = 28      zval = NaN
```

4.12 MANN–WHITNEY U test

The MANN–WHITNEY U test is designed for independent samples. Since the sample sizes no longer necessarily are equal, it is not possible to calculate pairwise differences. Therefore, the two samples are combined into a single total sample, and the rank values of all sample items are determined. The test statistic W is then the smaller sum of the ranks from the two samples.

Mental image



Visualization of U -test process on sample $X : 17, 34, 8, 1, 43, 4, 23$.

Figure 41: **1:** $X = (17, 34, 8, 1, 43, 4, 23) \xrightarrow{\text{sort}} (1, 4, 8, 17, 23, 34, 43) = X_{\text{sorted}}$.
2: $X_{\text{sorted}} = (1, 4, 8, 17, 23, 34, 43) \xrightarrow{\text{Rank}} (1, 2, 3, 4, 5, 6, 7) = X_{\text{ranked}}$.

Procedure MANN–WHITNEY U test

- Assumptions** two independent samples $X = (x_1, \dots, x_m)$ and $Y = (y_1, \dots, y_n)$ with $m < n$ and concepts (e.g. **median**) F resp. G .
- Null hypothesis** $H_0 : F = G$ (two-sided)
- Test statistics**

$$U := W - \frac{m * (m + n + 1)}{2} \quad \text{where} \quad W := \sum_{i=1}^m \text{Rank}(X_i)$$

- Decision** Reject H_0 , if $U < w_{\alpha/2}$ or $U > w_{1-\alpha/2}$.
 w_{α} is the critical value of the tabulated WILCOXON-distribution, $\triangleright$ WILCOXON : U-test

Example *Butterflies on sunny vs cloudy days.*

ZOEFL [29, p.104] gives the following example of a MANN–WHITNEY U test: As part of a biological study, butterflies were observed at various times within a fixed timeframe. The weather conditions were recorded, with a general classification into sunny and cloudy periods. The results are shown in the following table:

<i>butterflies on sunny vs cloudy days</i>								
sunny:	6	15	35	35	62	73	98	112
cloudy:	1	4	8	17	23	34	43	

An U test should explain whether the difference between the two median values of the samples is significant.

Solution along the columns of the rank table:

sun:	rank:	cloud:	rank:
6	3	1	1
15	5	4	2
35	9.5	8	4
35	9.5	17	6
62	12	23	7
73	13	34	8
98	14	43	11
112	15		
sum:	81		39

1. ranksum's $R_1 = 81$ and $R_2 = 39$.

2. U values

$$U_1 = R_1 - n_1(n_1 + 1) \cdot 0.5 = 11 \text{ and}$$

$$U_2 = R_2 - n_2(n_2 + 1) \cdot 0.5 = 45.$$

3. The tabulated critical value for
 $U = \min(U_1, U_2) = 11$ is 10.^a

4. Because $U = 11 > 10 = U_{crit}$,
 reject H_0 .

^ae.g. Mann-Whitney table

◦ *Solution* and explanation of this example with REDUCE

```
| X : [6, 15, 35, 35, 62, 73, 98, 112]$
| Y : [1, 4, 8, 17, 23, 34, 43]$
| ranksum(X,Y)$
```

U	p	Ho
11.0	0.97543	reject H0



View step-by-step solution *code*, run REDUCE.



View U-test in action *code*, run REDUCE.

General information

General mathematical information about the concept is here $\triangleright$ WIKIPEDIA : U-test
 Syntax and semantic of the implementation is here $\triangleright$ MATLAB : ranksum

4.12.1 Exercises

Exercise 217. (Redo the Example 4.12 using OCTAVE)

- Check the solution using `ranksum` function of `statistics` package of Octave.

```
octave:1> pkg load statistics
octave:2> x=[6, 15, 35, 35, 62, 73, 98, 112];
octave:3> y=[1, 4, 8, 17, 23, 34, 43];
octave:4> [p, h, stats] = ranksum(x, y)
ans : p = 0.050039      h = 0      stats = ranksum = 81
```

Exercise 218. (Redo the Example 4.12 using $\mathcal{R}$)

- Check the solution using `wilcox.test` function of $\mathcal{R}$.

```
R> X=c(6, 15, 35, 35, 62, 73, 98, 112)
R> Y=c(1, 4, 8, 17, 23, 34, 43)
R> wilcox.test(x = Y, y = X, alternative="two.sided",
               paired=FALSE, exact = F)
```

```
Wilcoxon rank sum test with continuity correction
data: Y and X
W = 11, p-value = 0.05598
alternative hypothesis: true location shift is not equal to 0
```

Exercise 219. (two sorts of alloy, cf. SPIEGEL [18, p. 447, P454])

See the example in Spiegel w.r.t. two sorts of alloy

```
A1 = (18.3, 16.4, 22.7, 17.8, 18.9, 25.3, 16.1, 24.2)
A2 = (12.6, 14.1, 20.5, 10.7, 15.9, 19.6, 12.9, 15.2, 11.8, 14.7)
```

Do a U -test to check if $A1$ and $A2$ come from the same population.

Result: reject H_0 at the 0.05 level.

4.13 PEARSON's Chi-squared test & Contingency Tables

Many experimental results can be represented in the form of so called ‘*four-field tables*’ or multi-field tables, also called *contingency tables*. The tests developed for this purpose do not require any specific parameters of the populations, they are therefore *non-parametric test* procedures. A very popular test is the *chi-square test* for four-field tables.

Mental image




Figure 42: Visualization of Four-field Table for an experimental outcome.
 • experiment outcome as frequency table $T = \begin{pmatrix} a & b \\ c & d \end{pmatrix} = \begin{pmatrix} 17 & 122 \\ 31 & 109 \end{pmatrix}$.

Procedure PEARSON's *independence test*

a	b
c	d

- Assumptions** The frequencies of the dichotomous characteristics A and B are counted on $n = a + b + c + d$ subjects forming a random sample; the observations are made independently.
- Null hypothesis** H_0 : the characteristics A and B are independent (two-sided)
- Test statistics**

$$\chi^2 := \frac{n \cdot (ad - bc)^2}{(a + b)(c + d)(a + c)(b + d)} \quad \text{is approximately } \chi_1^2 \text{ distributed.}$$

- Decision** Reject H_0 , if $\chi^2 > \chi_{1;1-\alpha/2}^2$.
 ▷ U-test

Example *Does vitamin C help against colds?*

We do the example in [7, p.153, 160]. Studies by the physician G. RITZEL from 1961, who tested the influence of vitamin C on cold prophylaxis on skier in a double-blind, randomized experiment, was given in the following four-field table.

	cold	no cold
C	17	122
placebo	31	109

Test the null hypothesis $Pr[cold] = Pr[no\ cold]$ for a 5% probability of significance.

Solution

1. $n = 279$.
2. test statistic $\chi = 4.1407$.
3. the p-value $1 - \chi(0.05; 1) = 0.0418$.
4. Because $p = 0.042 < 0.05 = \alpha$, accept H_0 .

◦ *Solution* and explanation of this example with REDUCE :

```
|  chisqTest(17,122, 31,109);
```

```
[  Chi      p      phi: ]
[ 4.1407  0.041864  0.12182 ]
```



View step-by-step solution *code*, run REDUCE.



View ChisqTest in action *code*, run REDUCE.

General information

General mathematical information about the concept is here ▷ Contingency table
 Syntax and semantic of the implementation is here ▷ MATLAB : **crosstab**

4.13.1 Exercises

Exercise 220. (red-green color blindness, cf. [7, p.161])

In a red-green color blindness study, the following four-field table was obtained.

	male	female
normal	: [8324	9031]
red-green-blind	: [725	40]

Check via a chi-squared test statistic for the dependence of blindness on gender.

DECISION: The χ^2 test statistic is 640.3.

The associated p-value p of the χ^2 distribution is 1.

With a 5color blindness of the gender can just be rejected.

Exercise 221. (Example from MATLAB, cf. signtest)

Do the example from the MATLAB documentation in REDUCE using the data from section ‘Medians of Paired Samples’.

- Test the hypothesis of zero median for the difference between the two paired samples using the exact methods.
- z-Statistic: For a large sample, signtest uses the z-statistic to approximate the p-value. Program the z-statistic in REDUCE.
- Check the z-statistic for the data of a.

Exercise 222. (check with $\mathcal{R}$)

Check the solution of the example 4.13 using `chisq.test` function of $\mathcal{R}$.

```
# Step 1: Creating a contingency table
R> data <- matrix(c(17, 122, 31, 109), nrow = 2)
# Step 2: Applying the chi-square test function
> chisq.test(data)

      Pearson's Chi-squared test with Yates' continuity correction
data:  data
X-squared = 4.1407, df = 1, p-value = 0.04186
```

or if one only wants the p- value

```
##### Check with R
R> x <- matrix(c(17, 122, 31, 109), ncol = 2)
> chisq.test(x)$p.value
[1] 0.04186438
```

4.14 FISHER test

The FISHER test is used for 2×2 contingency tables where the total frequency is less than or equal to 40. This test is also called the *exact* FISHER test because it calculates the exact probability according to the hypergeometric distribution.

Mental image



height	age of 31 US presidents	
small (X)	67,79,80,85,90	
tall (Y)	53,56,60,60,63,63,64,64,65,66,67,67,68,70,71,71,72,73,73,74,77,78,78,83,88,90	
Sample	$\geq \tilde{x}$	$< \tilde{x}$
X	1	4
Y	14	12

X = sample of small US presidents

Y = sample of tall US presidents

$\tilde{x}$: median of the total sample $X \cup Y$ is 70.

Figure 43: Motivation for exact FISHER test w.r.t. the natural death year before/after median. Is the death independent of age and height?

Procedure FISHER's independence test

1. **Assumptions** The frequencies of the dichotomous characteristics A and B are counted on $n := a + b + c + d$ subjects forming a random sample; the observations are made independently.
2. **Null hypothesis** H_0 : the characteristics A and B are independent (two-sided)

3. **Test statistics** given the corner distribution

		sum
a	b	a+b
c	d	c+d
a+c	b+d	n

, the test statistics

P is

$$P := \sum_{x=0}^{\min(a+b, a+c)} \frac{\binom{a+c}{x} \cdot \binom{b+d}{a+b-x}}{\binom{n}{a+b}} = \frac{(a+b)! \cdot (c+d)! \cdot (a+c)! \cdot (b+d)!}{n! \cdot a! \cdot b! \cdot c! \cdot d!}$$

4. **Decision** Reject H_0 , if $P < \text{hypINV}(\alpha/2)$ or $P > \text{hypINV}(1 - \alpha/2)$.

Example *Is natural death year independent of age and height?*

We do the example in [7, p.163, 160]. There was a random sample of 31 US presidents drawn, who died of natural causes. The whole sample was divided into tall (X) and small (Y) persons:

		<i>age at year of death</i>											
tall:	67	79	80	85	90								
small:	53	56	60	60	63	63	64	64	65	66	67	67	68
	70	71	71	72	73	73	74	77	78	78	83	88	90

Is the death independent of age and height?

Solution

1. To solve this, both sets of data are combined into a single sample and it is counted how many elements in the sample are above or below the median $\tilde{x} = 70$. This yields the following contingency table:

Sample	$\geq \tilde{x}$	$< \tilde{x}$
X	1	4
Y	14	12

2. H_0 : the death is independent of age and height

3. $n = 31$.

4. $\text{median}(X \cup Y) =: \tilde{x} = 70$.

5. test statistic $p = 0.1864$ (p-value)

6. Because $p = 0.1864 > 0.05 = \alpha$, reject H_0 . So the FISHER test tells us that there is no statistically significant difference in the association of height and age w.r.t. the natural death.

◦ *Solution* and explanation of this example with REDUCE :

| `fishertest(1,4, 14,12, 0.05);`

Ho	p
1	0.18638



View step-by-step solution *code*, run REDUCE.



View FisherTest in action *code*, run REDUCE.

General information

General mathematical information about the concept ▷ WIKIPEDIA : Fisher exact test
 Syntax and semantic of the implementation is here ▷ MATLAB : `fishertest`

4.14.1 Exercises

Exercise 223. (vitamin C, cf. [7, p.154])

Redo the example 4.13 with the Fisher test.

	Cold	notCold	
Vitamin C:	[17	122]	a+b 139
Placebo :	[31	109]	c+d 140
	48	231	
	a+c	b+d	

```
fishertest(17,122, 31,109, 0.05);
```

Exercise 224. (woman exampe, cf. the exact test in woman)

Do the ‘woman’ example in wikipedia.

Exercise 225. (students example, cf. students)

Do the ‘students’ example in WIKIPEDIA.

Exercise 226. (cancer example in cancer)

Do the ‘cancer’ example in WIKIPEDIA.

Exercise 227. (two formulas)

Calculate P for the following four-fields-table.

	Men	Women	Row Total	
Studying	0	10	: 10	a=0 b=10
Non-studying	12	2	: 14	c=12 d=2
				
Column Total	12	12	: 24	

Exercise 228. (check with $\mathcal{R}$)

Check the solution of the example 4.14 using `fisher.test` function of $\mathcal{R}$.


```
R> fisher.test(rbind(c(1,4),c(14,12)), alternative="less")$p.value
[1] 0.1863799
```

```
R> fisher.test(rbind(c(1,4),c(14,12)))$p.value
# default alternative="both.sided"
[1] 0.3325918
```

4.15 McNEMAR test

McNEMAR's test is a non-parametric test used to analyze paired nominal data. It is a test on a 2×2 contingency table and checks the marginal homogeneity of two dichotomous variables. The test requires one nominal variable with two categories (dichotomous) and one independent variable with two dependent groups, cf. *google* ▷ *McNemar test*

Mental image



Visualization of McNEMAR test: first split the outcome 'solved vs. not solved' w.r.t. the groups A and B .

Figure 44: **Left:** experiment outcome as frequency table $\begin{pmatrix} a' & b' \\ c' & d' \end{pmatrix} = \begin{pmatrix} 83 & 74 \\ 17 & 26 \end{pmatrix}$.

Right: experiment rearranged as new frequency table $\begin{pmatrix} a & b \\ c & d \end{pmatrix} = \begin{pmatrix} 63 & 20 \\ 11 & 6 \end{pmatrix}$.

⇒: McNEMAR needs 'splitted' table $[A \setminus B]$, not 'summary' table $[A \mid B]$.

Procedure McNEMAR's test

1. **Assumptions** The frequencies of the dichotomous characteristics A and B are counted on $n := a + b + c + d$ subjects forming a random sample; the observations are made independently.
2. **Null hypothesis** H_0 : the characteristics A and B are equally distributed. (two-sided)

3. **Test statistics** given the 'splitted' frequency table

A \cdot B	+	-
+	a	b
-	c	d

with $n > 40$, the test statistics

$$\chi := \frac{(b - c)^2}{b + c} \quad \text{is approximately } \chi_1^2 \text{ distributed.}$$

Remark: for $b + c < 40$ the YATES correction $\chi := \frac{(|b - c| - 1)^2}{b + c}$ is used.

4. **Decision** Reject H_0 , if $\chi > \text{chiINV}(1; 1 - \alpha)$.

Example *Survey of tea and coffee drinking*

We do the example of S. MANGIAFICO ▷ McNemar test. Consider a survey of tea and coffee drinking, in which each respondent is asked both if they drink coffee, and if they drink tea. Is coffee more popular than tea? That is, is it more common for someone to drink coffee and not tea than to drink tea and not coffee?

Coffee	Tea	
	Yes	No
Yes	37	17
No	9	25

Solution

1. $a = 37, b = 17, c = 9, d = 25$.
2. $n = 88$.
3. $\chi = 1.8846$.
4. p-value $p = 0.1698$
5. Because $p = 0.1698 > 0.05 = \alpha$, we reject H_0 .

◦ *Solution* and explanation of this example with REDUCE :

```
| mcnemarTest(37,17, 9,25, corrected)
```

```
[ chi      p      Ho ]
[ 1.8846  0.16981  1 ]
```

Remark. Neither coffee nor tea is more popular, specifically because neither the 9 nor the 17 in the table are large relative to the other.

$H_0 = 0$ means ‘reject H_0 ’.



View step-by-step solution *code*, run REDUCE.



View McNemarTest in action *code*, run REDUCE.

General information

General mathematical information is here ▷ WIKIPEDIA/GE : McNemar’s test
 Syntax and semantic of the implementation is here ▷ R : `mcnemar.test`

4.15.1 Exercises

Exercise 229. ('disease' example in wikipedia)

Do the 'disease' example in wikipedia disease

Exercise 230. (students example in wiki)

Do the example in wikipedia/GE studs

Exercise 231. (Presidents performance)

Redo the AGRESTI example in R-manual with REDUCE : agresti

Exercise 232. (EXCEL example in real-statistics)

Do the example in real-statistics/EXCEL with REDUCE: real.stats

Exercise 233. (MCNEMAR test for Fig.44)

Do a MCNEMAR test for the values of fig.44.

Remark. 1. See the implementation of `mcnemarTest` in `McNemar`.

2. The smaller the value of p , the greater the evidence for rejecting the null hypothesis.

3. *ibid.*: Python code excerpt

```
chi2_statistic = (abs(n_min - n_max) - corr) ** 2 / (n_min + n_max)
pvalue = chi2.sf(chi2_statistic, 1)
```

(!) This pvalue is not in coincidence with R, see below.

Exercise 234. (check example.4.15 with $\mathcal{R}$)

Check the solution of the example 4.15 using `mcnemar.test` function of $\mathcal{R}$.

```
R> Matrix = as.matrix(read.table(header=TRUE, row.names=1,
                                text=" Coffee Yes No   Yes 37 17 No 9 25"))
R> mcnemar.test(Matrix)
McNemar's Chi-squared test with continuity correction
McNemar's chi-squared = 1.8846, df = 1, p-value = 0.1698
```

5 Correlation and Bootstrap

We collect the well known measures of correlation using REDUCE. We translate the mathematical definitions into REDUCE code and demonstrate exemplary calls and show examples.

5.1 PEARSON's ρ Correlation coefficient

The correlation coefficient ρ of two random variables X and Y is a measure of their linear dependence. If each variable has n scalar observations, then the PEARSON correlation coefficient ρ is defined as $\frac{\text{cov}(X,Y)}{\sigma_X \cdot \sigma_Y}$.

Mental image




Figure 45: Visualization of PEARSON's correlation coefficient ρ : 6 typical cases.
 row1: from strong negative correlation ('-0.9') until no correlation ('0').
 row2: from mean positive correlation ('+0.5') until strong positive correlation ('+1').

Procedure PEARSON's correlation coefficient for linear correlated samples

1. **Assumptions** two independent random samples $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$ with means $\bar{X}$ resp. $\bar{Y}$
2. **correlation** ρ coefficient is given by

$$\rho(X, Y) := \frac{\sum_{i=1}^n (x_i - \bar{X}) \cdot (y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{X})^2} \cdot \sqrt{\sum_{i=1}^n (y_i - \bar{Y})^2}} = \frac{\text{cov}(X, Y)}{\sigma_X \cdot \sigma_Y}$$


Remark: for $n < 30$ the OLKIN-PRATT correction $\rho := \rho \cdot (1 + \frac{(1-\rho)^2}{2n-6})$ is used.

Example *Height and weight of 12 persons*

We do the example of Z. BOBBITT from *statology/pearson*. The dataset below shows the height and weight of 12 individuals. The scatterplot of this dataset depicts the values of these two variables. Verify: The Pearson correlation coefficient for these two variables is $\rho = 0.836$.

Solution by hand using ρ -formula.

Height (inches)	Weight (lbs)
60	120
65	140
72	180
70	184
74	190
63	160
66	155
68	158
67	170
69	190
70	210
64	150



1. $\bar{X} = 67.33$ and $\bar{Y} = 167.25$.
2. $cov(X, Y) = 916$.
3. $\sigma_X = 174.67$ and $\sigma_Y = 6874.25$.
4. $\rho = \frac{cov(X, Y)}{\sigma_X \cdot \sigma_Y} = 0.836$. Positive correlation: As X variable increases, the Y tends to increase as well.

◦ *Solution* and explanation of this example with REDUCE

```
| X := {60,65,72,70,74,63,66,68,67,69,70,64}$
| Y := {120,140,180,184,190,160,155,158,170,190,210,150}$
| corrccoef(X,Y); | 0.8359
```



View step-by-step solution *code*, run REDUCE.



View correlation coefficient function *code*, run REDUCE.

General Information

General mathematical information is here $\triangleright$ WIKIPEDIA : Pearson corrccoef
 Syntax and semantic of the implementation is here $\triangleright$ MATLAB : corrccoef

5.1.1 Exercises

Exercise 235. (daily calorie intake, cf. [29, p.125])

We use the data from ZOEFEL:

```
X=(3113, 3216, 3495, 2021, 1916, 2375, 2751, 2288, 1932, 2166, 2217,
    3482, 2036, 2581, 2639)
Y=(68, 74,73,53,59,60,57,47, 46, 50,55, 74,41,71, 67)
```

In a study of 15 randomly selected countries, using the variables ‘daily calorie intake’ X and ‘life expectancy’ Y of men, is there a positive correlation between calorie intake and life expectancy?

Hint: the test t statistics is

```
n = length(X)
r = corrcoef(X,Y)
t = abs(r) * sqrt(n-2)/sqrt(1 - r^2)
```

Exercise 236. (another check)

Verify that the `corrcoef` of $X = (1, 3, 4, 6, 8, 9, 11, 14)$ and $Y = (1, 2, 4, 4, 5, 7, 8, 9)$ is 0.977.

Exercise 237. (Check SPIEGEL example 14.12, cf. [18, p.361])

We have $X = (20, 5, 8, 10, 13, 7, 13, 5, 25, 14)$

and $Y = (2.35, 3.8, 3.5, 2.75, 3.25, 3.4, 2.9, 3.5, 2.25, 2.75)$

with $r = 0.9097$.

Exercise 238. (example 5.1 with $\mathcal{R}$ and Octave)

◦ Check the solution of the example 5.1 using `cor` function of $\mathcal{R}$.

```
R> X = c(60,65,72,70,74,63,66,68,67,69,70,64)
> Y = c(120,140,180,184,190,160,155,158,170,190,210,150)
> cor(X, Y, method = 'pearson')
[1] 0.8359452
```

◦ Check the solution using `corrcoef()` function of Octave:

```
octave:1> X = [60,65,72,70,74,63,66,68,67,69,70,64];
octave:2> Y = [120,140,180,184,190,160,155,158,170,190,210,150];
octave:3> corrcoef(X,Y)
      ans = 1.0000    0.8359
           0.8359    1.0000
```

Remark: (MatLAB) For two input arguments, result is a 2-by-2 matrix with ones along the diagonal and the correlation coefficients along the off-diagonal.

5.2 SPEARMAN'S ρ_S rank correlation coefficient

SPEARMAN's rank correlation coefficient ρ_S is a number ranging from -1 to 1 that indicates how strongly two *sets of ranks*(!) are correlated. The difference between the PEARSON correlation ρ and the SPEARMAN ρ_S correlation is that the PEARSON is most appropriate for measurements taken from an interval scale, while the SPEARMAN is more appropriate for measurements taken from ordinal scales.

Mental image




Figure 46: Visualization of SPEARMAN's rank coeff. ρ_S of two samples X and Y
Left: For roughly elliptically distributed data with no prominent outliers, Spearman's ρ_S and Pearson's ρ correlation give similar values.
Right: The Spearman correlation is less sensitive than Pearson's to strong outliers, because ρ_S limits the outlier to the value of its rank.

Figures and comments are cited from ▷ WIKIPEDIA : Spearman's ..

Procedure SPEARMAN's rank correlation coefficient

1. **Assumptions** two independent random samples $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$
2. **correlation** ρ_S rank coefficient by SPEARMAN is given by

$$\rho_S(X, Y) := \frac{\text{cov}(cX, cY)}{\sigma_{cX} \cdot \sigma_{cY}}$$

where cX resp. cY are the ranks of X and Y corresponding to their sample values.

Remark. If X and Y have no 'ties', then one uses the simpler formula

$$\rho_S := 1 - \frac{6 \cdot \sum_{i=1}^n d_i^2}{n \cdot (n^2 - 1)}$$

where $d_i := \text{Rank}(x_i) - \text{Rank}(y_i)$.

Example *Example from exceldemy*

We do the example by R.R. SUPROV in ▷ exceldemy with their data sets.
Calculate SPEARMAN's ρ_S in two ways using both formulas from above.

Solution along the columns of the 'rank' table:

Student Name	Math	Economics	R_{math}	$R_{\text{economics}}$
Ali	70	90	10	3
Beatriz	78	94	8	1
Charles	90	79	2	8
Diya	87	86	3	5
Eric	84	84	6	6
Fatima	86	83	4	7
Gabriel	91	88	1	4
Hanna	74	92	9	2
Rodriguez	83	76	7	9
Robert	85	75	5	10

$$1. R_{\text{math}} \equiv cX \text{ and } R_{\text{econ}} \equiv cY$$

$$2. \text{cov}(cX, cY) = -3.45.$$

$$3. \sigma_{cX} = 2.8722 \text{ and } \sigma_{cY} = 2.8722.$$

$$4. \rho_S = \frac{\text{cov}(cX, cY)}{\sigma_{cX} \cdot \sigma_{cY}} = -0.4181.$$

Negative correlation: As X variable increases, the Y tends to decrease.

◦ *Solution* and explanation of this example with REDUCE in ▷ Spearman.

```
| X := {70,78,90,87,84,86,91,74,83,85}$ -- math
| Y := {120,140,180,184,190,160,155,158,170,190,210,150}$ -- econ
| spearman(X,Y); | -0.4182
```



View step-by-step solution *code*, run REDUCE.



View Spearman's ρ in action *code*, run REDUCE.

General Information

General mathematical information about the concept is here ▷ WIKIPEDIA : spearman
Syntax and semantic of the implementation is here ▷ MATLAB : ...spearman

5.2.1 Exercises

Exercise 239. (redo example.5.2 using OCTAVE)Check the following solution using `spearman` function of `statistics` package of Octave.

```
octave:1> pkg load statistics
octave:2> X = [70,78,90,87,84,86,91,74,83,85];
octave:3> Y = [90,94,79,86,84,83,88,92,76,75];
octave:4> spearman(X, Y)
      ans = -0.4182
```

Exercise 240. (the Father–Son–example)

The heights of a sample of 12 fathers and their oldest sons are measured giving the following data, cf. SPIEGEL [18, p. 434, P 17.24]

```
Far : 65,63,67,64,68,62,70,66,68,67,69,71
Son : 68,66,68,65,69,66,68,65,71,67,68,70
```

- Calculate the Spearman rank correlation ρ .
- Check the solution of the example using $\mathcal{R}$.

```
R> Far = c(65,63,67,64,68,62,70,66,68,67,69,71)
R> Son = c(68,66,68,65,69,66,68,65,71,67,68,70)
R> cor.test(Far, Son, method = "spearman")
```

```
Spearman's rank correlation rho
data: Far and Son
S = 74.285, p-value = 0.005903
alternative hypothesis: true rho is not equal to 0
sample estimates:
rho = 0.7402623
```

View *code*, run REDUCE.**Exercise 241.** (significance)

Use the FACT

The test for significance of SPEARMAN's ρ_S use the statistics


$$T = |\rho_S| \cdot \frac{\sqrt{n-2}}{\sqrt{1-\rho_S^2}}$$

which is t -distributed with $df = n - 2$.and do a t test for the significance of the ρ_S .Result: $T = 2.905$

5.3 KENDALL'S τ rank correlation coefficient

KENDALL'S τ is a non-parametric rank correlation coefficient that measures the similarity between two rankings by assessing the number of concordant and discordant pairs. It ranges from -1 to $+1$, where $+1$ indicates identical rankings, -1 indicates the reverse of the other, and 0 indicates no relationship. $\triangleright$ KENDALL.

Mental image



Visualization of KENDALL'S rank `corrcoeff` τ of two samples X, Y .

All points in the gray area are *concordant* and all points in the white area are *discordant* with respect to point (X_1, Y_1) . With $n = 30$ points,

Figure 47: there are a total of $\binom{30}{2} = 435$ possible point pairs. In this example there are 395 concordant point pairs and 40 discordant point pairs, leading to a Kendall rank correlation coefficient of $\tau = 0.816$.

Figures and comments are cited from $\triangleright$ WIKIPEDIA : Kendall's ..

Procedure KENDALL'S rank correlation coefficient τ

1. **Assumptions** two random samples $X = (x_1, \dots, x_n)$ and $Y = (y_1, \dots, y_n)$
2. **correlation** τ rank coefficient by KENDALL is calculated by these steps:
 - a. The data for each sample X and Y is first converted into ranks..
 - b. Then examine every possible pair of observations to determine if the ranks of the pair are in the same order ('*concordant*') or a different order ('*discordant*').
 - c. The final value τ is based on the difference between the number C of concordant pairs and the number D of discordant pairs:

$$\tau := \frac{C - D}{C + D}$$

Positive correlation: As X variable increases, the Y tends also to increase.

General Information

- General mathematical information about the concept is here $\triangleright$ WIKIPEDIA : Kendall
- Syntax and semantic of the implementation is here $\triangleright$ MATLAB : `...corr()`

Example *Example : two doctors rank 6 patients.*

We do the example from $\triangleright$ *numiqo*. – Suppose two doctors rank 6 patients by descending physical health. One of the two doctors, in this case the female, is now defined as the reference (*doc1*) and the patients are sorted from 1 to 6.

Calculate KENDALL's τ using the formula from above.

Solution along the columns of the 'rank' table

in REDUCE:

	1	3			
1	1	3			
2	2	1	-		
3	3	4	+	+	
4	4	2	-	+	-
5	5	6	+	+	+
6	6	5	+	+	+

doc1:	doc2:	p:	q:
1	3	3	2
2	1	3	1
3	4	2	1
4	2	2	0
5	6	1	0
6	5	—	—
		11	4

$$\tau = \frac{7}{15}$$

1st column "-+-++" $\uparrow$ gives ...

... $(p, q) = (3, 2)$ in row 2 of matrix.

1. Why $p = 3$? – Because below $doc2 = 3$ are 3 items greater than $doc2 = 3$: 4, 6, 5.
2. Why $q = 2$? – Because below $doc2 = 3$ there are 2 items smaller than $doc2 = 3$: 1, 2.
3. $C := \Sigma_p = 3 + 3 + 2 + 2 + 1 = 11$. $D := \Sigma_q = 2 + 1 + 1 + 0 + 0 = 4$.
4. $\tau = \frac{C-D}{C+D} = \frac{7}{15} = 0.4666$.

Positive correlation: As *doc1* variable increases, the *doc2* tends also to increase.

- *Solution* and explanation of this example with REDUCE

```

| X := {1,2,3,4,5,6}$
| Y := {3,1,4,2,6,5}$
| kendall(X,Y);

```



View step-by-step solution *code*, run REDUCE.



View Kendall's tau τ in action *code*, run REDUCE.

5.3.1 Exercises

Exercise 242. (example.5.3.1 using $\mathcal{R}$)

- Check the following solution using `cor(..., "kendall")` function of R.

```
R> X=c(1,2,3,4,5,6)
R> Y=c(3,1,4,2,6,5)
R> cor(X,Y, method="kendall")
[1] 0.4666667
```

Exercise 243. (league standings by [29, p.129ff])

The following table compares the league standings X of the clubs with their stadium sizes Y . The aim is to check whether stadium size correlates with league position.

```
X = (1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18)
Y = (30,72,65,80,72,25,41,45,59,15,70,23,58,24,64,22,26,28) -- in 10^3
```

[We cite ZOEFEL, a.a.o.] According to KENDALL, the first step is to write down the ranking of the first variable X in ascending order and assign the other accordingly. In this context, one speaks of an anchor row X and a comparison row Y . In the example, the league position already noted in ascending order is the anchor row. The comparison row Y consists of the outermost positions in terms of stadium size. It is that the largest stadium is assigned to rank no 1.

With a positive correlation, the comparison series is expected to rise in a similarly monotonic manner as the anchor series. The number of disturbances ('inversion' w.r.t to this monotonicity) represents a progression for the strength of the correlation, since such a disturbance occurs when a lower ranking follows a higher ranking in the comparison series. E.g. this happens ten times for the club in first place and three times for the club in third place.

The number of inversions is entered in the column labeled I , the number of proversions is entered in the column labeled P , meaning that in the comparison series, a single-place entry follows – the higher one.

Now the sums of I and P are calculated and this gives the KENDALL τ as $\tau := \sum P - \sum I$

- Follow the algorithm by hand – but REDUCE is your friend.
- Make a 'video' a la example.5.3
- Verify

$$\tau = \frac{99 - 53}{99 + 53}$$

- Verify using $\mathcal{R}$: `cor(X,Y, method="kendall")` gives [1] -0.301641.

Exercise 244. (significance of Kendall- τ)

Use the FACT

The test for significance of KENDALL τ use the statistics

$$Z = |\tau| \sqrt{\frac{2 \cdot (2 \cdot n + 5)}{9 \cdot n \cdot (n - 1)}}$$

which is std.normal-distributed.

and do an F test for the significance of the τ .

Solution should be 1.74 for the ZOEFEL example.

5.4 ICC - Intraclass Correlation Coefficient

The intraclass correlation coefficient ICC is a descriptive statistic that measures how strongly units within the same group are similar to each other. It is used for quantitative measurements organized into groups and is often used to assess agreement between multiple raters or measurements, with a value between 0 (no agreement) and 1 (perfect agreement). An ICC helps determine reliability, like in cases where two raters evaluate the same set of patients or a single rater measures the same subjects multiple times. ▷ *Google*

Procedure `icc` *correlation coefficient*

1. **Assumptions** k correlating variables in n cases with observation matrix $X = (x_{ij})_{i=1..n}^{j=1..k}$
2. **correlation** `icc` coefficient is given by

$$icc(X) := \frac{MQ_{row} - MQ_{rest}}{MQ_{row} + (k - 1) \cdot MQ_{rest}}$$

where we use the following 9 terms

$$\begin{array}{lll} T_j := \sum_{i=1}^k x_{ij} & S := \sum_{j=1}^k T_j & SAQ_{total} := \sum_{i=1}^k \sum_{j=1}^n x_{ij}^2 - \frac{S^2}{k \cdot n} \\ SAQ_{row} := \frac{1}{k} \cdot \sum_{j=1}^n T_j^2 - \frac{S^2}{k \cdot n} & df_{row} := n - 1 & MQ_{row} := \frac{SAQ_{row}}{df_{row}} \\ SAQ_{rest} := SAQ_{total} - SAQ_{row} & df_{rest} := n \cdot (k - 1) & MQ_{rest} := \frac{SAQ_{rest}}{df_{rest}} \end{array}$$

Example *Example and text from ZOEFEL, [29, p.133 ff]*

Fourteen women were asked about their weight; the given answers were then compared to the actual measured weight. The measured values yield a mean of 57.6 *kg*, while the mean of the estimated values is 54.6 *kg*. The actual weight values are therefore significantly underestimated. The `icc` attempts to combine both aspects, the correlation and the differences in the means, into a single measure. `icc` only achieves high values, when both the direction and the levels of the variables involved match. We have the data

$$X = \left(\begin{array}{cccccccccccccc} 48 & 48 & 50 & 50 & 52 & 58 & 63 & 56 & 48 & 63 & 58 & 58 & 52 & 60 \\ 51 & 50 & 50 & 52 & 53 & 63 & 70 & 70 & 49 & 63 & 59 & 58 & 56 & 62 \end{array} \right) \begin{array}{l} \text{weight estimated} \\ \text{weight measured} \end{array}$$

Calculate the `icc` of X using the formulas from above.

Solution along the 9 constituting terms of the `icc`:

1. $k = 2$ and $n = 14$ and $T_1 = 99, T_2 = 98, \dots$ and $S = 1570$.
2. $SAQ_{total} = 1131.9$.
3. $SAQ_{row} = 976.9$ and $df_{row} = 13$ and $MQ_{row} = 75.1$.
4. $SAQ_{rest} = 155$ and $df_{rest} = 14$ and $MQ_{rest} = 11.07$.

5. $\text{icc}(X) = 0.743$.

◦ *Solution* and explanation of this example with REDUCE :

```
| X := {48,48,50,50,52,58,63,56,48,63,58,58,52,60}$
| Y := {51,50,50,52,53,63,70,70,49,63,59,58,56,62}$
| icc(X,Y); | 0.743
```



View step-by-step solution *code*, run REDUCE.



View `icc` in action *code*, run REDUCE.

General Information

General mathematical information about the concept ▷ WIKIPEDIA : ICC

Syntax and semantic of the implementation ▷ R : `icc`

5.4.1 Exercises

Exercise 245. (early ICC after PEARSON/FISHER)

Program the ‘early ICC formula’ $\triangleright$ `iccOld` in REDUCE and solve our example.5.4 with it. Compare both solutions.

You may use the following pseudocode:

```
### ----- ICCp -----
###   iccP(X,Y)
###   method = Intraclass Correlation after PEARSON/FISHER
###   .....
iccP(X,Y, Z,T,S,n) = do(
  n = dim(X),
  Xbar = 1/(2*n)*sum(i,1,n, X[i]+Y[i]),
  sigma2 = 1/(2*n)*( sum(i,1,n, (X[i]-Xbar)^2) +
                    sum(i,1,n, (Y[i]-Xbar)^2) ),
  r = 1/(n*sigma2)*sum(i,1,n, (X[i]-Xbar)*(Y[i]-Xbar)),
  float(r))
```

Exercise 246. (significance of ICC)

Use the fact, that the test for significance of ICC use the statistics

$$F = MQ_{row}/MQ_{rest}$$

which is F -distributed with $(n-1; n \cdot (k-1))$ degrees of freedom.

Do an F test for the significance of the ICC of example.5.4.

Exercise 247. (Check example.5.4 with $\mathcal{R}$)

Check the solution of example.5.4 using package `icc` of $\mathcal{R}$.

```
R> install.packages('irr')
R> library('irr')
R> data <- data.frame( X = c(48,48,50,50,52,58,63,56,48,63,58,58,52,60),
                      Y = c(51,50,50,52,53,63,70,70,49,63,59,58,56,62))
R> icc(data, model = "oneway", type = "agreement", unit = "single")
```

```
Single Score Intraclass Correlation
Model: oneway   Type : agreement
Subjects = 14   Raters = 2
ICC(1) = 0.743
```

```
F-Test, H0: r0 = 0 ; H1: r0 > 0; F(13,14) = 6.79 , p = 0.00053
95%-Confidence Interval for ICC Population Values: 0.385 < ICC < 0.909
```

5.5 regression - the linear regression line

Simple Linear regression models the relation between a dependent, or response, variable y and an independent, or predictor, variable x and considers the independent variable using the relation $y = ax + b$, where b is the y -intercept, and a is the slope (or regression coefficient), i.e. there is a linear relation between x and y ($x \sim y$).

Mental image




Figure 48: Visualization of linear regression (line):
 electrical resistance of a conductor as a function of temperature.
 —: regression line $y = 0.28 \cdot x + 70.8$ of the measurement $X \mapsto Y$ with
 ○: $(19, 25, 30, 36, 40, 45, 50) \mapsto (76.3, 77.8, 79.8, 80.8, 82.4, 83.9, 85.1)$.

Procedure *linear regression line $y = ax + b$*

1. **Assumptions** a 2-dimensional sample (X, Y) with a linear connection $y_i \approx a \cdot x_i + b$.
2. **regression line** coefficients a and b are calculated by

$$a := \frac{\sum_{i=1}^n (x_i - \bar{X}) \cdot (y_i - \bar{Y})}{\sum_{i=1}^n (x_i - \bar{X})^2}$$

$$b := \frac{1}{n} \cdot (\bar{Y} - a \cdot \bar{X})$$

where $\bar{X}$ resp. $\bar{Y}$ are the mean values of $X = (x_1, \dots, x_n)$ resp. $Y = (y_1, \dots, y_n)$.

Example *electrical resistance of a conductor as a function of temperature*

We do the example of figure.48, cf. [7, p.193]. We have the electrical resistance (Y) of a conductor as a function of temperature (X):

Temperature:	19	25	30	36	40	45	50
Resistance:	76.3	77.8	79.8	80.8	82.4	83.9	85.1

Calculate the equation $y = ax + b$ from the measurement.

Solution

1. $n = 7$.
2. $\text{mean}X = 35$ and $\text{mean}Y = 80.87$.
3. numerator of $a = \text{cov}(X, Y) = 210.5$
4. denominator of $a = \Sigma(X - \text{mean}X)^2 = 732$
5. $a = \frac{\text{numerator}}{\text{denominator}} = 0.287$
6. $b = 70.80$

◦ *Solution* and explanation of this example with REDUCE in

```
| X := {19,25,30,36,40,45,50}$
| Y := {76.3, 77.8, 79.8, 80.8, 82.4, 83.9, 85.1}$
| regression(X,Y);
```

a	b
0.28757	70.807



View step-by-step solution *code*, run REDUCE.



View regression line in action *code*, run REDUCE.

General Information

General mathematical information is here ▷ WIKIPEDIA : Linear regression
 Syntax and semantic of the implementation is here ▷ R : lm

5.5.1 Exercises

Exercise 248. (example from [18, p.361, P 14.12])

Is there a connection between grade point average (GPA) and hours of TV watched per week? There is a sample of 10 high school students:

```
TVhours = (20,5,8,10,13,7,13,5,25,14)
GPA      = (2.35, 3.8, 3.5, 2.75, 3.25, 3.4, 2.9, 3.5, 2.25, 2.75)
```

Result: The variable TVhours and GPA are negatively correlated. i.e. as more hours of TV are watched, the lower is the GPA :(

Exercise 249. (check example.5.5 with $\mathcal{R}$)

Check the solution using `lm` function of R:

```
R> temperature <- c(19,25,30,36,40,45,50)
R> resistance  <- c(76.3, 77.8, 79.8, 80.8, 82.4, 83.9, 85.1)
R> lm(resistance ~ temperature)
R> plot(temperature, resistance)+abline(lm(resistance ~ temperature))
```

Call:

```
lm(formula = resistance ~ temperature)
```


Coefficients:

```
(Intercept)  temperature
70.8065      0.2876
```

5.6 anova1 - One-way Analysis Of Variance

One-way ANOVA is a statistical test used to compare the means of three or more independent groups to determine if there is a statistically significant difference between them. It works by partitioning the total variability in the data into two sources: the variability *between* the groups and the variability *within* the groups. If the variability between groups is large compared to the variability within groups, it suggests that the group means are significantly different.

Mental image



Visualization of *One-way Analysis Of Variance* test of three groups.

A: A solid black square $\blacksquare$ is suggested as a general representative value such as mean of overall data.

B: It looks reasonable to divide the data into three groups and explain the data with three different means of groups: $\bullet$, $\star$, $\triangle$.

C: To evaluate the efficiency or validity of dividing three groups, the distances from group means to overall mean and the distances from group means to each data are compared. Distance between group means and overall mean (solid arrows) stands for the inter-group variance and distance between group means and each group data (dotted arrows) stands for the intra-group variances. — Figures and comments are cited from

T.K. KIM ▷ NIH : ANOVA

Procedure *One-way Analysis Of Variance*

1. **Assumptions** a data matrix ('table') $X = [x_{ij}]_{j=1..a}^{i=1..b}$, where x_{ij} denotes the observation no. j in group i . The rows of the data matrix are the 'treatments'.
2. **Null hypothesis** H_0 : the variances of all groups are equal (two-sided)
3. **Test statistics** $\text{anova1}(X)$ is F -distributed with $df_1 = a - 1$ and $df_2 = a \cdot (b - 1)$ degrees of freedom and is given by

$$F := \frac{MS_B}{MS_W} = \text{anova1}(X)$$

where we use the following ‘ANOVA1 table’ to calculate F

calculate from left to right:	Variation	df	Mean square	F
Between treatments	$V_B := \frac{1}{b} \cdot \sum_j T_j^2 - \frac{T^2}{ab}$	$a - 1$	$MS_B := \frac{V_B}{a-1}$	$F := \frac{SS_B}{SS_W}$
Within treatments	$V_W := V - V_B$	$a \cdot (b - 1)$	$MS_W := \frac{V_W}{a(b-1)}$	
Total treatments	$V := \sum_{j,k} x_{jk}^2 - \frac{T^2}{ab}$	$ab - 1$		ANOVA1

and the helper terms $T := \sum_{i,j} x_{ij}$ = the total sum of all values in X
and $T_j := \sum_{i=1}^a x_{ij}$ = the total sum of all values in the j th treatment (row) of X .

4. **Decision** Reject H_0 , if $\chi > \text{chiINV}(1; 1 - \alpha)$.

Example *Yields in bushels per acre*

We do the example 16.4 of M. SPIEGEL [18, p.415]. The table shows the yields in bushels per acre of a certain variety of wheat grown in a particular type of soil treated with chemicals A,B and C. Do an ANOVA1 analysis of the treatments.

A:	3	4	5	4
B:	2	4	3	3
C:	4	6	5	5

◦ *Solution* and explanation of this example with REDUCE :

```
| X :={{3,4,5,4},{2,4,3,3},{4,6,5,5}}$
| anova1(X);
```

Source	df	SS	MS	F
Between	2	8	4.0	6.0
Within	9	6	0.66667	-
Total	11	14	-	ANOVA



View step-by-step solution *code*, run REDUCE.



View ANOVA1 in action *code*, run REDUCE.

General Information

- General mathematical information is here ▷ WIKIPEDIA : One-way analysis of variance
- Syntax and semantic of the implementation is here ▷ MATLAB : `anova1`

5.6.1 Exercises

Exercise 250. (mineral water company, cf. [29, p.179])

A mineral water company claims that its mineral water lowers cholesterol levels. Eight test subjects drank this water for three weeks. The values at the beginning of the trial, as well as after one week, two weeks, and three weeks, are included in the following table.

```
A =(267,248,321,272,355,264,270,252)  -- values at begin of test
W1 =(238,232,307,295,348,260,266,249)  -- after 1 week
W2 =(191,246,295,270,330,262,295,220)  -- after 2 weeks
W3 =(206,207,282,269,275,281,263,219)  -- after 3 weeks
```

The aim is to determine whether these mean values differ significantly from one another, i.e., whether the mineral water actually has a significant effect on cholesterol levels. These are dependent samples.

Result: According to the F-table, this is a significant value at (3, 21) degrees of freedom ($p < 0.05$). Therefore, a significant decrease in cholesterol levels is observed over the time period.

Exercise 251. (blood pressure from numiqo, cf. anova)

Do the example in numiqo.

Exercise 252. (salaries of people from excel-easy, cf. salary)

Do the example in excel-easy.

Exercise 253. (Example from statsdirect, cf. statsdirect)

Do the example in www.statsdirect.com: *analysis_of_variance/one_way*.

Exercise 254. (Example from geeksforgeeks, cf. anova1)

Do the example in geeksforgeeks.org.

Exercise 255. (Example from kent.edu, cf. signtest)

Do the example in kent.edu : SPSS/OneWayANOVA.

Exercise 256. (Example from statsandr.com, cf. anova-by-hand)

Do the example in statsandr.com: *how-to-one-way-anova-by-hand*.

Exercise 257. (check with $\mathcal{R}$)

Check the solution of the example.5.6 using using `anova1` function of octave.


```
octave:1> pkg load statistics
octave:2> X=[3,4,5,4; 2,4,3,3; 4,6,5,5]
octave:3> anova1(X)
```

ANOVA Table					
Source	SS	df	MS	F	Prob>F
Groups	8.0000	2	4.0000	6.00	0.0221 = p value
Error	6.0000	9	0.6667		
Total	14.0000	11			

5.7 boot1 - the bootstrap method for dependent samples

The purpose of bootstrapping in statistics is to estimate the sampling distribution of a statistic by repeatedly resampling from the original data. This allows statisticians to perform hypothesis testing, calculate standard errors, and construct confidence intervals without needing to make strong assumptions about the underlying data distribution, a task that can be complex or impossible with traditional methods.

Mental image



Visualization of **bootstrap**.

Left: From sample $X = \{4, 5, 4, 7, 2\}$ with mean $\bar{X} = 4.4$ we draw a new ('bootstrapped') sample giving $X' = \{4, 2, 7, 5, 7\}$ with mean $\bar{X}' = 5$.

Middle: From the original sample X we draw ('resample') a new sample $X'' = \{4, 5, 5, 2, 2\}$ mean $\bar{X}'' = 3.6$

Right: From the original sample X we resample again a new sample $X''' = \{2, 4, 5, 4, 7\}$ with mean $\bar{X}''' = 4.4$.

From the resamples, the statistic T (here: *mean*) is calculated with $T = \text{mean}(\bar{X}', \bar{X}'', \bar{X}''') = 4.33$ to estimate the distribution of T .

▷ Figures are cited from ▷ G.J. SCOTT: The bootstrap

Procedure *the bootstrap*

1. **Assumptions** a data set (sample) $X = (x_1, \dots, x_n)$
2. **bootstrap** The bootstrap method is summarized by the following steps:
 - Choose a statistic T to be studied, e.g. $T := \text{mean}$.
 - a. Choose the number n of bootstrap samples to take.
 - b. For each bootstrap sample, draw a replacement sample X_i of size n .
 - c. Calculate the statistics T_i for the samples X_i .
 - d. Find a summary statistic T ('bootstrap statistic') for each of the n samples T_i .

Example *bootstrap with test statistic sd*

We follow the toy example of figure.50. Consider the sample $X = \{4, 5, 4, 7, 2\}$. Using the bootstrap samples X_1, X_2, X_3 from above, estimate the standard deviation of the bootstrap distribution.

Solution We follow the recipe bootstrap.

1. Choose $T = sd = \frac{\sqrt{\sum (X - \bar{X})^2}}{\sqrt{length(X) - 1}}$ with $length(X) = 5$.
2. $n = 3$
3. $T_1 := sd(X_1) = 2.1213, T_2 := sd(X_2) = 1.5165, T_3 := sd(X_3) = 1.8165$.
4. $T = mean(T_1, T_2, T_3) = mean(2.1213, 1.5165, 1.8165) = 1.8181$

◦ *Solution* and explanation of this example with REDUCE

```
| X := {4,5,4,7,2}$
| boot1(X,1000);
```

est.mean	est.sd
4.476	0.51367



View step-by-step solution *code*, run REDUCE.



View **boot1** in action *code*, run REDUCE.

General Information

General mathematical information is here ▷ WIKIPEDIA : Bootstrapping
 Syntax and semantic of the implementation is here ▷ MATLAB : **bootstrp**

5.7.1 Exercises

Exercise 258. (Marihuana, cf. [7, p.205])

“From a given sample, any number of further samples can be drawn using random numbers. Compare the number of random samples that meet a given condition with the total number of samples drawn. This allows a statement to be made about the significance level of the test.”

The following sample of memory test scores after smoking a regular cigarette and a Marijuana cigarette is considered.

```
person: 1  2  3  4  5  6  7  8  9
cigar: -3 10 -3  3  4 -3  2 -1 -1
marih.: 5 -17 -7 -3 -7 -9 -6  1 -3
```

The related samples are combinatorially interpreted as follows: First, the memory performances X after smoking a regular cigarette and Y after smoking a marijuana cigarette are combined into a new sample of 18 elements.

In our example, after inputting a sample of size n , 1000 pairs of samples of size $n/2$ are drawn, a counter counts the number of sample pairs whose sums S resp. T satisfy the condition $|T - S| \geq |D|$ where $D := \sum (X_i - Y_i)$.

Use the pseudocode to program a adapted bootstrap method

```
...
z = 0,                -- counter
n = dim(X),          -- number of samples
k = n/2,
...
    D = abs(sum(T-S)),
    test(abs(T[j]-S[j]>D), z=z+1) ),
estT = mean(T), -- bootstrap sample means
estS = mean(S), -- i.e. estimated values
p=0.5*z/IT,      -- calculate probability
print(("est.mean", "est.sd","p"),(estT, estS, p ))
...
```

Result: For the marijuana example dataset, the program randomly returns values around 0.015 as the significance level. The null hypothesis, that memory performance after smoking ordinary and marijuana cigarettes differs only by chance, can be rejected with 1.6% probability of error.

Exercise 259. (Bootstrap Samples of Observations from MATLAB, cf. bootstrp)


Do two bootstrap examples from MATLAB.

- Take bootstrap samples of patient data, compute the mean measurements for each data sample, and visualize the results.
- Create 50 bootstrap samples from the numbers 1 through 6. To create each sample, bootstrp randomly chooses with replacement from the numbers 1 through 6, six times. This process is similar to rolling a die six times. For each sample, the custom function *countfun* (to be programed by you) counts the number of 1s in the sample.

5.8 boot2 - the bootstrap method for independent samples

The bootstrap method for independent samples X and Y is a resampling technique that estimates the sampling distribution of a statistic θ ($\stackrel{e.g.}{=} \text{mean}$) without relying on strong assumptions like normality. It involves repeatedly drawing samples with replacement from each of the two original, independent samples X and Y , calculating the statistic of interest like the *difference in means for each resampled pair*, and using the resulting distribution of these statistics e.g. to calculate the estimated difference in means for X and Y .

Mental image



Visualization of **bootstrap2**. We generate new samples Z^{*1}, Z^{*2}, Z^{*B} directly from the population observations (leftmost histogram; *obs*).

If we choose $\theta \stackrel{p.d.}{=} \text{mean}$; then $\bar{X} = 3.83$ and $\bar{Y} = 2.1$.

Figure 51: We generate the distribution of sample means B -times, e.g. $B = 3$.

Z^{*1} : $\theta^{*X1} = \bar{X}^1 = \text{mean}(5.3, 4.3, 5.3) = 4.96$ and $\theta^{*Y1} = \bar{Y}^1 = 2.33$.

Z^{*2} : $\theta^{*X2} = \bar{X}^2 = 3.9$ and $\theta^{*Y2} = \bar{Y}^2 = 2.1$.

Z^{*B} : $\theta^{*XB} = \bar{X}^B = 2.83$ and $\theta^{*YB} = \bar{Y}^B = 1.53$.

Procedure *bootstrap2 method for independent samples*

1. **Assumptions** two data sets (samples) $X = (x_1, \dots, x_m)$ and $Y = (y_1, \dots, y_n)$.
2. **bootstrap2** The bootstrap2 method is summarized by the following steps:
 - Choose a statistic θ to be studied resp. estimated, e.g. $\theta := \text{mean}$.³¹
 - a. Choose the number B of bootstrap samples to take.
 - b. *Resample with replacement*: Create new ‘bootstrap’ samples by drawing observations with replacement from each of your original, independent samples.³²

³¹For independent samples, this is often the difference between the two sample means.

³²For example, if you have two samples, one of size m and one of size n , you will create b pairs of new samples, each with size m and n respectively.

- c. *Calculate the statistic:* For each of the B pairs of bootstrap samples, calculate the statistic θ_i of interest. This gives you a collection of B values for your statistic.
- d. *Use the bootstrap distribution:* The B calculated statistics form a *bootstrap distribution* approximating the true sampling distribution of the statistic θ .
- e. *Summarize the distribution:* Use the bootstrap distribution to estimate the difference in the means..

Example *the bootstrap mean of figure.51*

We follow the toy example of figure.51.

Calculate the bootstrap statistics θ^{*X} and θ^{*Y} and their difference.

Solution

1. The Bootstrap stats: $\theta^{*X} = \text{mean}(\theta^{*X1}, \theta^{*X2}, \theta^{*X3}) = 3.90$
2. $\theta^{*Y} = 1.98$
3. The difference is $d = \theta^{*X} - \theta^{*Y} = 1.92$, ergo X and Y are independent.

◦ *Solution* and explanation of this example with REDUCE

```
| X := {4.3, 2.1, 5.3}$
| Y := {2.4, 1.1, 2.8}$
| boot2(X,Y, 1000);
```

e.meanX	e.meanY	significance
1.8667	4.1333	0.055



View `boot2` in action *code*, run REDUCE.

General Information

General mathematical information is here ▷ WIKIPEDIA : Bootstrapping
 Syntax and semantic of the implementation is here ▷ MATLAB : `bootstrp`
 and ▷ R : `bootstrap`

5.8.1 Exercises

Exercise 260. (31 US presidents, cf. [7, p.207])

There was a random sample of 31 US presidents drawn, who died of natural causes. The whole sample was divided in tall (X) and small (Y) persons.

X = (85,79,90,67,80)

Y = (67,56,53,77,88,57,63,78,70,
66,83,67,60,74,65,64,73,71,
60,90,63,71,64,78,68,72)

The null hypothesis is: age and height are independent.
Use `boot2` to test the hypothesis.

Solution:  View 2nd example in *boot2*, run REDUCE.

Result: We get values around 0.013. With a 5% probability of error, the null hypothesis, that age and size differ only by chance, can be rejected.

Exercise 261. (Example from MATLAB, cf. `bootstrap`)

Do an example from the MATLAB documentation in REDUCE using the data from section ‘Medians of Paired Samples’.

5.9 bootCI - Confidence Interval using bootstrap

Bootstrapping is a statistical resampling technique that uses an existing dataset to estimate the properties of an estimator, such as constructing a confidence interval. These ‘*bootstrap samples*’ are then used to build an approximate sampling distribution for the statistic of interest, which helps in situations where analytical methods are complex or difficult to apply. For example, *to find a confidence interval*, you take an initial sample, then repeatedly create new bootstrap samples of the same size by randomly drawing from your original sample with replacement. You calculate the mean for each of these new samples and use the distribution of these means to find the confidence interval, such as the 2.5th and 97.5th percentiles for a 95% confidence interval.

Mental image




Figure 52: Visualization of **bootCI**: Instead of taking multiple samples directly from the whole employees population, we only draw a single, representative *sample* (5emp). We then generate e.g. 20 ‘new’ samples by *repeatedly* creating new bootstrap samples of the same size.

Picture found at [▷ ASKPYPHON : bootstrap-sampling](#)

Procedure **bootstrap confidence interval**

1. Choose a statistics T to be studied, e.g. $T := \text{mean}$.
2. Choose the number n of bootstrap samples to take.
3. For $i = 1, \dots, n$, draw a replacement sample X_i (‘bootstrap sample’) of size n .
4. Calculate the statistics T_i for the samples X_i , e.g. $T_i := \text{mean}(X_i)$.
5. For the k th percentile with $k \in \{1, \dots, 99\}$, the k percent confidence interval for the summary statistic $T := T_1, \dots, T_n$ ³³ of the n resamples T_i is defined by:

$$CI_k := \left(\text{percentile}(T, \frac{100 - k}{2}), \text{percentile}(T, \frac{100 + k}{2}) \right)$$

³³the ‘bootstrap statistic’, e.g. the vector of the bootstrap means T_i

Example The ‘employees’ example from *askpython*

We follow the example of Figure.52 by hand on the toy example by A. YADAV for the fictive ages sample $emp = (25, 30, 35, 40, 45, 50, 55, 60, 65, 70)$. Calculate the 95% confidence interval for the mean of emp .

Solution We do only $n = 3$ resamples to show the principle of the procedure.

1. $X := (25, 30, 35, 40, 45, 50, 55, 60, 65, 70) = emp$, $n = 3$. So we iterate 3-times:
2. $X_1 = (40, 60, 40, 35, 55, 65, 70, 45, 35, 25)$, $T_1 = mean(X_1) = 47$.
3. $X_2 = (70, 60, 40, 55, 55, 40, 50, 25, 50, 55)$, $T_2 = 50$.
4. $X_3 = (70, 70, 50, 70, 30, 50, 30, 30, 55, 60)$, $T_3 = 51.5$.
5. $T(X) := mean(T_1, T_2, T_3) = mean(47, 50, 51.5) = 51.16$.
6. $CI_{95} = (47; 51.5) = (CI.l, CI.h)$ – this interval catch 51.16 with no surprise.

◦ *Solution* and explanation of this example with REDUCE.

We check the solution using our function `bootci` and 100 resamples:

```
| X := {25, 30, 35, 40, 45, 50, 55, 60, 65, 70}$
| bootci(X,100);
```

est.Mean	CI.l	CI.h
48.135	39.5	50.5



View step-by-step solution *code*, run REDUCE.



View `bootCI` in action *code*, run REDUCE.

General Information

Syntax and semantic of the implementation is here `▷R : boot.ci`
 or here `▷MATLAB : bootci`

5.9.1 Exercises

Exercise 262. (Check the solution using $\mathcal{R}$)

Check the solution of example.5.9 using function `boot.ci` of $\mathcal{R}$ and 999 resamples:

```
R> X <- c(25, 30, 35, 40, 45, 50, 55, 60, 65, 70)
R> library(boot)
R> mean.fun <- function(d, i) { m <- mean(X[i]) }
R> X.boot <- boot(X, mean.fun, R = 999)
R> boot.ci(X.boot, type = c("perc"))           # method = "percentile"

      BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
      Based on 999 bootstrap replicates
CALL :      boot.ci(boot.out = X.boot, type = c("perc"))
Intervals :  Level      Percentile
              95%       (38, 56 )
Calculations and Intervals on Original Scale
```

Exercise 263. (Example from askPython)

Ten university students were recruited to play the seminal video game pac-man. After equal amounts of practice the students were then asked to play the first stage and the time to completion was recorded.

The following data in seconds were recorded.

```
X = (63.1, 64.6, 63.4, 65.6, 63.4, 65.7, 64.6, 63.1, 65.9, 63.3)
```

Calculate the 95% confidence interval CI, which catch the true population mean.

◦ The solution using Python is here ▷ *AskPython*.

Exercise 264. (cytokine responses)

In these data we have reason to believe that the confidence intervals will not always be symmetric about the mean. Let the data be the cytokine responses (in %) for a sample measured by 16 laboratories. For the purposes of this example we will assume independence among the samples.

```
X = (0.4043, 0.5958, 0.5876, 0.5939, 0.6053, 0.4649, 0.5722, 0.6234,
     0.5208, 0.6117, 0.6137, 0.7397, 0.6981, 0.7757, 0.7596, 0.532)
```

Do a bootCI for X with 100 and 1000 iterations.

Exercise 265. (A 95% Confidence Interval for Movie Run Time, cf. *usu.edu*)

Do the example ‘Movie Run Time’ of the *usu.edu*.

Exercise 266. (cholesterol)

Serum cholesterol (mmol/L) measured on a sample of 86 stroke patients (data of Markus et al., 1995):

```
X=(3.7, 4.8, 5.4, 5.6, 6.1, 6.4, 7.0, 7.6, 8.7,
    3.8, 4.9, 5.4, 5.6, 6.1, 6.5, 7.0, 7.6, 8.9,
    3.8, 4.9, 5.5, 5.7, 6.1, 6.5, 7.1, 7.6, 9.3,
    4.4, 4.9, 5.5, 5.7, 6.2, 6.6, 7.1, 7.7, 9.5,
    4.5, 5.0, 5.5, 5.7, 6.3, 6.7, 7.2, 7.8, 10.2,
    4.5, 5.1, 5.6, 5.8, 6.3, 6.7, 7.3, 7.8, 10.4,
    4.5, 5.1, 5.6, 5.8, 6.4, 6.8, 7.4, 7.8,
    4.7, 5.2, 5.6, 5.9, 6.4, 6.8, 7.4, 8.2,
    4.7, 5.3, 5.6, 6.0, 6.4, 7.0, 7.5, 8.3,
    4.8, 5.3, 5.6, 6.1, 6.4, 7.0, 7.5, 8.6)
```

Calculate the CI using `bootci(X,100);`.

Exercise 267. (CI for correlation coefficient after EFRON, cf. [7, p.209 ff])

The bootstrap method was first applied in determining the non-parametric determination of a correlation coefficient. In 1973, Efron determined the average grades of all law school graduates (great point average, GPA) and at the law school admission test (LSAT) at 82 American law schools. By means of a random sample with 15 students, the following values were obtained.

EFRON used the data LSAT and GPA

```
LSAT = (576,635,558,578,666, 580,555,661,651,605, 653,575,545,572,594)
```

```
GPA = (3.39,3.30,2.81,3.03,3.44, 3.07,3.00,3.43,3.36,3.13,
       3.12,2.74,2.76,2.88,2.96)
```

◦ EFRON developed the idea of repeatedly drawing a random sample from the given sample to determine the variance. The correlation coefficient of 1000 samples he drew had a variance of 0.254.


- First calculate the parametric Pearson correlation coefficient ρ .
- Second calculate the Confidence Interval CI for ρ ; remember: we need the inverse of the standard normal distribution to calculate quantiles.
- Now define the confidence interval function for Pearson ρ : `CIrho(rho,n,alpha)` and calculate the CI for ρ at level 68.3% for $\rho = 0.7763$.
- Third calculate the CI using bootstrap method. Therefore define a specialized bootstrap function `bootCIrho(X,Y, rho)` and use `bootCIrho` on the Efron data.

The parametric range of deviation is 0.2347 and the non-parametric bootstrap range is $2 \cdot \sigma = 0.27$, which is a approx. 0.05 greater.

5.10 jackknife - The Jackknife method

Jackknife is a statistical resampling technique that estimates the bias and variance of a statistic θ by repeatedly leaving out one observation at a time from the original sample $X = (x_1, x_2, \dots, x_n)$. This creates n smaller subsamples X_i , each with one observation removed. The statistic of interest is then calculated for each subsample, and the resulting set of estimates θ_i is used to derive bias-corrected estimates and confidence intervals for the original statistic. ▷ GOOGLE AI : Jackknife resampling

Mental image



Visualization of **jackknife**: The ‘leave-one-out’ jackknife algorithm.

Instead of taking multiple samples directly from the population, we only

Figure 53: have a single, representative *sample*. We then generate ‘new’ samples by repeatedly leaving out one observation at a time from the original *sample*. – Picture found at ▷ WIKIPEDIA : Jackknife resampling

Procedure jackknife’s resampling

- The jackknife method in words:

A *jackknife sample* is a ‘leave-one-out’ resample of the data.

If there are n observations, then there are n jackknife samples, each of size n . If the original data are $X = (x_1, x_2, \dots, x_n)$, then the i th jackknife sample is $X_i := (x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_n)$. You then compute n jackknife replicates. A *jackknife replicate* is the statistic of interest (e.g. estimate of the standard error) computed on a jackknife resample.

- The jackknife method is summarized by the following steps:

a. Compute a statistic, θ , on the original sample of size n , e.g. $\theta = \bar{X} = \text{mean}(X)$.

b. For $i = 1$ to n , repeat the following:

▷ Leave out the i th observation x_i from X to form the new i th jackknife sample X_i .

- ▷ Compute the i th jackknife replicate statistic, θ_i , by computing the statistic on the i th jackknife sample X_i , e.g. $\theta_i := \bar{X}_i$ as estimate of mean $\bar{X}$ of X
- ▷ Compute the mean of the jackknife replicates: $\bar{\theta}_i := \frac{1}{n} \cdot \sum_{i=1}^n \theta_i$.
Estimate the bias $J_i := n\theta - (n-1)\bar{\theta}_i$, the so-called ‘pseudovariabls’.
- ▷ Estimate the standard error $SE_\theta := \sqrt{\frac{n-1}{n} \sum_{i=1}^n (\theta_i - \bar{\theta})^2}$
- **Test statistics** Let σ_J be the standard deviation of the *pseudovariabls*’s J_i .
Then we have

$$T := \frac{\sqrt{n} \cdot (J_\theta - \theta)}{\sqrt{\sigma_J^2}} \quad \text{is approximately } t\text{-distributed.}$$

Example The ‘leave-one-out’ jackknife

We follow the jackknife procedure by hand on the toy example $X = (1, 2, 3, 4)$.
Calculate the standard error of the pseudovariabls J_i .

Solution

1. $X = (1, 2, 3, 4), n = 4$. So we iterate 4-times.
2. $stats(y) := mean(y) = \bar{y}$, defines the statistics of interest.
3. $\theta := stats(X) = mean(X) = 2.5$.
4. calculate pseudovariabls J_i :

$$\begin{aligned} J_1: & X_1 = (2, 3, 4) \rightarrow \theta \cdot n - (n-1) \cdot stats(X_1) = 2.5 \times 4 - 3 \times 3 = 1 \\ J_2: & X_2 = (1, 3, 4) \rightarrow \theta \cdot n - (n-1) \cdot stats(X_2) = 2.5 \times 4 - 3 \times 8/3 = 2 \\ J_3: & 1 \ 2 \ 4 \rightarrow \dots = 3 \\ J_4: & 1 \ 2 \ 3 \rightarrow \dots = 4 \end{aligned}$$

$$5. \bar{J} = mean(J_1, J_2, J_3, J_4) = mean(1, 2, 3, 4) = 2.5$$

$$6. \text{bias} = \bar{J} - \theta = 4 - 4 = 0$$

$$7. SE = \sqrt{\frac{variance(J)}{4}} = 0.6455$$

◦ *Solution* and explanation of this example with REDUCE

```
| X := {1,2,3,4}$
| jack(X, 0.05);
```

SE	CI.l	CI.h
0.6455	4.5543	0.44574



View step-by-step solution *code*, run REDUCE.



View *jackknife* in action *code*, run REDUCE.

General Information

General mathematical information is here $\triangleright$ WIKIPEDIA : Jackknife resampling
 Syntax and semantic of the implementation is here $\triangleright$ MATLAB : jackknife

5.10.1 Exercises

Exercise 268. (nonparametric ‘leave-one-out’ jackknife algorithm in $\mathcal{R}$, cf. jackknifing)
 The corresponding R code is partially found at *influentialpoints.com*:

```
R> ## The nonparametric 'leave-one-out' jackknife algorithm in R
R> X      <- c(1,2,3,4)          # assign data
R> stats <- function(y){mean(y)} # define statistics of interest as mean
R> n      <- length(X)
R> theta <- stats(X)             # the total mean of all of X
R>        # now calculate 'pseudo-values' (pv)
R> pv = i <- 0; while(i < n){i<-i+1 ; pv[i] <- theta*n - (n-1)*stats(X[-i]);
R>        print(X[-i]); print(pv[i])} # show process

R> J      <- mean(pv)            # estimate statistic as mean of pseudo-values
R> bias   <- mean(pv) - th       # estimate bias
R> se     <- sqrt(var(pv)/n)     # estimate standard error SE
R> se

R> # calculate the confidence interval
R> alpha <- 1 - 0.95
R> ci    <- qt(alpha/2,n-1,lower.tail=FALSE)*se
R> CI    <- c(J - ci, J + ci)    # Confidence Interval
R> CI

      The nonparametric 'leave-one-out' jackknife
[1] 2 3 4 # leave 1 out in X=(1,2,3,4)
[1] 1
[2] 1 3 4 # leave 2 out
[2] 2
[3] 1 2 4 # leave 3 out
[3] 3
[4] 1 2 3 # leave 4 out
[4] 4
[5] 0.6454972 SE
[6] 0.4457398 4.5542602 CI
```

Exercise 269. (birth weights from cf. [7, p.212])

The birth weights of 72 boys and 48 girls born in the first half of 1981 at the Freiburg (Germany) clinic were examined.

The 72 boys data were

$X = (3700, 3800, 3040, 3080, 3170, \dots, 2880, 3150, 3160)$

The 48 girls data were

$Y = (3700, 3800, 3040, 3080, 3170, \dots, 2880, 3150, 3160)$

The 120 combined data is

$Z = (2990, 3700, 3800, 3040, 3080, 3350, 2930, 3355, 3360, 3170, \dots, 3220, 3850, 3080, 3440, 3840, 3320, 2640, 3020, 2850, 3700, 4630, 3670, 3400, 3580, 3270, 3360, 3600, 3220, 4060, 2670, 3550, 3690, 2510, 1980, 3700, 3730, 3170, 3580, 3720, 2910, 3500, 3430, 3510, 1520, 3420, 4630, 3470, 3400, 3290, 3250, 2420, 3580, 3100, 3480, 3830, 3050, 3680, 3330, 3650, 3200, 3810, 3350, 3390, 3550, 3350, 3200, 4190, 3050, 4020, 2880, 3020, 3010, 3560, 3030, 3200, 2830, 3120, 3000, 4010, 3050, 3980, 3530, 3245, 3415, 3660, 4350, 3740, 3220, 3350, 3350, 3780, 3620, 3310, 3630, 3290, 3800, 4130, 3380, 3100, 3430, 3620, 3060, 2610, 3560, 3720, 3030, 3630, 3400, 3580, 3700, 3470, 2810, 2330, 1960, 3300, 2880, 3660, 3150, 3160, 3300)$

From this sample of size 120, there are 120 jackknife samples generated and their mean is determined. This is intended to provide a non-parametric 95% confidence interval for the sample mean.

The available point estimate for the mean is 3352.8.

The non-parametric 95% confidence interval for the mean is [3267.9, 3437.7].

Task: verify the results above using

```
.. 3352.8    jack estimation of mean, theta = jac
.. [3267.88, 3437.71]    CI, our code
.. [3267.89, 3437.69]    CI jknife,  $\mathcal{R}$  code
.. [3297.3, 3458.1]    parametric CI
```

6 Bibliography

References

- [1] BEUCHER, O. (2005): *Wahrscheinlichkeitsrechnung und Statistik mit MATLAB*. Berlin: Springer.
- [2] BOGNAR, M.: *Probability Distribution Applets*.
url: <https://mabognar.github.io/apps/>
- [3] DALGAARD, P. (2002): *Introductory Statistics with R*. New York: Springer.
- [4] DIALEKT-PROJEKT (2000): *Statistik interaktiv! - Deskriptive Statistik*. Berlin: Springer.
- [5] GRABMEIER, J., KALTOFEN, E., WEISPFENNING, V., ED. (2003): *Computer Algebra Handbook*. Berlin: Springer.
- [6] HEIBERGER, R.M., HOLLAND, B. (2004): *Statistical Analysis and Data Display - An Intermediate Course with Examples in S-Plus, R and SAS*. New York: Springer.
- [7] HERRMANN, D. (1991): *Statistik in C*. Braunschweig: Vieweg.
- [8] LANE, D.M. et al. (2024): *Online Statistics Education*.
url: <https://onlinestatbook.com/2/index.html>
url: https://stats.libretexts.org/Bookshelves/Introductory_Statistics
url: https://onlinestatbook.com/Online_Statistics_Education.epub
- [9] LIBREOFFICE (2025): *LibreOffice Calc*.
url: <https://www.libreoffice.org/download/download-libreoffice/>
- [10] LINDNER, W. (2025): *Statistics with EIGENMATH*.
url: <https://lindnerdrwg.github.io/Statistics-with-Eigenmath.pdf>
- [11] LINDNER, W. (2026): *Statistics - A Maxima Companion*.
url: <https://lindnerdrwg.github.io/Statistics-with-Maxima.pdf>
- [12] LINDNER, W. (2025): *Memos, e.g. APOS Theory and A.C.E. Teaching Cycle*.
url: <https://lindnerdrwg.github.io/>
- [13] MACCALLUM, M.A.H. & WRIGHT, F.J. (1991): *Algebraic Computing with REDUCE*. Oxford: Oxford University Press.
- [14] PRESS, W.H., FLANERY, B.P., TEUCHOLSKY, S.A., VETTERLING, W.T. (1988): *Numerical Recipes in C*. Cambridge: Cambridge University Press.

- [15] ROLFES, H. (202?): *Significant.Statistics - Intro to Probability and Statistics*.
url: https://significantstatistics.com/index.php/Graduate_Level:_Intro_to_Probability_and_Statistics
- [16] ROSE, C. & SMITH, M.D. (2002): *Mathematical Statistics with MATHEMATICA*.
New York: Springer.
- [17] SOFTMAKER (2024): *PlanMaker Free 2024*.
url: https://www.freeoffice.com/en/?option=com_content
- [18] SPIEGEL, M.R. & STEPHENS, L.J. (^{4th}2011): *Statistics*. New York: McGraw-Hill.
- [19] STEEB, W-H. & LEWIEN, D. (1992): *Algorithms and Computations with REDUCE*.
Mannheim: B.I.-Wissenschaftsverlag.
- [20] STOUTEMYER, D. (2010 ff): *REDUCE Lessons. - Edited by F.J. WRIGHT*
url: <https://reduce-algebra.sourceforge.io/tutorials/lessons.zip>
- [21] VENABLES, W.N., RIPLEY, B.D. (^{3rd}1999): *Modern Applied Statistics with S-Plus*.
New York: Springer.
- [22] UEBERBERG, J. (1992): *Einführung in die Computeralgebra mit REDUCE*.
Mannheim: B.I.-Wissenschaftsverlag.
- [23] WEIGT, G. (2025): *EIGENMATH Command List*.
url: <https://georgeweigt.github.io/>
- [24] WEIGT, G. (2025): *EIGENMATH Manual*.
url: <https://georgeweigt.github.io/eigenmath.pdf>
- [25] WOOLLETT, E. (2024): *Statsitics with REDUCE*.
url: <https://home.csulb.edu/~woollett/stat.html>
- [26] WRIGHT, F.J. (2018): *REDUCE Calculus Tutorial*.
url: <https://reduce-algebra.sourceforge.io/tutorials/calculus-tutorial.php>
- [27] WRIGHT, F.J. (2018): *REDUCE Vectors and Linear Algebra Tutorial*.
url: <https://reduce-algebra.sourceforge.io/tutorials/veclinalg-tutorial.php>
- [28] YAU, C.: *Elementary Statistics with R*.
url: <https://www.r-tutor.com/elementary-statistics>
- [29] ZÖFEL, P. (1991): *Statistik verstehen*. Bonn: Addison-Wesley.
- [30] WESTER, M.J., ED. (1999): *Computer Algebra Systems - A Practical Guide*.
New York: McGraw-Hill. Wiley: Chichester.

The following texts are excerpts from the WESTER book:

- [31] GRÄBE, H.G. (1999): *About the Polynomial System Solve Facility of Axiom, Macsyma, Maple, Mathematica, MuPAD, and Reduce.*
url: <https://www.informatik.uni-leipzig.de/~graebe/ComputerAlgebra/Publications/WesterBook.pdf>
- [32] KOEPF, W. (1999): *Efficient Computation of Chebyshev Polynomials in CA.*
url: <https://www.mathematik.uni-kassel.de/~koepf/Publikationen/cheby.pdf>
- [33] WESTER, M.J. (1999): *A Critique of the Mathematical Abilities of CA Systems.*
url: <https://math.unm.edu/~wester/cas/book/Wester.pdf>
- [34] WESTER, M.J. (1999): *Computer Algebra Synonyms.*
url: <https://math.unm.edu/~wester/cas/synonyms.pdf>
- [35] WESTER, M.J. (1999): *REDUCE test files for 'A Critique ...'.*
url: https://math.unm.edu/~wester/cas_review.html



Dr. Wolfgang Lindner
Leichlingen, Germany
dr.w.g.Lindner@gmail.com
2026